Python code for Artificial Intelligence: Foundations of Computational Agents

David L. Poole and Alan K. Mackworth

Version 0.7.4 of May 28, 2018.

This document and all the code can be downloaded from http://artint.info/AIPython/ or from http://aipython.org

The authors and publisher of this book have used their best efforts in preparing this book. These efforts include the development, research and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or the documentation contained in this book. The author and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

Contents

Python for Artificial Intelligence

1.1   Why Python?

We use Python because Python programs can be close to pseudo-code. It is designed for humans to read.

Python is reasonably efficient. Efficiency is usually not a problem for small examples. If your Python code is not efficient enough, a general procedure to improve it is to find out what is taking most the time, and implement just that part more efficiently in some lower-level language. Most of these lower-level languages interoperate with Python nicely. This will result in much less programming and more efficient code (because you will have more time to optimize) than writing everything in a low-level language. You will not have to do that for the code here if you are using it for course projects.

1.2   Getting Python

You need Python 3 (`http://python.org/`) and matplotlib (`http://matplotlib.org/`) that runs with Python 3. This code is *not* compatible with Python 2 (e.g., with Python 2.7).

Download and istall the latest Python 3 release from `http://python.org/`. This should also install *pip3*. You can install matplotlib using

```
pip3 install matplotlib
```

in a terminal shell (not in Python). That should "just work". If not, try using `pip` instead of `pip3`.

The command `python` or `python3` should then start the interactive python shell. You can quit Python with a control-D or with `quit()`.

To upgrade matplotlib to the latest version (which you should do if you install a new version of Python) do:

```
pip3 install --upgrade matplotlib
```

We recommend using the enhanced interactive python **ipython** (http://ipython.org/). To install ipython after you have installed python do:

```
pip3 install ipython
```

1.3   Running Python

We assume that everything is done with an interactive Python shell. You can either do this with an IDE, such as IDLE that comes with standard Python distributions, or just running ipython3 (or perhaps just ipython) from a shell.

Here we describe the most simple version that uses no IDE. If you download the zip file, and cd to the "aipython" folder where the .py files are, you should be able to do the following, with user input following : . The first ipython3 command is in the operating system shell (note that the -i is important to enter interactive mode):

```
$ ipython3 -i searchGeneric.py
Python 3.6.5 (v3.6.5:f59c0932b4, Mar 28 2018, 05:52:31)
Type 'copyright', 'credits' or 'license' for more information
IPython 6.2.1 -- An enhanced Interactive Python. Type '?' for help.
Testing problem 1:
7 paths have been expanded and 4 paths remain in the frontier
Path found: a --> b --> c --> d --> g
Passed unit test

In [1]: searcher2 = AStarSearcher(searchProblem.acyclic_delivery_problem) #A*

In [2]: searcher2.search()  # find first path
16 paths have been expanded and 5 paths remain in the frontier
Out[2]: o103 --> o109 --> o119 --> o123 --> r123

In [3]: searcher2.search()  # find next path
21 paths have been expanded and 6 paths remain in the frontier
Out[3]: o103 --> b3 --> b4 --> o109 --> o119 --> o123 --> r123

In [4]: searcher2.search()  # find next path
28 paths have been expanded and 5 paths remain in the frontier
Out[4]: o103 --> b3 --> b1 --> b2 --> b4 --> o109 --> o119 --> o123 --> r123

In [5]: searcher2.search()  # find next path
No (more) solutions. Total of 33 paths expanded.
```

```
In [6]:
```

You can then interact at the last prompt.

There are many textbooks for Python. The best source of information about python is `https://www.python.org/`. We will be using Python 3; please download the latest release. The documentation is at `https://docs.python.org/3/`.

The rest of this chapter is about what is special about the code for AI tools. We will only use the Standard Python Library and matplotlib. All of the exercises can be done (and should be done) without using other libraries; the aim is for you to spend your time thinking about how to solve the problem rather than searching for pre-existing solutions.

1.4 Pitfalls

It is important to know when side effects occur. Often AI programs consider what would happen or what may have happened. In many such cases, we don't want side effects. When an agent acts in the world, side effects are appropriate.

In Python, you need to be careful to understand side effects. For example, the inexpensive function to add an element to a list, namely *append*, changes the list. In a functional language like Lisp, adding a new element to a list, without changing the original list, is a cheap operation. For example if *x* is a list containing *n* elements, adding an extra element to the list in Python (using *append*) is fast, but it has the side effect of changing the list *x*. To construct a new list that contains the elements of *x* plus a new element, without changing the value of *x*, entails copying the list, or using a different representation for lists. In the searching code, we will use a different representation for lists for this reason.

1.5 Features of Python

1.5.1 Lists, Tuples, Sets, Dictionaries and Comprehensions

We make extensive uses of lists, tuples, sets and dictionaries (dicts). See
`https://docs.python.org/3/library/stdtypes.html`

One of the nice features of Python is the use of list comprehensions (and also tuple, set and dictionary comprehensions).

>   (*fe* for *e* in *iter* if *cond*)

enumerates the values *fe* for each *e* in *iter* for which *cond* is true. The "if cond" part is optional, but the "for" and "in" are not optional. Here *e* has to be a variable, *iter* is an iterator, which can generate a stream of data, such as a list, a set, a range object (to enumerate integers between ranges) or a file. *cond*

is an expression that evaluates to either True or False for each *e*, and *fe* is an expression that will be evaluated for each value of *e* for which *cond* returns *True*.

The result can go in a list or used in another iteration, or can be called directly using *next*. The procedure *next* takes an iterator returns the next element (advancing the iterator) and raises a StopIteration exception if there is no next element. The following shows a simple example, where user input is prepended with >>>

```
>>> [e*e for e in range(20) if e%2==0]
[0, 4, 16, 36, 64, 100, 144, 196, 256, 324]
>>> a = (e*e for e in range(20) if e%2==0)
>>> next(a)
0
>>> next(a)
4
>>> next(a)
16
>>> list(a)
[36, 64, 100, 144, 196, 256, 324]
>>> next(a)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
StopIteration
```

Notice how *list*(*a*) continued on the enumeration, and got to the end of it.

Comprehensions can also be used for dictionaries. The following code creates an index for list *a*:

```
>>> a = ["a","f","bar","b","a","aaaaa"]
>>> ind = {a[i]:i for i in range(len(a))}
>>> ind
{'a': 4, 'f': 1, 'bar': 2, 'b': 3, 'aaaaa': 5}
>>> ind['b']
3
```

which means that 'b' is the 3rd element of the list.

The assignment of *ind* could have also be written as:

```
>>> ind = {val:i for (i,val) in enumerate(a)}
```

where *enumerate* returns an iterator of (*index*, *value*) pairs.

1.5.2 Functions as first-class objects

Python can create lists and other data structures that contain functions. There is an issue that tricks many newcomers to Python. For a local variable in a function, the function uses the last value of the variable when the function is

called, not the value of the variable when the function was defined (this is called "late binding"). This means if you want to use the value a variable has when the function is created, you need to save the current value of that variable. Whereas Python uses "late binding" by default, the alternative that newcomers often expect is "early binding", where a function uses the value a variable had when the function was defined, can be easily implemented.

Consider the following programs designed to create a list of 5 functions, where the *i*th function in the list is meant to add *i* to its argument:[1]

_____ pythonDemo.py — Some tricky examples _____

```python
11  fun_list1 = []
12  for i in range(5):
13      def fun1(e):
14          return e+i
15      fun_list1.append(fun1)
16
17  fun_list2 = []
18  for i in range(5):
19      def fun2(e,iv=i):
20          return e+iv
21      fun_list2.append(fun2)
22
23  fun_list3 = [lambda e: e+i for i in range(5)]
24
25  fun_list4 = [lambda e,iv=i: e+iv for i in range(5)]
26
27  i=56
```

Try to predict, and then test to see the output, of the output of the following calls, remembering that the function uses the latest value of any variable that is not bound in the function call:

_____ pythonDemo.py — (continued) _____

```python
29  # in Shell do
30  ## ipython -i pythonDemo.py
31  # Try these (copy text after the comment symbol and paste in the Python prompt):
32  # print([f(10) for f in fun_list1])
33  # print([f(10) for f in fun_list2])
34  # print([f(10) for f in fun_list3])
35  # print([f(10) for f in fun_list4])
```

In the first for-loop, the function *fun* uses *i*, whose value is the last value it was assigned. In the second loop, the function *fun*2 uses *iv*. There is a separate *iv* variable for each function, and its value is the value of *i* when the function was defined. Thus *fun*1 uses late binding, and *fun*2 uses early binding. *fun_list*3 and *fun_list*4 are equivalent to the first two (except *fun_list*4 uses a different *i* variable).

_____

[1]Numbered lines are Python code available in the code-directory, `aipython`. The name of the file is given in the gray text above the listing. The numbers correspond to the line numbers in that file.

One of the advantages of using the embedded definitions (as in *fun*1 and *fun*2 above) over the lambda is that is it possible to add a `__doc__` string, which is the standard for documenting functions in Python, to the embedded definitions.

1.5.3 Generators and Coroutines

Python has generators which can be used for a form of coroutines.

The *yield* command returns a value that is obtained with *next*. It is typically used to enumerate the values for a *for* loop or in generators.

A version of the built-in *range*, with 2 or 3 arguments (and positive steps) can be implemented as:

_____pythonDemo.py — (continued)_____

```
37  def myrange(start, stop, step=1):
38      """enumerates the values from start in steps of size step that are
39      less than stop.
40      """
41      assert step>0, "only positive steps implemented in myrange"
42      i = start
43      while i<stop:
44          yield i
45          i += step
46
47  print("myrange(2,30,3):",list(myrange(2,30,3)))
```

Note that the built-in *range* is unconventional in how it handles a single argument, as the single argument acts as the second argument of the function. Note also that the built-in range also allows for indexing (e.g., $range(2, 30, 3)[2]$ returns 8), which the above implementation does not. However *myrange* also works for floats, which the built-in range does not.

Exercise 1.1 Implement a version of *myrange* that acts like the built-in version when there is a single argument. (Hint: make the second argument have a default value that can be recognized in the function.)

Yield can be used to generate the same sequence of values as in the example of Section 1.5.1:

_____pythonDemo.py — (continued)_____

```
49  def ga(n):
50      """generates square of even nonnegative integers less than n"""
51      for e in range(n):
52          if e%2==0:
53              yield e*e
54  a = ga(20)
```

The sequence of next(a), and list(a) gives exactly the same results as the comprehension in Section 1.5.1.

It is straightforward to write a version of the built-in *enumerate*. Let's call it *myenumerate*:

---
pythonDemo.py — (continued)

```
56  def myenumerate(enum):
57      for i in range(len(enum)):
58          yield i,enum[i]
```

Exercise 1.2 Write a version of *enumerate* where the only iteration is "for val in enum". Hint: keep track of the index.

1.6 Useful Libraries

1.6.1 Timing Code

In order to compare algorithms, we often want to compute how long a program takes; this is called the **runtime** of the program. The most straightforward way to compute runtime is to use *time.perf_counter()*, as in:

```
import time
start_time = time.perf_counter()
compute_for_a_while()
end_time = time.perf_counter()
print("Time:", end_time - start_time, "seconds")
```

If this time is very small (say less than 0.2 second), it is probably very inaccurate, and it may be better to run your code many times to get a more accurate count. For this you can use *timeit* (https://docs.python.org/3/library/timeit.html). To use timeit to time the call to *foo.bar(aaa)* use:

```
import timeit
time = timeit.timeit("foo.bar(aaa)",
                     setup="from __main__ import foo,aaa", number=100)
```

The setup is needed so that Python can find the meaning of the names in the string that is called. This returns the number of seconds to execute *foo.bar(aaa)* 100 times. The variable *number* should be set so that the runtime is at least 0.2 seconds.

You should not trust a single measurement as that can be confounded by interference from other processes. *timeit.repeat* can be used for running *timit* a few (say 3) times. Usually the minimum time is the one to report, but you should be explicit and explain what you are reporting.

1.6.2 Plotting: Matplotlib

The standard plotting for Python is matplotlib (http://matplotlib.org/). We will use the most basic plotting using the pyplot interface.

Here is a simple example that uses everything we will use.

```
60  import matplotlib.pyplot as plt
61
62  def myplot(min,max,step,fun1,fun2):
63      plt.ion() # make it interactive
64      plt.xlabel("The x axis")
65      plt.ylabel("The y axis")
66      plt.xscale('linear') # Makes a 'log' or 'linear' scale
67      xvalues = range(min,max,step)
68      plt.plot(xvalues,[fun1(x) for x in xvalues],
69                  label="The first fun")
70      plt.plot(xvalues,[fun2(x) for x in xvalues], linestyle='--',color='k',
71                  label=fun2.__doc__) # use the doc string of the function
72      plt.legend(loc="upper right") # display the legend
73
74  def slin(x):
75      """y=2x+7"""
76      return 2*x+7
77  def sqfun(x):
78      """y=(x-40)^2/10-20"""
79      return (x-40)**2/10-20
80
81  # Try the following:
82  # from pythonDemo import myplot, slin, sqfun
83  # import matplotlib.pyplot as plt
84  # myplot(0,100,1,slin,sqfun)
85  # plt.legend(loc="best")
86  # import math
87  # plt.plot([41+40*math.cos(th/10) for th in range(50)],
88  #          [100+100*math.sin(th/10) for th in range(50)])
89  # plt.text(40,100,"ellipse?")
90  # plt.xscale('log')
```

At the end of the code are some commented-out commands you should try in
interactive mode. Cut from the file and paste into Python (and remember to
remove the comments symbol and leading space).

1.7   Utilities

1.7.1   Display

In this distribution, to keep things simple and to only use standard Python, we
use a text-oriented tracing of the code. A graphical depiction of the code could
override the definition of *display* (but we leave it as a project).

The method *self*.*display* is used to trace the program. Any call

$$self.display(level, to\_print \ldots)$$

where the level is less than or equal to the value for *max_display_level* will be printed. The *to_print . . .* can be anything that is accepted by the built-in *print* (including any keyword arguments).

The definition of *display* is:

_____ utilities.py — AIFCA utilities _____
```python
11  class Displayable(object):
12      max_display_level = 1 # can be overridden in subclasses
13
14      def display(self,level,*args,**nargs):
15          """print the arguments if level is less than or equal to the
16          current max_display_level.
17          level is an integer.
18          the other arguments are whatever arguments print can take.
19          """
20          if level <= self.max_display_level:
21              print(*args, **nargs) ##if error you are using Python2 not Python3
```

Note that *args* gets a tuple of the positional arguments, and *nargs* gets a dictionary of the keyword arguments). This will not work in Python 2, and will give an error.

Any class that wants to use *display* can be made a subclass of *Displayable*.

To change the maximum display level to say 3, for a class do:

> *Classname.max_display_level* = 3

which will make calls to *display* in that class print when the value of *level* is less than-or-equal to 3. The default display level is 1. It can also be changed for individual objects (the object value overrides the class value).

The value of *max_display_level* by convention is:

0 display nothing

1 display solutions (nothing that happens repeatedly)

2 also display the values as they change (little detail through a loop)

3 also display more details

4 and above even more detail

In order to implement more sophisticated visualizations of the algorithm, we add a **visualize** "decorator" to the methods to be visualized. The following code ignores the decorator:

_____ utilities.py — (continued) _____
```python
23  def visualize(func):
24      return func
```

1.7.2   Argmax

Python has a built-in *max* function that takes a generator (or a list or set) and returns the maximum value. The *argmax* method returns the index of an element that has the maximum value. If there are multiple elements with the maximum value, one if the indexes to that value is returned at random. This assumes a generator of (*element*, *value*) pairs, as for example is generated by the built-in *enumerate*.

───────────────────── utilities.py — (continued) ─────────────────────

```
25  import random
26
27  def argmax(gen):
28      """gen is a generator of (element,value) pairs, where value is a real.
29      argmax returns an element with maximal value.
30      If there are multiple elements with the max value, one is returned at random.
31      """
32      maxv = float('-Infinity')    # negative infinity
33      maxvals = []     # list of maximal elements
34      for (e,v) in gen:
35          if v>maxv:
36              maxvals,maxv = [e], v
37          elif v==maxv:
38              maxvals.append(e)
39      return random.choice(maxvals)
40
41  # Try:
42  # argmax(enumerate([1,6,3,77,3,55,23]))
```

Exercise 1.3 Change argmax to have an optinal argument that specifies whether you want the "first", "last" or a "random" index of the maximum value returned. If you want the first or the last, you don't need to keep a list of the maximum elements.

1.7.3   Probability

For many of the simulations, we want to make a variable True with some probability. *flip*(*p*) returns True with probability *p*, and otherwise returns False.

───────────────────── utilities.py — (continued) ─────────────────────

```
44  def flip(prob):
45      """return true with probability prob"""
46      return random.random() < prob
```

1.7.4   Dictionary Union

The function *dict_union*(*d*1, *d*2) returns the union of dictionaries *d*1 and *d*2. If the values for the keys conflict, the values in *d*2 are used. This is similar to *dict*(*d*1, ∗ ∗ *d*2), but that only works when the keys of *d*2 are strings.

—————————— utilities.py — (continued) ——————————

```
48  def dict_union(d1,d2):
49      """returns a dictionary that contains the keys of d1 and d2.
50      The value for each key that is in d2 is the value from d2,
51      otherwise it is the value from d1.
52      This does not have side effects.
53      """
54      d = dict(d1)   # copy d1
55      d.update(d2)
56      return d
```

1.8   Testing Code

It is important to test code early and test it often. We include a simple form of
unit tests. The value of the current module is in `__name__` and if the module is
run at the top-level, it's value is `"__main__"`. See `https://docs.python.org/3/`
`library/__main__.html`.

The following code tests `argmax` and `dict_union`, but only when if `utilities`
is loaded in the top-level. If it is loaded in a module the test code is not run.

In your code you should do more substantial testing than we do here, in
particular testing the boundary cases.

—————————— utilities.py — (continued) ——————————

```
58  def test():
59      """Test part of utilities"""
60      assert argmax(enumerate([1,6,55,3,55,23])) in [2,4]
61      assert dict_union({1:4, 2:5, 3:4},{5:7, 2:9}) == {1:4, 2:9, 3:4, 5:7}
62      print("Passed unit test in utilities")
63
64  if __name__ == "__main__":
65      test()
```

Chapter 2

Agents and Control

This implements the controllers described in Chapter 2.

In this version the higher-levels call the lower-levels. A more sophisticated version may have them run concurrently (either as coroutines or in parallel). The higher-levels calling the lower-level works in simulated environments when there is a single agent, and where the lower-level are written to make sure they return (and don't go on forever), and the higher level doesn't take too long (as the lower-levels will wait until called again).

2.1 Representing Agents and Environments

An agent observes the world, and carries out actions in the environment, it also has an internal state that it updates. The environment takes in actions of the agents, updates it internal state and returns the percepts.

In this implementation, the state of the agent and the state of the environment are represented using standard Python variables, which are updated as the state changes. The percepts and the actions are represented as variable-value dictionaries.

An agent implements the $go(n)$ method, where n is an integer. This means that the agent should run for n time steps.

In the following code `raise NotImplementedError()` is a way to specify an abstract method that needs to be overidden in any implemented agent or environment.

_____ agents.py — Agent and Controllers _____
```
11  import random
12
13  class Agent(object):
14      def __init__(self,env):
```

```
15          """set up the agent"""
16          self.env=env
17
18      def go(self,n):
19          """acts for n time steps"""
20          raise NotImplementedError("go") # abstract method
```

The environment implements a *do*(*action*) method where *action* is a variable-value dictionary. This returns a percept, which is also a variable-value dictionary. The use of dictionaries allows for structured actions and percepts.

Note that *Environment* is a subclass of *Displayable* so that it can use the *display* method described in Section 1.7.1.

_____ agents.py — (continued) _____

```
22  from utilities import Displayable
23  class Environment(Displayable):
24      def initial_percepts(self):
25          """returns the initial percepts for the agent"""
26          raise NotImplementedError("initial_percepts") # abstract method
27
28      def do(self,action):
29          """does the action in the environment
30          returns the next percept """
31          raise NotImplementedError("do") # abstract method
```

2.2   Paper buying agent and environment

> To run the demo, in folder "aipython", load "agents.py", using e.g.,
> `ipython -i agents.py`, and copy and paste the commented-out
> commands at the bottom of that file. This requires Python 3 with
> `matplotlib`.

This is an implementation of the paper buying example.

2.2.1   The Environment

The environment state is given in terms of the *time* and the amount of paper in *stock*. It also remembers the in-stock history and the price history. The percepts are the price and the amount of paper in stock. The action of the agent is the number to buy.

Here we assume that the prices are obtained from the *prices* list plus a random integer in range [0, *max_price_addon*) plus a linear "inflation". The agent cannot access the price model; it just observes the prices and the amount in stock.

_____ agents.py — (continued) _____

```
33  class TP_env(Environment):
```

```
34    prices = [234, 234, 234, 234, 255, 255, 275, 275, 211, 211, 211,
35    234, 234, 234, 234, 199, 199, 275, 275, 234, 234, 234, 234, 255,
36    255, 260, 260, 265, 265, 265, 265, 270, 270, 255, 255, 260, 260,
37    265, 265, 150, 150, 265, 265, 270, 270, 255, 255, 260, 260, 265,
38    265, 265, 265, 270, 270, 211, 211, 255, 255, 260, 260, 265, 265,
39    260, 265, 270, 270, 205, 255, 255, 260, 260, 265, 265, 265, 265,
40    270, 270]
41    max_price_addon = 20 # maximum of random value added to get price
42
43    def __init__(self):
44        """paper buying agent"""
45        self.time=0
46        self.stock=20
47        self.stock_history = [] # memory of the stock history
48        self.price_history = [] # memory of the price history
49
50    def initial_percepts(self):
51        """return initial percepts"""
52        self.stock_history.append(self.stock)
53        price = self.prices[0]+random.randrange(self.max_price_addon)
54        self.price_history.append(price)
55        return {'price': price,
56                'instock': self.stock}
57
58    def do(self, action):
59        """does action (buy) and returns percepts (price and instock)"""
60        used = pick_from_dist({6:0.1, 5:0.1, 4:0.2, 3:0.3, 2:0.2, 1:0.1})
61        bought = action['buy']
62        self.stock = self.stock+bought-used
63        self.stock_history.append(self.stock)
64        self.time += 1
65        price = (self.prices[self.time%len(self.prices)] # repeating pattern
66                +random.randrange(self.max_price_addon) # plus randomness
67                +self.time//2)                          # plus inflation
68        self.price_history.append(price)
69        return {'price': price,
70                'instock': self.stock}
```

The *pick_from_dist* method takes in a *item : probability* dictionary, and returns one of the items in proportion to its probability.

———————————— _agents.py — (continued) ————————————

```
72  def pick_from_dist(item_prob_dist):
73      """ returns a value from a distribution.
74      item_prob_dist is an item:probability dictionary, where the
75          probabilities sum to 1.
76      returns an item chosen in proportion to its probability
77      """
78      ranreal = random.random()
79      for (it,prob) in item_prob_dist.items():
80          if ranreal < prob:
```

```
81            return it
82        else:
83            ranreal -= prob
84    raise RuntimeError(str(item_prob_dist)+" is not a probability distribution")
```

2.2.2   The Agent

The agent does not have access to the price model but can only observe the current price and the amount in stock. It has to decide how much to buy.

The belief state of the agent is an estimate of the average price of the paper, and the total amount of money the agent has spent.

_____ agents.py — (continued) _____
```
86  class TP_agent(Agent):
87      def __init__(self, env):
88          self.env = env
89          self.spent = 0
90          percepts = env.initial_percepts()
91          self.ave = self.last_price = percepts['price']
92          self.instock = percepts['instock']
93
94      def go(self, n):
95          """go for n time steps
96          """
97          for i in range(n):
98              if self.last_price < 0.9*self.ave and self.instock < 60:
99                  tobuy = 48
100             elif self.instock < 12:
101                 tobuy = 12
102             else:
103                 tobuy = 0
104             self.spent += tobuy*self.last_price
105             percepts = env.do({'buy': tobuy})
106             self.last_price = percepts['price']
107             self.ave = self.ave+(self.last_price-self.ave)*0.05
108             self.instock = percepts['instock']
```

Set up an environment and an agent. Uncomment the last lines to run the agent for 90 steps, and determine the average amount spent.

_____ agents.py — (continued) _____
```
110 env = TP_env()
111 ag = TP_agent(env)
112 #ag.go(90)
113 #ag.spent/env.time ## average spent per time period
```

2.2.3   Plotting

The following plots the price and number in stock history:

_____agents.py — (continued) _____

```
115  import matplotlib.pyplot as plt
116
117  class Plot_prices(object):
118      """Set up the plot for history of price and number in stock"""
119      def __init__(self, ag,env):
120          self.ag = ag
121          self.env = env
122          plt.ion()
123          plt.xlabel("Time")
124          plt.ylabel("Number in stock.                    Price.")
125
126      def plot_run(self):
127          """plot history of price and instock"""
128          num = len(env.stock_history)
129          plt.plot(range(num),env.stock_history,label="In stock")
130          plt.plot(range(num),env.price_history,label="Price")
131          #plt.legend(loc="upper left")
132          plt.draw()
133
134  # pl = Plot_prices(ag,env)
135  # ag.go(90); pl.plot_run()
```

2.3 Hierarchical Controller

> To run the hierarchical controller, in folder "aipython", load "agentTop.py", using e.g., ipython -i agentTop.py, and copy and paste the commands near the bottom of that file. This requires Python 3 with matplotlib.

In this implementation, each layer, including the top layer, implements the environment class, because each layer is seen as an environment from the layer above.

We arbitrarily divide the environment and the body, so that the environment just defines the walls, and the body includes everything to do with the agent. Note that the named locations are part of the (top-level of the) agent, not part of the environment, although they could have been.

2.3.1 Environment

The environment defines the walls.

_____agentEnv.py — Agent environment _____

```
11  import math
12  from agents import Environment
13
14  class Rob_env(Environment):
```

```
15    def __init__(self,walls = {}):
16        """walls is a set of line segments
17              where each line segment is of the form ((x0,y0),(x1,y1))
18        """
19        self.walls = walls
```

2.3.2   Body

The body defines everything about the agent body.

───────── *agentEnv.py — (continued)* ─────────
```
21  import math
22  from agents import Environment
23  import matplotlib.pyplot as plt
24  import time
25
26  class Rob_body(Environment):
27      def __init__(self, env, init_pos=(0,0,90)):
28          """ env is the current environment
29          init_pos is a triple of (x-position, y-position, direction)
30              direction is in degrees; 0 is to right, 90 is straight-up, etc
31          """
32          self.env = env
33          self.rob_x, self.rob_y, self.rob_dir = init_pos
34          self.turning_angle = 18 # degrees that a left makes
35          self.whisker_length = 6 # length of the whisker
36          self.whisker_angle = 30 # angle of whisker relative to robot
37          self.crashed = False
38          # The following control how it is plotted
39          self.plotting = True    # whether the trace is being plotted
40          self.sleep_time = 0.05  # time between actions (for real-time plotting)
41          # The following are data structures maintained:
42          self.history = [(self.rob_x, self.rob_y)] # history of (x,y) positions
43          self.wall_history = []  # history of hitting the wall
44
45      def percepts(self):
46          return {'rob_x_pos':self.rob_x, 'rob_y_pos':self.rob_y,
47                  'rob_dir':self.rob_dir, 'whisker':self.whisker() , 'crashed':self.crashed}
48      initial_percepts = percepts # use percept function for initial percepts too
49
50      def do(self,action):
51          """ action is {'steer':direction}
52          direction is 'left', 'right' or 'straight'
53          """
54          if self.crashed:
55              return self.percepts()
56          direction = action['steer']
57          compass_deriv = {'left':1,'straight':0,'right':-1}[direction]*self.turning_angle
58          self.rob_dir = (self.rob_dir + compass_deriv +360)%360 # make in range [0,360)
59          rob_x_new = self.rob_x + math.cos(self.rob_dir*math.pi/180)
```

```
60          rob_y_new = self.rob_y + math.sin(self.rob_dir*math.pi/180)
61          path = ((self.rob_x,self.rob_y),(rob_x_new,rob_y_new))
62          if any(line_segments_intersect(path,wall) for wall in self.env.walls):
63              self.crashed = True
64              if self.plotting:
65                  plt.plot([self.rob_x],[self.rob_y],"r*",markersize=20.0)
66                  plt.draw()
67          self.rob_x, self.rob_y = rob_x_new, rob_y_new
68          self.history.append((self.rob_x, self.rob_y))
69          if self.plotting and not self.crashed:
70              plt.plot([self.rob_x],[self.rob_y],"go")
71              plt.draw()
72              plt.pause(self.sleep_time)
73          return self.percepts()
```

This detects if the whisker and the wall intersect. It's value is returned as a percept.

<div align="center">_____agentEnv.py — (continued) _____</div>

```
75      def whisker(self):
76          """returns true whenever the whisker sensor intersects with a wall
77          """
78          whisk_ang_world = (self.rob_dir-self.whisker_angle)*math.pi/180
79              # angle in radians in world coordinates
80          wx = self.rob_x + self.whisker_length * math.cos(whisk_ang_world)
81          wy = self.rob_y + self.whisker_length * math.sin(whisk_ang_world)
82          whisker_line = ((self.rob_x,self.rob_y),(wx,wy))
83          hit = any(line_segments_intersect(whisker_line,wall)
84                      for wall in self.env.walls)
85          if hit:
86              self.wall_history.append((self.rob_x, self.rob_y))
87              if self.plotting:
88                  plt.plot([self.rob_x],[self.rob_y],"ro")
89                  plt.draw()
90          return hit
91
92  def line_segments_intersect(linea,lineb):
93      """returns true if the line segments, linea and lineb intersect.
94      A line segment is represented as a pair of points.
95      A point is represented as a (x,y) pair.
96      """
97      ((x0a,y0a),(x1a,y1a)) = linea
98      ((x0b,y0b),(x1b,y1b)) = lineb
99      da, db = x1a-x0a, x1b-x0b
100     ea, eb = y1a-y0a, y1b-y0b
101     denom = db*ea-eb*da
102     if denom==0:   # line segments are parallel
103         return False
104     cb = (da*(y0b-y0a)-ea*(x0b-x0a))/denom # position along line b
105     if cb<0 or cb>1:
106         return False
```

```
107        ca = (db*(y0b-y0a)-eb*(x0b-x0a))/denom # position along line a
108        return 0<=ca<=1
109
110  # Test cases:
111  # assert line_segments_intersect(((0,0),(1,1)),((1,0),(0,1)))
112  # assert not line_segments_intersect(((0,0),(1,1)),((1,0),(0.6,0.4)))
113  # assert line_segments_intersect(((0,0),(1,1)),((1,0),(0.4,0.6)))
```

2.3.3 Middle Layer

The middle layer acts like both a controller (for the environment layer) and an environment for the upper layer. It has to tell the environment how to steer. Thus it calls *env.do*(·). It also is told the position to go to and the timeout. Thus it also has to implement *do*(·).

_____agentMiddle.py — Middle Layer _____
```
11  from agents import Environment
12  import math
13
14  class Rob_middle_layer(Environment):
15      def __init__(self,env):
16          self.env=env
17          self.percepts = env.initial_percepts()
18          self.straight_angle = 11 # angle that is close enough to straight ahead
19          self.close_threshold = 2 # distance that is close enough to arrived
20          self.close_threshold_squared = self.close_threshold**2 # just compute it once
21
22      def initial_percepts(self):
23          return {}
24
25      def do(self, action):
26          """action is {'go_to':target_pos,'timeout':timeout}
27          target_pos is (x,y) pair
28          timeout is the number of steps to try
29          returns {'arrived':True} when arrived is true
30              or {'arrived':False} if it reached the timeout
31          """
32          if 'timeout' in action:
33              remaining = action['timeout']
34          else:
35              remaining = -1  # will never reach 0
36          target_pos = action['go_to']
37          arrived = self.close_enough(target_pos)
38          while not arrived and remaining != 0:
39              self.percepts = self.env.do({"steer":self.steer(target_pos)})
40              remaining -= 1
41              arrived = self.close_enough(target_pos)
42          return {'arrived':arrived}
```

This determines how to steer depending on whether the goal is to the right or the left of where the robot is facing.

```
                          ___agentMiddle.py — (continued)___
44      def steer(self,target_pos):
45          if self.percepts['whisker']:
46              self.display(3,'whisker on', self.percepts)
47              return "left"
48          else:
49              gx,gy = target_pos
50              rx,ry = self.percepts['rob_x_pos'],self.percepts['rob_y_pos']
51              goal_dir = math.acos((gx-rx)/math.sqrt((gx-rx)*(gx-rx)
52                                              +(gy-ry)*(gy-ry)))*180/math.pi
53              if ry>gy:
54                  goal_dir = -goal_dir
55              goal_from_rob = (goal_dir - self.percepts['rob_dir']+540)%360-180
56              assert -180 < goal_from_rob <= 180
57              if goal_from_rob > self.straight_angle:
58                  return "left"
59              elif goal_from_rob < -self.straight_angle:
60                  return "right"
61              else:
62                  return "straight"
63
64      def close_enough(self,target_pos):
65          gx,gy = target_pos
66          rx,ry = self.percepts['rob_x_pos'],self.percepts['rob_y_pos']
67          return (gx-rx)**2 + (gy-ry)**2 <= self.close_threshold_squared
```

2.3.4   Top Layer

The top layer treats the middle layer as its environment. Note that the top layer is an environment for us to tell it what to visit.

```
                          ___agentTop.py — Top Layer___
11  from agentMiddle import Rob_middle_layer
12  from agents import Environment
13
14  class Rob_top_layer(Environment):
15      def __init__(self, middle, timeout=200, locations = {'mail':(-5,10),
16                      'o103':(50,10), 'o109':(100,10),'storage':(101,51)} ):
17          """middle is the middle layer
18          timeout is the number of steps the middle layer goes before giving up
19          locations is a loc:pos dictionary
20              where loc is a named location, and pos is an (x,y) position.
21          """
22          self.middle = middle
23          self.timeout = timeout # number of steps before the middle layer should give up
24          self.locations = locations
25
```

```
26    def do(self,plan):
27        """carry out actions.
28        actions is of the form {'visit':list_of_locations}
29        It visits the locations in turn.
30        """
31        to_do = plan['visit']
32        for loc in to_do:
33            position = self.locations[loc]
34            arrived = self.middle.do({'go_to':position, 'timeout':self.timeout})
35            self.display(1,"Arrived at",loc,arrived)
```

2.3.5   Plotting

The following is used to plot the locations, the walls and (eventually) the movement of the robot. It can either plot the movement if the robot as it is going (with the default *env.plotting = True*), or not plot it as it is going (setting *env.plotting = False*; in this case the trace can be plotted using *pl.plot_run()*).

```
_____ agentTop.py — (continued) _____
37  import matplotlib.pyplot as plt
38
39  class Plot_env(object):
40      def __init__(self, body,top):
41          """sets up the plot
42          """
43          self.body = body
44          plt.ion()
45          plt.clf()
46          plt.axes().set_aspect('equal')
47          for wall in body.env.walls:
48              ((x0,y0),(x1,y1)) = wall
49              plt.plot([x0,x1],[y0,y1],"-k",linewidth=3)
50          for loc in top.locations:
51              (x,y) = top.locations[loc]
52              plt.plot([x],[y],"k<")
53              plt.text(x+1.0,y+0.5,loc) # print the label above and to the right
54          plt.plot([body.rob_x],[body.rob_y],"go")
55          plt.draw()
56
57      def plot_run(self):
58          """plots the history after the agent has finished.
59          This is typically only used if body.plotting==False
60          """
61          xs,ys = zip(*self.body.history)
62          plt.plot(xs,ys,"go")
63          wxs,wys = zip(*self.body.wall_history)
64          plt.plot(wxs,wys,"ro")
65          #plt.draw()
```

The following code plots the agent as it acts in the world:

_____ agentTop.py — (continued) _____
```
67  from agentEnv import Rob_body, Rob_env
68
69  env = Rob_env({((20,0),(30,20)), ((70,-5),(70,25))})
70  body = Rob_body(env)
71  middle = Rob_middle_layer(body)
72  top = Rob_top_layer(middle)
73
74  # try:
75  # pl=Plot_env(body,top)
76  # top.do({'visit':['o109','storage','o109','o103']})
77  # You can directly control the middle layer:
78  # middle.do({'go_to':(30,-10), 'timeout':200})
79  # Can you make it crash?
```

Exercise 2.1 The following code implements a robot trap. Write a controller that can escape the "trap" and get to the goal. See textbook for hints.

_____ agentTop.py — (continued) _____
```
81  # Robot Trap for which the current controller cannot escape:
82  trap_env = Rob_env({((10,-21),(10,0)), ((10,10),(10,31)), ((30,-10),(30,0)),
83                      ((30,10),(30,20)), ((50,-21),(50,31)), ((10,-21),(50,-21)),
84                      ((10,0),(30,0)), ((10,10),(30,10)), ((10,31),(50,31))})
85  trap_body = Rob_body(trap_env,init_pos=(-1,0,90))
86  trap_middle = Rob_middle_layer(trap_body)
87  trap_top = Rob_top_layer(trap_middle,locations={'goal':(71,0)})
88
89  # Robot trap exercise:
90  # pl=Plot_env(trap_body,trap_top)
91  # trap_top.do({'visit':['goal']})
```

Chapter 3

Searching for Solutions

3.1 Representing Search Problems

A search problem consists of:

- a start node

- a neighbors function that given a node, returns an enumeration of the arcs from the node

- a specification of a goal in terms of a Boolean function that takes a node and returns true if the node is a goal

- a (optional) heuristic function that, given a node, returns a non-negative real number. The heuristic function defaults to zero.

As far as the searcher is concerned a node can be anything. If multiple-path pruning is used, a node must be hashable. In the simple examples, it is a string, but in more complicated examples (in later chapters) it can be a tuple, a frozen set, or a Python object.

In the following code `raise NotImplementedError()` is a way to specify that this is an abstract method that needs to be overridden to define an actual search problem.

_____searchProblem.py — representations of search problems _____
```
11  class Search_problem(object):
12      """A search problem consists of:
13      * a start node
14      * a neighbors function that gives the neighbors of a node
15      * a specification of a goal
16      * a (optional) heuristic function.
```

```
17 |    The methods must be overridden to define a search problem."""
18 |
19 |    def start_node(self):
20 |        """returns start node"""
21 |        raise NotImplementedError("start_node") # abstract method
22 |
23 |    def is_goal(self,node):
24 |        """is True if node is a goal"""
25 |        raise NotImplementedError("is_goal") # abstract method
26 |
27 |    def neighbors(self,node):
28 |        """returns a list of the arcs for the neighbors of node"""
29 |        raise NotImplementedError("neighbors") # abstract method
30 |
31 |    def heuristic(self,n):
32 |        """Gives the heuristic value of node n.
33 |        Returns 0 if not overridden."""
34 |        return 0
```

The neighbors is a list of arcs. A (directed) arc consists of a *from_node* node and a *to_node* node. The arc is the pair ⟨*from_node*, *to_node*⟩, but can also contain a non-negative *cost* (which defaults to 1) and can be labeled with an *action*.

_____ searchProblem.py — (continued) _____

```
36 | class Arc(object):
37 |    """An arc has a from_node and a to_node node and a (non-negative) cost"""
38 |    def __init__(self, from_node, to_node, cost=1, action=None):
39 |        assert cost >= 0, ("Cost cannot be negative for"+
40 |                           str(from_node)+"->"+str(to_node)+", cost: "+str(cost))
41 |        self.from_node = from_node
42 |        self.to_node = to_node
43 |        self.action = action
44 |        self.cost=cost
45 |
46 |    def __repr__(self):
47 |        """string representation of an arc"""
48 |        if self.action:
49 |            return str(self.from_node)+" --"+str(self.action)+"--> "+str(self.to_node)
50 |        else:
51 |            return str(self.from_node)+" --> "+str(self.to_node)
```

3.1.1   Explicit Representation of Search Graph

The first representation of a search problem is from an explicit graph (as opposed to one that is generated as needed).

An **explicit graph** consists of

- a list or set of nodes

- a list or set of arcs

- a start node

- a list or set of goal nodes

- (optionally) a dictionary that maps a node to a heuristic value for that node

To define a search problem, we need to define the start node, the goal predicate, the neighbors function and the heuristic function.

```
_____searchProblem.py — (continued)_____
53  class Search_problem_from_explicit_graph(Search_problem):
54      """A search problem consists of:
55      * a list or set of nodes
56      * a list or set of arcs
57      * a start node
58      * a list or set of goal nodes
59      * a dictionary that maps each node into its heuristic value.
60      """
61
62      def __init__(self, nodes, arcs, start=None, goals=set(), hmap={}):
63          self.neighs = {}
64          self.nodes = nodes
65          for node in nodes:
66              self.neighs[node]=[]
67          self.arcs = arcs
68          for arc in arcs:
69              self.neighs[arc.from_node].append(arc)
70          self.start = start
71          self.goals = goals
72          self.hmap = hmap
73
74      def start_node(self):
75          """returns start node"""
76          return self.start
77
78      def is_goal(self,node):
79          """is True if node is a goal"""
80          return node in self.goals
81
82      def neighbors(self,node):
83          """returns the neighbors of node"""
84          return self.neighs[node]
85
86      def heuristic(self,node):
87          """Gives the heuristic value of node n.
88          Returns 0 if not overridden in the hmap."""
89          if node in self.hmap:
90              return self.hmap[node]
91          else:
92              return 0
```

```
93
94      def __repr__(self):
95          """returns a string representation of the search problem"""
96          res=""
97          for arc in self.arcs:
98              res += str(arc)+". "
99          return res
```

The following is used for the depth-first search implementation below.

_____searchProblem.py — (continued) _____

```
101     def neighbor_nodes(self,node):
102         """returns an iterator over the neighbors of node"""
103         return (path.to_node for path in self.neighs[node])
```

3.1.2   Paths

A searcher will return a path from the start node to a goal node. A Python list
is not a suitable representation for a path, as many search algorithms consider
multiple paths at once, and these paths should share initial parts of the path.
If we wanted to do this with Python lists, we would need to keep copying the
list, which can be expensive if the list is long. An alternative representation is
used here in terms of a recursive data structure that can share subparts.

A path is either:

- a node (representing a path of length 0) or

- a path, *initial* and an arc, where the *from_node* of the arc is the node at the
  end of *initial*.

These cases are distinguished in the following code by having *arc = None* if the
path has length 0, in which case *initial* is the node of the path.

_____searchProblem.py — (continued) _____

```
105 class Path(object):
106     """A path is either a node or a path followed by an arc"""
107
108     def __init__(self,initial,arc=None):
109         """initial is either a node (in which case arc is None) or
110         a path (in which case arc is an object of type Arc)"""
111         self.initial = initial
112         self.arc=arc
113         if arc is None:
114             self.cost=0
115         else:
116             self.cost = initial.cost+arc.cost
117
118     def end(self):
119         """returns the node at the end of the path"""
120         if self.arc is None:
```

```
121              return self.initial
122          else:
123              return self.arc.to_node
124
125      def nodes(self):
126          """enumerates the nodes for the path.
127          This starts at the end and enumerates nodes in the path backwards."""
128          current = self
129          while current.arc is not None:
130              yield current.arc.to_node
131              current = current.initial
132          yield current.initial
133
134      def initial_nodes(self):
135          """enumerates the nodes for the path before the end node.
136          This starts at the end and enumerates nodes in the path backwards."""
137          if self.arc is not None:
138              for nd in self.initial.nodes(): yield nd # could be "yield from"
139
140      def __repr__(self):
141          """returns a string representation of a path"""
142          if self.arc is None:
143              return str(self.initial)
144          elif self.arc.action:
145              return (str(self.initial)+"\n --"+str(self.arc.action)
146                      +"--> "+str(self.arc.to_node))
147          else:
148              return str(self.initial)+" --> "+str(self.arc.to_node)
```

3.1.3 Example Search Problems

The first search problem is one with 5 nodes where the least-cost path is one with many arcs. See Figure 3.1. Note that this example is used for the unit tests, so the test (in searchGeneric) will need to be changed if this is changed.

_____ searchProblem.py — (continued) _____

```
150  problem1 = Search_problem_from_explicit_graph(
151      {'a','b','c','d','g'},
152      [Arc('a','b',1), Arc('a','c',3), Arc('b','d',3), Arc('b','c',1),
153          Arc('c','d',1), Arc('c','g',3), Arc('d','g',1)],
154      start = 'a',
155      goals = {'g'})
```

The second search problem is one with 8 nodes where many paths do not lead to the goal. See Figure 3.2.

_____ searchProblem.py — (continued) _____

```
157  problem2 = Search_problem_from_explicit_graph(
158      {'a','b','c','d','e','g','h','j'},
159      [Arc('a','b',1), Arc('b','c',3), Arc('b','d',1), Arc('d','e',3),
```

Figure 3.1: problem1
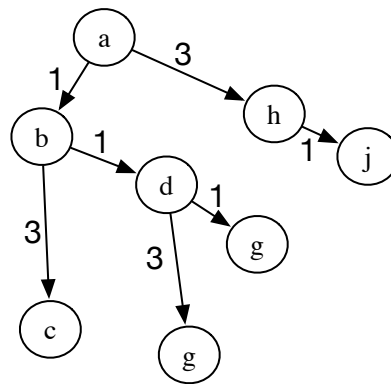
Figure 3.2: problem2

```
160        Arc('d','g',1), Arc('a','h',3), Arc('h','j',1)],
161     start = 'a',
162     goals = {'g'})
```

The third search problem is a disconnected graph (contains no arcs), where the start node is a goal node. This is a boundary case to make sure that weird cases work.

_____ searchProblem.py — (continued) _____

```
164  problem3 = Search_problem_from_explicit_graph(
165     {'a','b','c','d','e','g','h','j'},
166     [],
167     start = 'g',
168     goals = {'k','g'})
```

The acyclic_delivery_problem is the delivery problem described in Example 3.4 and shown in Figure 3.2 of the textbook.

_____ searchProblem.py — (continued) _____

```
170  acyclic_delivery_problem = Search_problem_from_explicit_graph(
171      {'mail','ts','o103','o109','o111','b1','b2','b3','b4','c1','c2','c3',
172       'o125','o123','o119','r123','storage'},
173      [Arc('ts','mail',6),
174         Arc('o103','ts',8),
175         Arc('o103','b3',4),
176         Arc('o103','o109',12),
177         Arc('o109','o119',16),
178         Arc('o109','o111',4),
179         Arc('b1','c2',3),
180         Arc('b1','b2',6),
181         Arc('b2','b4',3),
182         Arc('b3','b1',4),
183         Arc('b3','b4',7),
184         Arc('b4','o109',7),
185         Arc('c1','c3',8),
186         Arc('c2','c3',6),
187         Arc('c2','c1',4),
188         Arc('o123','o125',4),
189         Arc('o123','r123',4),
190         Arc('o119','o123',9),
191         Arc('o119','storage',7)],
192      start = 'o103',
193      goals = {'r123'},
194      hmap = {
195          'mail' : 26,
196          'ts' : 23,
197          'o103' : 21,
198          'o109' : 24,
199          'o111' : 27,
200          'o119' : 11,
201          'o123' : 4,
202          'o125' : 6,
203          'r123' : 0,
204          'b1' : 13,
205          'b2' : 15,
206          'b3' : 17,
207          'b4' : 18,
208          'c1' : 6,
209          'c2' : 10,
210          'c3' : 12,
211          'storage' : 12
212          }
213      )
```

The cyclic_delivery_problem is the delivery problem described in Example 3.8 and shown in Figure 3.6 of the textbook. This is the same as acyclic_delivery_problem, but almost every arc also has its inverse.

——————————————— searchProblem.py — (continued) ———————————————

```
215  cyclic_delivery_problem = Search_problem_from_explicit_graph(
```

```
216       {'mail','ts','o103','o109','o111','b1','b2','b3','b4','c1','c2','c3',
217        'o125','o123','o119','r123','storage'},
218        [ Arc('ts','mail',6), Arc('mail','ts',6),
219          Arc('o103','ts',8), Arc('ts','o103',8),
220          Arc('o103','b3',4),
221          Arc('o103','o109',12), Arc('o109','o103',12),
222          Arc('o109','o119',16), Arc('o119','o109',16),
223          Arc('o109','o111',4), Arc('o111','o109',4),
224          Arc('b1','c2',3),
225          Arc('b1','b2',6), Arc('b2','b1',6),
226          Arc('b2','b4',3), Arc('b4','b2',3),
227          Arc('b3','b1',4), Arc('b1','b3',4),
228          Arc('b3','b4',7), Arc('b4','b3',7),
229          Arc('b4','o109',7),
230          Arc('c1','c3',8), Arc('c3','c1',8),
231          Arc('c2','c3',6), Arc('c3','c2',6),
232          Arc('c2','c1',4), Arc('c1','c2',4),
233          Arc('o123','o125',4), Arc('o125','o123',4),
234          Arc('o123','r123',4), Arc('r123','o123',4),
235          Arc('o119','o123',9), Arc('o123','o119',9),
236          Arc('o119','storage',7), Arc('storage','o119',7)],
237     start = 'o103',
238     goals = {'r123'},
239     hmap = {
240         'mail' : 26,
241         'ts' : 23,
242         'o103' : 21,
243         'o109' : 24,
244         'o111' : 27,
245         'o119' : 11,
246         'o123' : 4,
247         'o125' : 6,
248         'r123' : 0,
249         'b1' : 13,
250         'b2' : 15,
251         'b3' : 17,
252         'b4' : 18,
253         'c1' : 6,
254         'c2' : 10,
255         'c3' : 12,
256         'storage' : 12
257         }
258     )
```

3.2 Generic Searcher and Variants

> To run the search demos, in folder "aipython", load "searchGeneric.py" , using e.g., ipython -i searchGeneric.py, and copy and paste the example queries at the bottom of that file. This requires Python 3.

3.2.1 Searcher

A *Searcher* for a problem can be asked repeatedly for the next path. To solve a problem, we can construct a *Searcher* object for the problem and then repeatedly ask for the next path using *search*. If there are no more paths, *None* is returned.

_____ searchGeneric.py — Generic Searcher, including depth-first and A* _____

```python
from utilities import Displayable, visualize

class Searcher(Displayable):
    """returns a searcher for a problem.
    Paths can be found by repeatedly calling search().
    This does depth-first search unless overridden
    """
    def __init__(self, problem):
        """creates a searcher from a problem
        """
        self.problem = problem
        self.initialize_frontier()
        self.num_expanded = 0
        self.add_to_frontier(Path(problem.start_node()))
        super().__init__()

    def initialize_frontier(self):
        self.frontier = []

    def empty_frontier(self):
        return self.frontier == []

    def add_to_frontier(self,path):
        self.frontier.append(path)

    @visualize
    def search(self):
        """returns (next) path from the problem's start node
        to a goal node.
        Returns None if no path exists.
        """
        while not self.empty_frontier():
            path = self.frontier.pop()
            self.display(2, "Expanding:",path,"(cost:",path.cost,")")
            self.num_expanded += 1
```

```
46          if self.problem.is_goal(path.end()): # solution found
47              self.display(1, self.num_expanded, "paths have been expanded and",
48                      len(self.frontier), "paths remain in the frontier")
49              self.solution = path # store the solution found
50              return path
51          else:
52              neighs = self.problem.neighbors(path.end())
53              self.display(3,"Neighbors are", neighs)
54              for arc in reversed(neighs):
55                  self.add_to_frontier(Path(path,arc))
56              self.display(3,"Frontier:",self.frontier)
57      self.display(1,"No (more) solutions. Total of",
58                  self.num_expanded,"paths expanded.")
```

Note that this reverses the neigbours so that it implements depth-first search in an intutive manner (expanding the first neighbor first). This might not be required for other methods.

Exercise 3.1 When it returns a path, the algorithm can be used to find another path by calling *search*() again. However, it does not find other paths that go through one goal node to another. Explain why, and change the code so that it can find such paths when *search*() is called again.

3.2.2   Frontier as a Priority Queue

In many of the search algorithms, such as A^* and other best-first searchers, the frontier is implemented as a priority queue. Here we use the Python's built-in priority queue implementations, heapq.

Following the lead of the Python documentation, http://docs.python.org/ 3.3/library/heapq.html, a frontier is a list of triples. The first element of each triple is the value to be minimized. The second element is a unique index which specifies the order when the first elements are the same, and the third element is the path that is on the queue. The use of the unique index ensures that the priority queue implementation does not compare paths; whether one path is less than another is not defined. It also lets us control what sort of search (e.g., depth-first or breadth-first) occurs when the value to be minimized does not give a unique next path.

The variable *frontier_index* is the total number of elements of the frontier that have been created. As well as being used as a unique index, it is useful for statistics, particularly in conjunction with the current size of the frontier.

_____ searchGeneric.py — (continued) _____

```
60  import heapq       # part of the Python standard library
61  from searchProblem import Path
62
63  class FrontierPQ(object):
64      """A frontier consists of a priority queue (heap), frontierpq, of
65          (value, index, path) triples, where
66      * value is the value we want to minimize (e.g., path cost + h).
```

```
67      * index is a unique index for each element
68      * path is the path on the queue
69      Note that the priority queue always returns the smallest element.
70      """
71
72      def __init__(self):
73          """constructs the frontier, initially an empty priority queue
74          """
75          self.frontier_index = 0 # the number of items ever added to the frontier
76          self.frontierpq = [] # the frontier priority queue
77
78      def empty(self):
79          """is True if the priority queue is empty"""
80          return self.frontierpq == []
81
82      def add(self, path, value):
83          """add a path to the priority queue
84          value is the value to be minimized"""
85          self.frontier_index += 1 # get a new unique index
86          heapq.heappush(self.frontierpq,(value, -self.frontier_index, path))
87
88      def pop(self):
89          """returns and removes the path of the frontier with minimum value.
90          """
91          (_,_,path) = heapq.heappop(self.frontierpq)
92          return path
```

The following methods are used for finding and printing information about
the frontier.

_____ searchGeneric.py — (continued) _____
```
94      def count(self,val):
95          """returns the number of elements of the frontier with value=val"""
96          return sum(1 for e in self.frontierpq if e[0]==val)
97
98      def __repr__(self):
99          """string representation of the frontier"""
100         return str([(n,c,str(p)) for (n,c,p) in self.frontierpq])
101
102     def __len__(self):
103         """length of the frontier"""
104         return len(self.frontierpq)
105
106     def __iter__(self):
107         """iterate through the paths in the frontier"""
108         for (_,_,path) in self.frontierpq:
109             yield path
```

3.2.3 A^* Search

For an A^* **Search** the frontier is implemented using the FrontierPQ class.

_____ searchGeneric.py — (continued) _____

```python
111  class AStarSearcher(Searcher):
112      """returns a searcher for a problem.
113      Paths can be found by repeatedly calling search().
114      """
115
116      def __init__(self, problem):
117          super().__init__(problem)
118
119      def initialize_frontier(self):
120          self.frontier = FrontierPQ()
121
122      def empty_frontier(self):
123          return self.frontier.empty()
124
125      def add_to_frontier(self,path):
126          """add path to the frontier with the appropriate cost"""
127          value = path.cost+self.problem.heuristic(path.end())
128          self.frontier.add(path, value)
```

Testing:

_____ searchGeneric.py — (continued) _____

```python
130  import searchProblem
131
132  def test(SearchClass):
133      print("Testing problem 1:")
134      schr1 = SearchClass(searchProblem.problem1)
135      path1 = schr1.search()
136      print("Path found:",path1)
137      assert list(path1.nodes()) == ['g','d','c','b','a'], "Shortest path not found in problem1"
138      print("Passed unit test")
139
140  if __name__ == "__main__":
141      #test(Searcher)
142      test(AStarSearcher)
143
144  # example queries:
145  # searcher1 = Searcher(searchProblem.acyclic_delivery_problem) # DFS
146  # searcher1.search() # find first path
147  # searcher1.search() # find next path
148  # searcher2 = AStarSearcher(searchProblem.acyclic_delivery_problem) # A*
149  # searcher2.search() # find first path
150  # searcher2.search() # find next path
151  # searcher3 = Searcher(searchProblem.cyclic_delivery_problem) # DFS
152  # searcher3.search() # find first path with DFS. What do you expect to happen?
153  # searcher4 = AStarSearcher(searchProblem.cyclic_delivery_problem) # A*
154  # searcher4.search() # find first path
```

Exercise 3.2 Change the code so that it implements (i) best-first search and (ii) lowest-cost-first search. For each of these methods compare it to A^* in terms of the number of paths expanded, and the path found.

Exercise 3.3 In the *add* method in *FrontierPQ* what does the "-" in front of *frontier_index* do? When there are multiple paths with the same f-value, which search method does this act like? What happens if the "-" is removed? When there are multiple paths with the same value, which search method does this act like? Does it work better with or without the "-"? What evidence did you base your conclusion on?

Exercise 3.4 The searcher acts like a Python iterator, in that it returns one value (here a path) and then returns other values (paths) on demand, but does not implement the iterator interface. Change the code so it implements the iterator interface. What does this enable us to do?

3.2.4 Multiple Path Pruning

> To run the multiple-path pruning demo, in folder "aipython", load "searchMPP.py" , using e.g., ipython -i searchMPP.py, and copy and paste the example queries at the bottom of that file.

The following implements A^* with multiple-path pruning. It overrides *search()* in *Searcher*.

---
searchMPP.py — Searcher with multiple-path pruning

```
11 from searchGeneric import AStarSearcher, visualize
12 from searchProblem import Path
13
14 class SearcherMPP(AStarSearcher):
15     """returns a searcher for a problem.
16     Paths can be found by repeatedly calling search().
17     """
18     def __init__(self, problem):
19         super().__init__(problem)
20         self.explored = set()
21
22     @visualize
23     def search(self):
24         """returns next path from an element of problem's start nodes
25         to a goal node.
26         Returns None if no path exists.
27         """
28         while not self.empty_frontier():
29             path = self.frontier.pop()
30             if path.end() not in self.explored:
31                 self.display(2, "Expanding:",path,"(cost:",path.cost,")")
32                 self.explored.add(path.end())
33                 self.num_expanded += 1
34                 if self.problem.is_goal(path.end()):
35                     self.display(1, self.num_expanded, "paths have been expanded and",
```

```
36                              len(self.frontier), "paths remain in the frontier")
37                          self.solution = path # store the solution found
38                          return path
39                      else:
40                          neighs = self.problem.neighbors(path.end())
41                          self.display(3,"Neighbors are", neighs)
42                          for arc in neighs:
43                              self.add_to_frontier(Path(path,arc))
44                          self.display(3,"Frontier:",self.frontier)
45              self.display(1,"No (more) solutions. Total of",
46                          self.num_expanded,"paths expanded.")
47
48  from searchGeneric import test
49  if __name__ == "__main__":
50      test(SearcherMPP)
51
52  import searchProblem
53  # searcherMPPcdp = SearcherMPP(searchProblem.cyclic_delivery_problem)
54  # print(searcherMPPcdp.search()) # find first path
```

Exercise 3.5 Implement a searcher that implements cycle pruning instead of multiple-path pruning. You need to decide whether to check for cycles when paths are added to the frontier or when they are removed. (Hint: either method can be implemented by only changing one or two lines in SearcherMPP.) Compare no pruning, multiple path pruning and cycle pruning for the cyclic delivery problem. Which works better in terms of number of paths expanded, computational time or space?

3.3   Branch-and-bound Search

> To run the demo, in folder "aipython", load "searchBranchAndBound.py", and copy and paste the example queries at the bottom of that file.

Depth-first search methods do not need an a priority queue, but can use a list as a stack. In this implementation of branch-and-bound search, we call *search* to find an optimal solution with cost less than bound. This uses depth-first search to find a path to a goal that extends *path* with cost less than the bound. Once a path to a goal has been found, that path is remembered as the *best_path*, the bound is reduced, and the search continues.

_____ searchBranchAndBound.py — Branch and Bound Search _____
```
11  from searchProblem import Path
12  from searchGeneric import Searcher
13  from utilities import Displayable, visualize
14
15  class DF_branch_and_bound(Searcher):
16      """returns a branch and bound searcher for a problem.
```

```
17          An optimal path with cost less than bound can be found by calling search()
18          """
19      def __init__(self, problem, bound=float("inf")):
20          """creates a searcher than can be used with search() to find an optimal path.
21          bound gives the initial bound. By default this is infinite - meaning there
22          is no initial pruning due to depth bound
23          """
24          super().__init__(problem)
25          self.best_path = None
26          self.bound = bound
27
28      @visualize
29      def search(self):
30          """returns an optimal solution to a problem with cost less than bound.
31          returns None if there is no solution with cost less than bound."""
32          self.frontier = [Path(self.problem.start_node())]
33          self.num_expanded = 0
34          while self.frontier:
35              path = self.frontier.pop()
36              if path.cost+self.problem.heuristic(path.end()) < self.bound:
37                  self.display(3,"Expanding:",path,"cost:",path.cost)
38                  self.num_expanded += 1
39                  if self.problem.is_goal(path.end()):
40                      self.best_path = path
41                      self.bound = path.cost
42                      self.display(2,"New best path:",path," cost:",path.cost)
43                  else:
44                      neighs = self.problem.neighbors(path.end())
45                      self.display(3,"Neighbors are", neighs)
46                      for arc in reversed(list(neighs)):
47                          self.add_to_frontier(Path(path, arc))
48          self.display(1,"Number of paths expanded:",self.num_expanded)
49          self.solution = self.best_path
50          return self.best_path
```

Note that this code used *reversed* in order to expand the neighbors of a node
in the left-to-right order one might expect. It does this because *pop()* removes
the rightmost element of the list. Note that reversed only works on lists and
tuples, but the neighbours can be generated.

Here is a unit test and some queries:

_____ searchBranchAndBound.py — (continued) _____

```
52  from searchGeneric import test
53  if __name__ == "__main__":
54      test(DF_branch_and_bound)
55
56  # Example queries:
57  import searchProblem
58  # searcherb1 = DF_branch_and_bound(searchProblem.acyclic_delivery_problem)
59  # print(searcherb1.search())     # find optimal path
60  # searcherb2 = DF_branch_and_bound(searchProblem.cyclic_delivery_problem, bound=100)
```

```
61  # print(searcherb2.search())      # find optimal path
```

Exercise 3.6 Implement a branch-and-bound search uses recursion. Hint: you don't need an explicit frontier, but can do a recursive call for the children.

Exercise 3.7 After the branch-and-bound search found a solution, Sam ran search again, and noticed a different count. Sam hypothesized that this count was related to the number of nodes that an $A*$ search would use (either expand or be added to the frontier). Or maybe, Sam thought, the count for a number of nodes when the bound is slightly above the optimal path case is related to how A^* would work. Is there relationship between these counts? Are there different things that it could count so they are related? Try to find the most specific statement that is true, and explain why it is true.

To test the hypothesis, Sam wrote the following code, but isn't sure it is helpful:

———————————searchTest.py — code that may be useful to compare A* and branch-and-bound ———————————

```
11  from searchGeneric import Searcher, AStarSearcher
12  from searchBranchAndBound import DF_branch_and_bound
13  from searchMPP import SearcherMPP
14
15  DF_branch_and_bound.max_display_level = 1
16  Searcher.max_display_level = 1
17
18  def run(problem,name):
19      print("\n\n*******",name)
20
21      print("\nA*:")
22      asearcher = AStarSearcher(problem)
23      print("Path found:",asearcher.search()," cost=",asearcher.solution.cost)
24      print("there are",asearcher.frontier.count(asearcher.solution.cost),
25            "elements remaining on the queue with f-value=",asearcher.solution.cost)
26
27      print("\nA* with MPP:"),
28      msearcher = SearcherMPP(problem)
29      print("Path found:",msearcher.search()," cost=",msearcher.solution.cost)
30      print("there are",msearcher.frontier.count(msearcher.solution.cost),
31            "elements remaining on the queue with f-value=",msearcher.solution.cost)
32
33      bound = asearcher.solution.cost+0.01
34      print("\nBranch and bound (with too-good initial bound of", bound,")")
35      tbb = DF_branch_and_bound(problem,bound) # cheating!!!!
36      print("Path found:",tbb.search()," cost=",tbb.solution.cost)
37      print("Rerunning B&B")
38      print("Path found:",tbb.search())
39
40      bbound = asearcher.solution.cost*2+10
41      print("\nBranch and bound (with not-very-good initial bound of", bbound, ")")
42      tbb2 = DF_branch_and_bound(problem,bbound) # cheating!!!!
43      print("Path found:",tbb2.search()," cost=",tbb2.solution.cost)
44      print("Rerunning B&B")
```

```
45      print("Path found:",tbb2.search())
46
47      print("\nDepth-first search: (Use ^C if it goes on forever)")
48      tsearcher = Searcher(problem)
49      print("Path found:",tsearcher.search()," cost=",tsearcher.solution.cost)
50
51
52
53  import searchProblem
54  from searchTest import run
55  if __name__ == "__main__":
56      run(searchProblem.problem1,"Problem 1")
57  #   run(searchProblem.acyclic_delivery_problem,"Acyclic Delivery")
58  #   run(searchProblem.cyclic_delivery_problem,"Cyclic Delivery")
59  # also test some graphs with cycles, and some with multiple least-cost paths
```

Chapter 4

Reasoning with Constraints

4.1 Constraint Satisfaction Problems

4.1.1 Constraints

A **variable** is a string or any value that is printable and can be the key of a Python dictionary.

A **constraint** consists of a tuple (or list) of variables and a condition.

- The tuple (or list) of variables is called the **scope**.

- The **condition** is a Boolean function that takes the same number of arguments as there are variables in the scope. The condition must have a `__name__` property that gives a printable name of the function; built-in functions and functions that are defined using *def* have such a property; for other functions you may need to define this property.

_____cspProblem.py — Representations of a Constraint Satisfaction Problem _____
```python
11  from utilities import Displayable, dict_union
12
13  class Constraint(object):
14      """A Constraint consists of
15      * scope: a tuple of variables
16      * condition: a function that can applied to a tuple of values
17      for the variables
18      """
19      def __init__(self, scope, condition):
20          self.scope = scope
21          self.condition = condition
22
23      def __repr__(self):
```

```
24            return self.condition.__name__ + str(self.scope)
```

An **assignment** is a *variable:value* dictionary.

If *con* is a constraint, *con.holds*(*assignment*) returns True or False depending on whether the condition is true or false for that assignment. The assignment *assignment* must assigns a value to every variable in the scope of the constraint *con* (and could also assign values other variables); *con.holds* gives an error if not all variables in the scope of *con* are assigned in the assignment. It ignores variables in *assignment* that are not in the scope of the constraint.

In Python, the ∗ notation is used for unpacking a tuple. For example, $F(*(1,2,3))$ is the same as $F(1,2,3)$. So if t has value $(1,2,3)$, then $F(*t)$ is the same as $F(1,2,3)$.

─────────────────────── _cspProblem.py — (continued)_ ───────────────────────

```
26      def holds(self,assignment):
27          """returns the value of Constraint con evaluated in assignment.
28
29          precondition: all variables are assigned in assignment
30          """
31          return self.condition(*tuple(assignment[v] for v in self.scope))
```

4.1.2 CSPs

A constraint satisfaction problem (CSP) requires:

- *domains*: a dictionary that maps variables to the set of possible values. Thus *domains*[*var*] is the domain of variable *var*.

- *constaraints*: a set or list of constraints.

Other properties are inferred from these:

- *variables* is the set of variables. The variables can be enumerated by using "for var in domains" because iterating over a dictionary gives the keys, which in this case are the variables.

- *var_to_const* is a mapping from variables to set of constraints, such that *var_to_const*[*var*] is the set of constraints with *var* in the scope.

─────────────────────── _cspProblem.py — (continued)_ ───────────────────────

```
33  class CSP(Displayable):
34      """A CSP consists of
35      * domains, a dictionary that maps each variable to its domain
36      * constraints, a list of constraints
37      * variables, a set of variables
38      * var_to_const, a variable to set of constraints dictionary
39      """
40      def __init__(self,domains,constraints):
```

```
41        """domains is a variable:domain dictionary
42        constraints is a list of constriants
43        """
44        self.variables = set(domains)
45        self.domains = domains
46        self.constraints = constraints
47        self.var_to_const = {var:set() for var in self.variables}
48        for con in constraints:
49            for var in con.scope:
50                self.var_to_const[var].add(con)
51
52    def __str__(self):
53        """string representation of CSP"""
54        return str(self.domains)
55
56    def __repr__(self):
57        """more detailed string representation of CSP"""
58        return "CSP("+str(self.domains)+", "+str([str(c) for c in self.constraints])+")"
```

csp.consistent(*assignment*) returns true if the assignment is consistent with each of the constraints in *csp* (i.e., all of the constraints that can be evaluated evaluate to true). Note that this is a local consistency with each constraint; it does *not* imply the CSP is consistent or has a solution.

─────────────── cspProblem.py — (continued) ───────────────

```
60    def consistent(self,assignment):
61        """assignment is a variable:value dictionary
62        returns True if all of the constraints that can be evaluated
63                evaluate to True given assignment.
64        """
65        return all(con.holds(assignment)
66                    for con in self.constraints
67                    if all(v in assignment for v in con.scope))
```

4.1.3 Examples

In the following code *ne_*, when given a number, returns a function that is true when its argument is not that number. For example, if $f = ne_(3)$, then $f(2)$ is True and $f(3)$ is False. That is, $ne_(x)(y)$ is true when $x \neq y$. Allowing a function of multiple arguments to use its arguments one at a time is called **currying**, after the logician Haskell Curry. The use of a condition in constraints requires that the function with a single argument has a name.

─────────────── cspExamples.py — Example CSPs ───────────────

```
11  from cspProblem import CSP, Constraint
12  from operator import lt,ne,eq,gt
13
14  def ne_(val):
15      """not equal value"""
```

```
16        # nev = lambda x: x != val # alternative definition
17        # nev = partial(neq,val)  # another alternative definition
18        def nev(x):
19            return val != x
20        nev.__name__ = str(val)+"!="   # name of the function
21        return nev
```

Similarly $is\_(x)(y)$ is true when $x = y$.

cspExamples.py — (continued)

```
23  def is_(val):
24      """is a value"""
25      # isv = lambda x: x == val # alternative definition
26      # isv = partial(eq,val)   # another alternative definition
27      def isv(x):
28          return val == x
29      isv.__name__ = str(val)+"=="
30      return isv
```

The CSP, *csp*0 has variables *X*, *Y* and *Z*, each with domain $\{1, 2, 3\}$. The constraints are $X < Y$ and $Y < Z$.

cspExamples.py — (continued)

```
32  csp0 = CSP({'X':{1,2,3},'Y':{1,2,3}, 'Z':{1,2,3}},
33          [ Constraint(('X','Y'),lt),
34            Constraint(('Y','Z'),lt)])
```

The CSP, *csp*1 has variables *A*, *B* and *C*, each with domain $\{1, 2, 3, 4\}$. The constraints are $A < B$, $B \neq 2$ and $B < C$. This is slightly more interesting than *csp*0 as it has more solutions. This example is used in the unit tests, and so if it is changed, the unit tests need to be changed.

cspExamples.py — (continued)

```
36  C0 = Constraint(('A','B'),lt)
37  C1 = Constraint(('B',),ne_(2))
38  C2 = Constraint(('B','C'),lt)
39  csp1 = CSP({'A':{1,2,3,4},'B':{1,2,3,4}, 'C':{1,2,3,4}},
40          [C0, C1, C2])
```

The next CSP, *csp*2 is Example 4.9 of the textbook; the domain consistent network is shown in Figure 4.1.

cspExamples.py — (continued)

```
42  csp2 = CSP({'A':{1,2,3,4},'B':{1,2,3,4}, 'C':{1,2,3,4},
43          'D':{1,2,3,4}, 'E':{1,2,3,4}},
44          [ Constraint(('B',),ne_(3)),
45            Constraint(('C',),ne_(2)),
46            Constraint(('A','B'),ne),
47            Constraint(('B','C'),ne),
48            Constraint(('C','D'),lt),
49            Constraint(('A','D'),eq),
50            Constraint(('A','E'),gt),
```

Figure 4.1: Domain-consistent constraint network (*csp*2).

```
51              Constraint(('B','E'),gt),
52              Constraint(('C','E'),gt),
53              Constraint(('D','E'),gt),
54              Constraint(('B','D'),ne)])
```

The following example is another scheduling problem (but with multiple answers). This is the same a scheduling 2 in the original AIspace.org consistency app.

_____cspExamples.py — (continued) _____

```
56  csp3 = CSP({'A':{1,2,3,4},'B':{1,2,3,4}, 'C':{1,2,3,4},
57             'D':{1,2,3,4}, 'E':{1,2,3,4}},
58          [Constraint(('A','B'), ne),
59           Constraint(('A','D'), lt),
60           Constraint(('A','E'), lambda a,e: (a-e)%2 == 1), # A-E is odd
61           Constraint(('B','E'), lt),
62           Constraint(('D','C'), lt),
63           Constraint(('C','E'), ne),
64           Constraint(('D','E'), ne)])
```

The following example is another abstract scheduling problem. What are the solutions?

_____cspExamples.py — (continued) _____

Figure 4.2: A crossword puzzle to be solved

Words:
ant, big, bus, car, has, book, buys, hold, lane, year, ginger, search, symbol, syntax.

```
66  def adjacent(x,y):
67      """True when x and y are adjacent numbers"""
68      return abs(x-y) == 1
69
70  csp4 = CSP({'A':{1,2,3,4,5},'B':{1,2,3,4,5}, 'C':{1,2,3,4,5},
71              'D':{1,2,3,4,5}, 'E':{1,2,3,4,5}},
72           [Constraint(('A','B'),adjacent),
73            Constraint(('B','C'),adjacent),
74            Constraint(('C','D'),adjacent),
75            Constraint(('D','E'),adjacent),
76            Constraint(('A','C'),ne),
77            Constraint(('B','D'),ne),
78            Constraint(('C','E'),ne)])
```

The following examples represent the crossword shown in Figure 4.2.

```
                        ___cspExamples.py — (continued)___
80  def meet_at(p1,p2):
81      """returns a function that is true when the words meet at the postions p1, p2
82      """
83      def meets(w1,w2):
84          return w1[p1] == w2[p2]
85      meets.__name__ = "meet_at("+str(p1)+','+str(p2)+')'
86      return meets
87
88  crossword1 = CSP({'one_across':{'ant', 'big', 'bus', 'car', 'has'},
89                    'one_down':{'book', 'buys', 'hold', 'lane', 'year'},
90                    'two_down':{'ginger', 'search', 'symbol', 'syntax'},
91                    'three_across':{'book', 'buys', 'hold', 'land', 'year'},
92                    'four_across':{'ant', 'big', 'bus', 'car', 'has'}},
93                   [Constraint(('one_across','one_down'),meet_at(0,0)),
94                    Constraint(('one_across','two_down'),meet_at(2,0)),
95                    Constraint(('three_across','two_down'),meet_at(2,2)),
```

```
96                      Constraint(('three_across','one_down'),meet_at(0,2)),
97                      Constraint(('four_across','two_down'),meet_at(0,4))])
```

In an alternative representation of a crossword (the "dual" representation), the variables represent letters, and the constraints are that adjacent sequences of letters form words.

——————————————————cspExamples.py — (continued) ——————————————————

```
99   words = {'ant', 'big', 'bus', 'car', 'has','book', 'buys', 'hold',
100          'lane', 'year', 'ginger', 'search', 'symbol', 'syntax'}
101
102  def is_word(*letters, words=words):
103      """is true if the letters concatenated form a word in words"""
104      return "".join(letters) in words
105
106  letters = ["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l",
107    "m", "n", "o", "p", "q", "r", "s", "t", "u", "v", "w", "x", "y",
108    "z"]
109
110  crossword1d = CSP({'p00':letters, 'p10':letters, 'p20':letters, # first row
111                     'p01':letters, 'p21':letters, # second row
112                     'p02':letters, 'p12':letters, 'p22':letters, 'p32':letters, # third row
113                     'p03':letters, 'p23':letters, #fourth row
114                     'p24':letters, 'p34':letters, 'p44':letters, # fifth row
115                     'p25':letters # sixth row
116                     },
117                    [Constraint(('p00', 'p10', 'p20'), is_word), #1-across
118                     Constraint(('p00', 'p01', 'p02', 'p03'), is_word), # 1-down
119                     Constraint(('p02', 'p12', 'p22', 'p32'), is_word), # 3-across
120                     Constraint(('p20', 'p21', 'p22', 'p23', 'p24', 'p25'), is_word), # 2-down
121                     Constraint(('p24', 'p34', 'p44'), is_word) # 4-across
122                     ])
```

Unit tests

The following defines a unit test for solvers on example csp1.

——————————————————cspExamples.py — (continued) ——————————————————

```
124  def test(CSP_solver, csp=csp1,
125             solutions=[{'A': 1, 'B': 3, 'C': 4}, {'A': 2, 'B': 3, 'C': 4}]):
126      """CSP_solver is a solver that finds a solution to a CSP.
127      CSP_solver takes a csp and returns a solution.
128      csp has to be a CSP, where solutions is the list of all solutions.
129      This tests whether the solution returned by CSP_solver is a solution.
130      """
131      print("Testing csp with",CSP_solver.__doc__)
132      sol0 = CSP_solver(csp)
133      print("Solution found:",sol0)
134      assert sol0 in solutions, "Solution not found for "+str(csp)
135      print("Passed unit test")
```

Exercise 4.1 Modify *test* so that instead of taking in a list of solutions, it checks whether the returned solution actually is a solution.

Exercise 4.2 Propose a test that is appropriate for CSPs with no solutions. Assume that the test designer knows there are no solutions. Consider what a CSP solver should return if there are no solutions to the CSP.

4.2 Solving a CSP using Search

> To run the demo, in folder "aipython", load "cspSearch.py", and copy and paste the example queries at the bottom of that file.

The first solver searches through the space of partial assignments. This takes in a CSP problem and an optional variable ordering, which is a list of the variables in the CSP. It then constructs a search space that can be solved using the search methods of the previous chapter. In this search space:

- A node is a *variable* : *value* dictionary.

- An arc corresponds to an assignment of a value to the next variable. This assumes a static ordering; the next variable chosen to split does not depend on the context. If no variable ordering is given, this makes no attempt to choose a good ordering.

_____ cspSearch.py — Representations of a Search Problem from a CSP. _____

```
11  from cspProblem import CSP, Constraint
12  from searchProblem import Arc, Search_problem
13  from utilities import dict_union
14
15  class Search_from_CSP(Search_problem):
16      """A search problem directly from the CSP.
17
18      A node is a variable:value dictionary"""
19      def __init__(self, csp, variable_order=None):
20          self.csp=csp
21          if variable_order:
22              assert set(variable_order) == set(csp.variables)
23              assert len(variable_order) == len(csp.variables)
24              self.variables = variable_order
25          else:
26              self.variables = list(csp.variables)
27
28      def is_goal(self, node):
29          return len(node)==len(self.csp.variables)
30
31      def start_node(self):
32          return {}
```

The *neighbors(node)* method uses the fact that the length of the node, which is the number of variables already assigned, is the index of the next variable to split on.

_____cspSearch.py — (continued)_____

```
34    def neighbors(self, node):
35        """returns a list of the neighboring nodes of node.
36        """
37        var = self.variables[len(node)] # the next variable
38        res = []
39        for val in self.csp.domains[var]:
40            new_env = dict_union(node,{var:val}) #dictionary union
41            if self.csp.consistent(new_env):
42                res.append(Arc(node,new_env))
43        return res
```

_____cspSearch.py — (continued)_____

```
45   from cspExamples import csp1,csp2,test, crossword1, crossword1d
46   from searchGeneric import Searcher
47
48   def dfs_solver(csp):
49       """depth-first search solver"""
50       path = Searcher(Search_from_CSP(csp)).search()
51       if path is not None:
52           return path.end()
53       else:
54           return None
55
56   if __name__ == "__main__":
57       test(dfs_solver)
58
59   ## Test Solving CSPs with Search:
60   searcher1 = Searcher(Search_from_CSP(csp1))
61   #print(searcher1.search()) # get next solution
62   searcher2 = Searcher(Search_from_CSP(csp2))
63   #print(searcher2.search()) # get next solution
64   searcher3 = Searcher(Search_from_CSP(crossword1))
65   #print(searcher3.search()) # get next solution
66   searcher4 = Searcher(Search_from_CSP(crossword1d))
67   #print(searcher4.search()) # get next solution (warning: slow)
```

Exercise 4.3 What would happen if we constructed the new assignment by assigning *node*[*var*] = *val* (with side effects) instead of using dictionary union? Give an example of where this could give a wrong answer. How could the algorithm be changed to work with side effects? (Hint: think about what information needs to be in a node).

Exercise 4.4 Change neighbors so that it returns an iterator of values rather than a list. (Hint: use *yield*.)

4.3 Consistency Algorithms

> To run the demo, in folder "aipython", load "cspConsistency.py", and
> copy and paste the commented-out example queries at the bottom of
> that file.

A *Con_solver* is used to simplify a CSP using arc consistency.

_____cspConsistency.py — Arc Consistency and Domain splitting for solving a CSP _____
```python
11  from utilities import Displayable
12
13  class Con_solver(Displayable):
14      def __init__(self, csp, **kwargs):
15          """a CSP solver that uses arc consistency
16          * csp is the CSP to be solved
17          * kwargs is the keyword arguments for Displayable superclass
18          """
19          super().__init__(**kwargs) # Or Displayable.__init__(self,**kwargs)
20          self.csp = csp
```

The following implementation of arc consistency maintains the set *to_do* of
(variable, constraint) pairs that are to be checked. It takes in a domain dic-
tionary and returns a new domain dictionary. It needs to be careful to avoid
side effects (by copying the *domains* dictionary and the *to_do* set).

_____cspConsistency.py — (continued) _____
```python
22      def make_arc_consistent(self, orig_domains=None, to_do=None):
23          """Makes this CSP arc-consistent using generalized arc consistency
24          orig_domains is the original domains
25          to_do is a set of (variable,constraint) pairs
26          returns the reduced domains (an arc-consistent variable:domain dictionary)
27          """
28          if orig_domains is None:
29              orig_domains = self.csp.domains
30          if to_do is None:
31              to_do = {(var, const) for const in self.csp.constraints
32                      for var in const.scope}
33          else:
34              to_do = to_do.copy() # use a copy of to_do
35          domains = orig_domains.copy()
36          self.display(2,"Performing AC with domains", domains)
37          while to_do:
38              var, const = self.select_arc(to_do)
39              self.display(3, "Processing arc (", var, ",", const, ")")
40              other_vars = [ov for ov in const.scope if ov != var]
41              if len(other_vars)==0:
42                  new_domain = {val for val in domains[var]
43                              if const.holds({var:val})}
44              elif len(other_vars)==1:
45                  other = other_vars[0]
```

```
46          new_domain = {val for val in domains[var]
47                        if any(const.holds({var: val,other:other_val})
48                            for other_val in domains[other])}
49       else: # general case
50          new_domain = {val for val in domains[var]
51                        if self.any_holds(domains, const, {var: val}, other_vars)}
52       if new_domain != domains[var]:
53          self.display(4, "Arc: (", var, ",", const, ") is inconsistent")
54          self.display(3, "Domain pruned", "dom(", var, ") =", new_domain,
55                        " due to ", const)
56          domains[var] = new_domain
57          add_to_do = self.new_to_do(var, const) - to_do
58          to_do |= add_to_do    # set union
59          self.display(3, " adding", add_to_do if add_to_do else "nothing", "to to_do.")
60       self.display(4, "Arc: (", var, ",", const, ") now consistent")
61    self.display(2, "AC done. Reduced domains", domains)
62    return domains
63
64  def new_to_do(self, var, const):
65      """returns new elements to be added to to_do after assigning
66      variable var in constraint const.
67      """
68      return {(nvar, nconst) for nconst in self.csp.var_to_const[var]
69              if nconst != const
70              for nvar in nconst.scope
71              if nvar != var}
```

The following selects an arc. Any element of *to_do* can be selected. The selected element needs to be removed from *to_do*. The default implementation just selects which ever element *pop* method for sets returns. A user interface could allow the user to select an arc. Alternatively a more sophisticated selection could be employed (or just a stack or a queue).

_____cspConsistency.py — (continued) _____

```
73  def select_arc(self, to_do):
74      """Selects the arc to be taken from to_do .
75      * to_do is a set of arcs, where an arc is a (variable,constraint) pair
76      the element selected must be removed from to_do.
77      """
78      return to_do.pop()
```

The function *any_holds* is useful to go beyond unary and binary constraints. It allows us to use constraints involving an arbitrary number of variables. (Note that it also works for unary and binary constraints; the cases where *len(other_vars)* is 0 or 1 are not actually required, but are there for efficiency and because they are easier to understand.) *any_holds* is a recursive function that tries to finds an assignment of values to the other variables (*other_vars*) that satisfies constraint *const* given the assignment in *env*. The integer variable *ind* specifies which index to *other_vars* needs to be checked next. As soon as one assignment returns *True*, the algorithm returns *True*. Note that it has side effects with respect to

env; it changes the values of the variables in *other_vars*. It should only be called when the side effects have no ill effects.

```
                          cspConsistency.py — (continued)
80      def any_holds(self, domains, const, env, other_vars, ind=0):
81          """returns True if Constraint const holds for an assignment
82          that extends env with the variables in other_vars[ind:]
83          env is a dictionary
84          Warning: this has side effects and changes the elements of env
85          """
86          if ind == len(other_vars):
87              return const.holds(env)
88          else:
89              var = other_vars[ind]
90              for val in domains[var]:
91                  # env = dict_union(env,{var:val}) # no side effects!
92                  env[var] = val
93                  if self.any_holds(domains, const, env, other_vars, ind + 1):
94                      return True
95              return False
```

4.3.1   Direct Implementation of Domain Splitting

The following is a direct implementation of domain splitting with arc consistency that uses recursion. It finds one solution if one exists or returns False if there are no solutions.

```
                          cspConsistency.py — (continued)
97      def solve_one(self, domains=None, to_do=None):
98          """return a solution to the current CSP or False if there are no solutions
99          to_do is the list of arcs to check
100         """
101         if domains is None:
102             domains = self.csp.domains
103         new_domains = self.make_arc_consistent(domains, to_do)
104         if any(len(new_domains[var]) == 0 for var in domains):
105             return False
106         elif all(len(new_domains[var]) == 1 for var in domains):
107             self.display(2, "solution:", {var: select(
108                 new_domains[var]) for var in new_domains})
109             return {var: select(new_domains[var]) for var in domains}
110         else:
111             var = select(x for x in self.csp.variables if len(new_domains[x]) > 1)
112             if var:
113                 dom1, dom2 = partition_domain(new_domains[var])
114                 self.display(3, "...splitting", var, "into", dom1, "and", dom2)
115                 new_doms1 = copy_with_assign(new_domains, var, dom1)
116                 new_doms2 = copy_with_assign(new_domains, var, dom2)
117                 to_do = self.new_to_do(var, None)
```

```
118              return self.solve_one(new_doms1, to_do) or self.solve_one(new_doms2, to_do)
119
120  def partition_domain(dom):
121      """partitions domain dom into two.
122      """
123      split = len(dom) // 2
124      dom1 = set(list(dom)[:split])
125      dom2 = dom - dom1
126      return dom1, dom2
```

The domains are implemented as a dictionary that maps each variables to its domain. Assigning a value in Python has side effects which we want to avoid. *copy_with_assign* takes a copy of the domains dictionary, perhaps allowing for a new domain for a variable. It creates a copy of the CSP with an (optional) assignment of a new domain to a variable. Only the domains are copied.

---
cspConsistency.py — (continued)

```
128  def copy_with_assign(domains, var=None, new_domain={True, False}):
129      """create a copy of the domains with an assignment var=new_domain
130      if var==None then it is just a copy.
131      """
132      newdoms = domains.copy()
133      if var is not None:
134          newdoms[var] = new_domain
135      return newdoms
```

---
cspConsistency.py — (continued)

```
137  def select(iterable):
138      """select an element of iterable. Returns None if there is no such element.
139
140      This implementation just picks the first element.
141      For many of the uses, which element is selected does not affect correctness,
142      but may affect efficiency.
143      """
144      for e in iterable:
145          return e # returns first element found
```

Exercise 4.5 Implement of *solve_all* that is like *solve_one* but returns the set of all solutions.

Exercise 4.6 Implement *solve_enum* that enumerates the solutions. It should use Python's *yield* (and perhaps *yield from*).

Unit test:

---
cspConsistency.py — (continued)

```
147  from cspExamples import test
148  def ac_solver(csp):
149      "arc consistency (solve_one)"
150      return Con_solver(csp).solve_one()
```

```
151  if __name__ == "__main__":
152      test(ac_solver)
```

4.3.2 Domain Splitting as an interface to graph searching

An alternative implementation is to implement domain splitting in terms of
the search abstraction of Chapter 3.

A node is domains dictionary.

_____cspConsistency.py — (continued) _____

```
154  from searchProblem import Arc, Search_problem
155
156  class Search_with_AC_from_CSP(Search_problem,Displayable):
157      """A search problem with arc consistency and domain splitting
158
159      A node is a CSP """
160      def __init__(self, csp):
161          self.cons = Con_solver(csp) #copy of the CSP
162          self.domains = self.cons.make_arc_consistent()
163
164      def is_goal(self, node):
165          """node is a goal if all domains have 1 element"""
166          return all(len(node[var])==1 for var in node)
167
168      def start_node(self):
169          return self.domains
170
171      def neighbors(self,node):
172          """returns the neighboring nodes of node.
173          """
174          neighs = []
175          var = select(x for x in node if len(node[x])>1)
176          if var:
177              dom1, dom2 = partition_domain(node[var])
178              self.display(2,"Splitting", var, "into", dom1, "and", dom2)
179              to_do = self.cons.new_to_do(var,None)
180              for dom in [dom1,dom2]:
181                  newdoms = copy_with_assign(node,var,dom)
182                  cons_doms = self.cons.make_arc_consistent(newdoms,to_do)
183                  if all(len(cons_doms[v])>0 for v in cons_doms):
184                      # all domains are non-empty
185                      neighs.append(Arc(node,cons_doms))
186                  else:
187                      self.display(2,"...",var,"in",dom,"has no solution")
188          return neighs
```

Exercise 4.7 When splitting a domain, this code splits the domain into half, ap-
proximately in half (without any effort to make a sensible choice). Does it work
better to split one element from a domain?

Unit test:

```
_____cspConsistency.py — (continued) _____
190  from cspExamples import test
191  from searchGeneric import Searcher
192
193  def ac_search_solver(csp):
194      """arc consistency (search interface)"""
195      sol = Searcher(Search_with_AC_from_CSP(csp)).search()
196      if sol:
197          return {v:select(d) for (v,d) in sol.end().items()}
198
199  if __name__ == "__main__":
200      test(ac_search_solver)
```

Testing:

```
_____cspConsistency.py — (continued) _____
202  from cspExamples import csp1, csp2, crossword1, crossword1d
203
204  ## Test Solving CSPs with Arc consistency and domain splitting:
205  #Con_solver.max_display_level = 4 # display details of AC (0 turns off)
206  #Con_solver(csp1).solve_one()
207  #searcher1d = Searcher(Search_with_AC_from_CSP(csp1))
208  #print(searcher1d.search())
209  #Searcher.max_display_level = 2 # display search trace (0 turns off)
210  #searcher2c = Searcher(Search_with_AC_from_CSP(csp2))
211  #print(searcher2c.search())
212  #searcher3c = Searcher(Search_with_AC_from_CSP(crossword1))
213  #print(searcher3c.search())
214  #searcher5c = Searcher(Search_with_AC_from_CSP(crossword1d))
215  #print(searcher5c.search())
```

4.4 Solving CSPs using Stochastic Local Search

> To run the demo, in folder "aipython", load "cspSLS.py", and copy
> and paste the commented-out example queries at the bottom of that
> file. This assumes Python 3. Some of the queries require matplotlib.

This implements both the two-stage choice, the any-conflict algorithm and a random choice of variable (and a probabilistic mix of the three).

Given a CSP, the stochastic local searcher (*SLSearcher*) creates the data structures:

- *variables_to_select* is the set of all of the variables with domain-size greater than one. For a variable not in this set, we cannot pick another value from that variable.

- *var_to_constraints* maps from a variable into the set of constraints it is involved in. Note that the inverse mapping from constraints into variables is part of the definition of a constraint.

---

cspSLS.py — Stochastic Local Search for Solving CSPs

```
11  from cspProblem import CSP, Constraint
12  from searchProblem import Arc, Search_problem
13  from utilities import Displayable
14  import random
15  import heapq
16
17  class SLSearcher(Displayable):
18      """A search problem directly from the CSP..
19
20      A node is a variable:value dictionary"""
21      def __init__(self, csp):
22          self.csp = csp
23          self.variables_to_select = {var for var in self.csp.variables
24                                      if len(self.csp.domains[var]) > 1}
25          # Create assignment and conflicts set
26          self.current_assignment = None # this will trigger a random restart
27          self.number_of_steps = 1 #number of steps after the initialization
```

restart creates a new total assignment, and constructs the set of conflicts (the constraints that are false in this assignment).

---

cspSLS.py — (continued)

```
29      def restart(self):
30          """creates a new total assignment and the conflict set
31          """
32          self.current_assignment = {var:random_sample(dom) for
33                                     (var,dom) in self.csp.domains.items()}
34          self.display(2,"Initial assignment",self.current_assignment)
35          self.conflicts = set()
36          for con in self.csp.constraints:
37              if not con.holds(self.current_assignment):
38                  self.conflicts.add(con)
39          self.display(2,"Number of conflicts",len(self.conflicts))
40          self.variable_pq = None
```

The *search* method is the top-level searching algorithm. It can either be used to start the search or to continue searching. If there is no current assignment, it must create one. Note that, when counting steps, a restart is counted as one step.

This method selects one of two implementations. The argument *pob_best* is the probability of selecting a best variable (one involving the most conflicts). When the value of *prob_best* is positive, the algorithm needs to maintain a priority queue of variables and the number of conflicts (using *search_with_var_pq*). If the probability of selecting a best variable is zero, it does not need to maintain this priority queue (as implemented in *search_with_any_conflict*).

The argument *prob_anycon* is the probability that the any-conflict strategy is used (which selects a variable at random that is in a conflict), assuming that it is not picking a best variable. Note that for the probability parameters, any value less that zero acts like probability zero and any value greater than 1 acts like probability 1. This means that when *prob_anycon* = 1.0, a best variable is chosen with probability *prob_best*, otherwise a variable in any conflict is chosen. A variable is chosen at random with probability 1 − *prob_anycon* − *prob_best* as long as that is positive.

This returns the number of steps needed to find a solution, or *None* if no solution is found. If there is a solution, it is in *self.current_assignment*.

```
_____ cspSLS.py — (continued) _____

42      def search(self,max_steps, prob_best=1.0, prob_anycon=1.0):
43          """
44          returns the number of steps or None if these is no solution.
45          If there is a solution, it can be found in self.current_assignment
46
47          max_steps is the maximum number of steps it will try before giving up
48          prob_best is the probability that a best varaible (one in most conflict) is selected
49          prob_anycon is the probability that a variabe in any conflict is selected
50          (otherwise a variable is chosen at random)
51          """
52          if self.current_assignment is None:
53              self.restart()
54              self.number_of_steps += 1
55              if not self.conflicts:
56                  return self.number_of_steps
57          if prob_best > 0: # we need to maintain a variable priority queue
58              return self.search_with_var_pq(max_steps, prob_best, prob_anycon)
59          else:
60              return self.search_with_any_conflict(max_steps, prob_anycon)
```

Exercise 4.8 This does an initial random assignment but does not do any random restarts. Implement a searcher that takes in the maximum number of walk steps (corresponding to existing *max_steps*) and the maximum number of restarts, and returns the total number of steps for the first solution found. (As in *search*, the solution found can be extracted from the variable *self.current_assignment*).

4.4.1 Any-conflict

If the probability of picking a best variable is zero, the implementation need to keeps track of which variables are in conflicts.

```
_____ cspSLS.py — (continued) _____

62      def search_with_any_conflict(self, max_steps, prob_anycon=1.0):
63          """Searches with the any_conflict heuristic.
64          This relies on just maintaining the set of conflicts;
65          it does not maintain a priority queue
66          """
```

```
67          self.variable_pq = None # we are not maintaining the priority queue.
68                             # This ensures it is regenerated if needed.
69      for i in range(max_steps):
70          self.number_of_steps +=1
71          if random.random() < prob_anycon:
72              con = random_sample(self.conflicts) # pick random conflict
73              var = random_sample(con.scope) # pick variable in conflict
74          else:
75              var = random_sample(self.variables_to_select)
76          if len(self.csp.domains[var]) > 1:
77              val = random_sample(self.csp.domains[var] -
78                                  {self.current_assignment[var]})
79              self.display(2,self.number_of_steps,": Assigning",var,"=",val)
80              self.current_assignment[var]=val
81              for varcon in self.csp.var_to_const[var]:
82                  if varcon.holds(self.current_assignment):
83                      if varcon in self.conflicts:
84                          self.conflicts.remove(varcon)
85                  else:
86                      if varcon not in self.conflicts:
87                          self.conflicts.add(varcon)
88              self.display(2,"   Number of conflicts",len(self.conflicts))
89          if not self.conflicts:
90              self.display(1,"Solution found:", self.current_assignment,
91                              "in", self.number_of_steps,"steps")
92              return self.number_of_steps
93      self.display(1,"No solution in",self.number_of_steps,"steps",
94                  len(self.conflicts),"conflicts remain")
95      return None
```

Exercise 4.9 This makes no attempt to find the best alternative value for a variable. Modify the code so that after selecting a variable it selects a value the reduces the number of conflicts by the most. Have a parameter that specifies the probability that the best value is chosen.

4.4.2   Two-Stage Choice

This is the top-level searching algorithm that maintains a priority queue of variables ordered by (the negative of) the number of conflicts, so that the variable with the most conflicts is selected first. If there is no current priority queue of variables, one is created.

The main complexity here is to maintain the priority queue. This uses the dictionary *var_differential* which specifies how much the values of variables should change. This is used with the updatable queue (page 68) to find a variable with the most conflicts.

─────────────── cspSLS.py — (continued) ───────────────

```
97  def search_with_var_pq(self,max_steps, prob_best=1.0, prob_anycon=1.0):
98      """search with a priority queue of variables.
```

```
99              This is used to select a variable with the most conflicts.
100             """
101         if not self.variable_pq:
102             self.create_pq()
103         pick_best_or_con = prob_best + prob_anycon
104         for i in range(max_steps):
105             self.number_of_steps +=1
106             randnum = random.random()
107             ## Pick a variable
108             if randnum < prob_best: # pick best variable
109                 var,oldval = self.variable_pq.top()
110             elif randnum < pick_best_or_con: # pick a variable in a conflict
111                 con = random_sample(self.conflicts)
112                 var = random_sample(con.scope)
113             else: #pick any variable that can be selected
114                 var = random_sample(self.variables_to_select)
115             if len(self.csp.domains[var]) > 1: # var has other values
116                 ## Pick a value
117                 val = random_sample(self.csp.domains[var] -
118                                     {self.current_assignment[var]})
119                 self.display(2,"Assigning",var,val)
120                 ## Update the priority queue
121                 var_differential = {}
122                 self.current_assignment[var]=val
123                 for varcon in self.csp.var_to_const[var]:
124                     self.display(3,"Checking",varcon)
125                     if varcon.holds(self.current_assignment):
126                         if varcon in self.conflicts: #was incons, now consis
127                             self.display(3,"Became consistent",varcon)
128                             self.conflicts.remove(varcon)
129                             for v in varcon.scope: # v is in one fewer conflicts
130                                 var_differential[v] = var_differential.get(v,0)-1
131                     else:
132                         if varcon not in self.conflicts: # was consis, not now
133                             self.display(3,"Became inconsistent",varcon)
134                             self.conflicts.add(varcon)
135                             for v in varcon.scope: # v is in one more conflicts
136                                 var_differential[v] = var_differential.get(v,0)+1
137                 self.variable_pq.update_each_priority(var_differential)
138                 self.display(2,"Number of conflicts",len(self.conflicts))
139             if not self.conflicts: # no conflicts, so solution found
140                 self.display(1,"Solution found:", self.current_assignment,"in",
141                             self.number_of_steps,"steps")
142                 return self.number_of_steps
143         self.display(1,"No solution in",self.number_of_steps,"steps",
144                     len(self.conflicts),"conflicts remain")
145         return None
```

create_pq creates an updatable priority queue of the variables, ordered by the number of conflicts they participate in. The priority queue only includes variables in conflicts and the value of a variable is the *negative* of the number of

conflicts the variable is in. This ensures that the priority queue, which picks
the minimum value, picks a variable with the most conflicts.

```
                            cspSLS.py — (continued)
147    def create_pq(self):
148        """Create the variable to number-of-conflicts priority queue.
149        This is needed to select the variable in the most conflicts.
150
151        The value of a variable in the priority queue is the negative of the
152        number of conflicts the variable appears in.
153        """
154        self.variable_pq = Updatable_priority_queue()
155        var_to_number_conflicts = {}
156        for con in self.conflicts:
157            for var in con.scope:
158                var_to_number_conflicts[var] = var_to_number_conflicts.get(var,0)+1
159        for var,num in var_to_number_conflicts.items():
160            if num>0:
161                self.variable_pq.add(var,-num)
```

```
                            cspSLS.py — (continued)
163  def random_sample(st):
164      """selects a random element from set st"""
165      return random.sample(st,1)[0]
```

Exercise 4.10 This makes no attempt to find the best alternative value for a vari-
able. Modify the code so that after selecting a variable it selects a value the reduces
the number of conflicts by the most. Have a parameter that specifies the probabil-
ity that the best value is chosen.

Exercise 4.11 These implementations always select a value for the variable se-
lected that is different from its current value (if that is possible). Change the code
so that it does not have this restriction (so it can leave the value the same). Would
you expect this code to be faster? Does it work worse (or better)?

4.4.3   Updatable Priority Queues

An **updatable priority queue** is a priority queue, where key-value pairs can be
stored, and the pair with the smallest key can be found and removed quickly,
and where the values can be updated. This implementation follows the idea
of `http://docs.python.org/3.5/library/heapq.html`, where the updated ele-
ments are marked as removed. This means that the priority queue can be used
unmodified. However, this might be expensive if changes are more common
than popping (as might happen if the probability of choosing the best is close
to zero).

In this implementation, the equal values are sorted randomly. This is achieved
by having the elements of the heap being $[val, rand, elt]$ triples, where the sec-
ond element is a random number. Note that Python requires this to be a list,
not a tuple, as the tuple cannot be modified.

_____ cspSLS.py — (continued) _____

```python
class Updatable_priority_queue(object):
    """A priority queue where the values can be updated.
    Elements with the same value are ordered randomly.

    This code is based on the ideas described in
    http://docs.python.org/3.3/library/heapq.html
    It could probably be done more efficiently by
    shuffling the modified element in the heap.
    """
    def __init__(self):
        self.pq = [] # priority queue of [val,rand,elt] triples
        self.elt_map = {} # map from elt to [val,rand,elt] triple in pq
        self.REMOVED = "*removed*" # a string that won't be a legal element
        self.max_size=0

    def add(self,elt,val):
        """adds elt to the priority queue with priority=val.
        """
        assert val <= 0,val
        assert elt not in self.elt_map, elt
        new_triple = [val, random.random(),elt]
        heapq.heappush(self.pq, new_triple)
        self.elt_map[elt] = new_triple

    def remove(self,elt):
        """remove the element from the priority queue"""
        if elt in self.elt_map:
            self.elt_map[elt][2] = self.REMOVED
            del self.elt_map[elt]

    def update_each_priority(self,update_dict):
        """update values in the priority queue by subtracting the values in
        update_dict from the priority of those elements in priority queue.
        """
        for elt,incr in update_dict.items():
            if incr != 0:
                newval = self.elt_map.get(elt,[0])[0] - incr
                assert newval <= 0, str(elt)+":"+str(newval+incr)+"-"+str(incr)
                self.remove(elt)
                if newval != 0:
                    self.add(elt,newval)

    def pop(self):
        """Removes and returns the (elt,value) pair with minimal value.
        If the priority queue is empty, IndexError is raised.
        """
        self.max_size = max(self.max_size, len(self.pq)) # keep statistics
        triple = heapq.heappop(self.pq)
        while triple[2] == self.REMOVED:
```

```
216            triple = heapq.heappop(self.pq)
217        del self.elt_map[triple[2]]
218        return triple[2], triple[0] # elt, value
219
220    def top(self):
221        """Returns the (elt,value) pair with minimal value, without removing it.
222        If the priority queue is empty, IndexError is raised.
223        """
224        self.max_size = max(self.max_size, len(self.pq)) # keep statistics
225        triple = self.pq[0]
226        while triple[2] == self.REMOVED:
227            heapq.heappop(self.pq)
228            triple = self.pq[0]
229        return triple[2], triple[0] # elt, value
230
231    def empty(self):
232        """returns True iff the priority queue is empty"""
233        return all(triple[2] == self.REMOVED for triple in self.pq)
```

4.4.4   Plotting Runtime Distributions

Runtime_distribution uses matplotlib to plot runtime distributions. Here the
runtime is a misnomer as we are only plotting the number of steps, not the
time. Computing the runtime is non-trivial as many of the runs have a very
short runtime. To compute the time accurately would require running the same
code, with the same random seed, multiple times to get a good estimate of the
runtime. This is left as an exercise.

_____ cspSLS.py — (continued) _____

```
235  import matplotlib.pyplot as plt
236
237  class Runtime_distribution(object):
238      def __init__(self, csp, xscale='log'):
239          """Sets up plotting for csp
240          xscale is either 'linear' or 'log'
241          """
242          self.csp = csp
243          plt.ion()
244          plt.xlabel("Number of Steps")
245          plt.ylabel("Cumulative Number of Runs")
246          plt.xscale(xscale) # Makes a 'log' or 'linear' scale
247
248      def plot_run(self,num_runs=100,max_steps=1000, prob_best=1.0, prob_anycon=1.0):
249          stats = []
250          SLSearcher.max_display_level, temp_mdl = 0, SLSearcher.max_display_level # no display
251          for i in range(num_runs):
252              searcher = SLSearcher(self.csp)
253              num_steps = searcher.search(max_steps, prob_best, prob_anycon)
254              if num_steps:
```

```
255                  stats.append(num_steps)
256          stats.sort()
257          if prob_best >= 1.0:
258              label = "P(best)=1.0"
259          else:
260              p_ac = min(prob_anycon, 1-prob_best)
261              label = "P(best)=%.2f, P(ac)=%.2f" % (prob_best, p_ac)
262          plt.plot(stats,range(len(stats)),label=label)
263          plt.legend(loc="upper left")
264          #plt.draw()
265          SLSearcher.max_display_level= temp_mdl #restore display
```

4.4.5   Testing

```
_____ cspSLS.py — (continued) _____

267  from cspExamples import test
268  def sls_solver(csp,prob_best=0.7):
269      """stochastic local searcher (prob_best=0.7)"""
270      se0 = SLSearcher(csp)
271      se0.search(1000,prob_best)
272      return se0.current_assignment
273  def any_conflict_solver(csp):
274      """stochastic local searcher (any-conflict)"""
275      return sls_solver(csp,0)
276
277  if __name__ == "__main__":
278      test(sls_solver)
279      test(any_conflict_solver)
280
281  from cspExamples import csp1, csp2, crossword1
282
283  ## Test Solving CSPs with Search:
284  #se1 = SLSearcher(csp1); print(se1.search(100))
285  #se2 = SLSearcher(csp2); print(se2.search(1000,1.0)) # greedy
286  #se2 = SLSearcher(csp2); print(se2.search(1000,0)) # any_conflict
287  #se2 = SLSearcher(csp2); print(se2.search(1000,0.7)) # 70% greedy; 30% any_conflict
288  #SLSearcher.max_display_level=2 #more detailed display
289  #se3 = SLSearcher(crossword1); print(se3.search(100),0.7)
290  #p = Runtime_distribution(csp2)
291  #p.plot_run(1000,1000,0) # any_conflict
292  #p.plot_run(1000,1000,1.0) # greedy
293  #p.plot_run(1000,1000,0.7) # 70% greedy; 30% any_conflict
```

Exercise 4.12   Modify this to plot the runtime, instead of the number of steps.
To measure runtime use *timeit* (https://docs.python.org/3.5/library/timeit.
html). Small runtimes are inaccurate, so timeit can run the same code multiple times. Stochastic local algorithms give different runtimes each time called.
To make the timing meaningful, you need to make sure the random seed is the same for each repeated call (see random.getstate and random.setstate in https:

`//docs.python.org/3.5/library/random.html`). Because the runtime for different seeds can vary a great deal, for each seed, you should start with 1 iteration and multiplying it by, say 10, until the time is greater than 0.2 seconds. Make sure you plot the average time for each run. Before you start, try to estimate the total runtime, so you will be able to tell if there is a problem with the algorithm stopping.

Chapter 5

Propositions and Inference

5.1 Representing Knowledge Bases

A clause consists of a head (an atom) and a body. A body is represented as a list of atoms. Atoms are represented as strings.

_____ logicProblem.py — Representations Logics _____
```
11  class Clause(object):
12      """A definite clause"""
13
14      def __init__(self,head,body=[]):
15          """clause with atom head and lost of atoms body"""
16          self.head=head
17          self.body = body
18
19      def __str__(self):
20          """returns the string representation of a clause.
21          """
22          if self.body:
23              return self.head + " <- " + " & ".join(self.body) + "."
24          else:
25              return self.head + "."
```

An askable atom can be asked of the user. The user can respond in English or French or just with a "y".

_____ logicProblem.py — (continued) _____
```
27  class Askable(object):
28      """An askable atom"""
29
30      def __init__(self,atom):
31          """clause with atom head and lost of atoms body"""
```

```
32              self.atom=atom
33
34      def __str__(self):
35          """returns the string representation of a clause."""
36          return "askable " + self.atom + "."
37
38  def yes(ans):
39      """returns true if the answer is yes in some form"""
40      return ans.lower() in ['yes', 'yes.', 'oui', 'oui.', 'y', 'y.'] # bilingual
```

A knowledge base is a list of clauses and askables. In order to make top-down inference faster, this creates a dictionary that maps each atoms into the set of clauses with that atom in the head.

─────────────────── logicProblem.py — (continued) ───────────────────

```
42  from utilities import Displayable
43
44  class KB(Displayable):
45      """A knowledge base consists of a set of clauses.
46      This also creates a dictionary to give fast access to the clauses with an atom in head.
47      """
48      def __init__(self, statements=[]):
49          self.statements = statements
50          self.clauses = [c for c in statements if isinstance(c, Clause)]
51          self.askables = [c.atom for c in statements if isinstance(c, Askable)]
52          self.atom_to_clauses = {} # dictionary giving clauses with atom as head
53          for c in self.clauses:
54              if c.head in self.atom_to_clauses:
55                  self.atom_to_clauses[c.head].add(c)
56              else:
57                  self.atom_to_clauses[c.head] = {c}
58
59      def clauses_for_atom(self,a):
60          """returns set of clauses with atom a as the head"""
61          if a in self.atom_to_clauses:
62              return self.atom_to_clauses[a]
63          else:
64              return set()
65
66      def __str__(self):
67          """returns a string representation of this knowledge base.
68          """
69          return '\n'.join([str(c) for c in self.statements])
```

Here is a trivial example (I think therefore I am) using in the unit tests:

─────────────────── logicProblem.py — (continued) ───────────────────

```
71  triv_KB = KB([
72      Clause('i_am', ['i_think']),
73      Clause('i_think'),
74      Clause('i_smell', ['i_exist'])
75      ])
```

Here is a representation of the electrical domain of the textbook:

_____logicProblem.py — (continued)_____

```
77  elect = KB([
78      Clause('light_l1'),
79      Clause('light_l2'),
80      Clause('ok_l1'),
81      Clause('ok_l2'),
82      Clause('ok_cb1'),
83      Clause('ok_cb2'),
84      Clause('live_outside'),
85      Clause('live_l1', ['live_w0']),
86      Clause('live_w0', ['up_s2','live_w1']),
87      Clause('live_w0', ['down_s2','live_w2']),
88      Clause('live_w1', ['up_s1', 'live_w3']),
89      Clause('live_w2', ['down_s1','live_w3' ]),
90      Clause('live_l2', ['live_w4']),
91      Clause('live_w4', ['up_s3','live_w3' ]),
92      Clause('live_p_1', ['live_w3']),
93      Clause('live_w3', ['live_w5', 'ok_cb1']),
94      Clause('live_p_2', ['live_w6']),
95      Clause('live_w6', ['live_w5', 'ok_cb2']),
96      Clause('live_w5', ['live_outside']),
97      Clause('lit_l1', ['light_l1', 'live_l1', 'ok_l1']),
98      Clause('lit_l2', ['light_l2', 'live_l2', 'ok_l2']),
99      Askable('up_s1'),
100     Askable('down_s1'),
101     Askable('up_s2'),
102     Askable('down_s2'),
103     Askable('up_s3'),
104     Askable('down_s2')
105     ])
106
107 # print(kb)
```

5.2 Bottom-up Proofs

fixed_point computes the fixed point of the knowledge base *kb*.

_____logicBottomUp.py — Bottom-up Proof Procedure for Definite Clauses_____

```
11  from logicProblem import yes
12
13  def fixed_point(kb):
14      """Returns the fixed point of knowledge base kb.
15      """
16      fp = ask_askables(kb)
17      added = True
18      while added:
19          added = False # added is true when an atom was added to fp this iteration
```

```
20          for c in kb.clauses:
21              if c.head not in fp and all(b in fp for b in c.body):
22                  fp.add(c.head)
23                  added = True
24                  kb.display(2,c.head,"added to fp due to clause",c)
25      return fp
26
27  def ask_askables(kb):
28      return {at for at in kb.askables if yes(input("Is "+at+" true? "))}
```

Testing:

```
_____logicBottomUp.py — (continued)_____
30  from logicProblem import triv_KB
31  def test():
32      fp = fixed_point(triv_KB)
33      assert fp == {'i_am','i_think'}, "triv_KB gave result "+str(fp)
34      print("Passed unit test")
35  if __name__ == "__main__":
36      test()
37
38  from logicProblem import elect
39  # elect.max_display_level=3 # give detailed trace
40  # fixed_point(elect)
```

Exercise 5.1 It is not very user-friendly to ask all of the askables up-front. Implement ask-the-user so that questions are only asked if useful, and are not re-asked. For example, if there is a clause $h \leftarrow a \wedge b \wedge c \wedge d \wedge e$, where c and e are askable, c and e only need to be asked if a, b, d are all in *fp* and they have not been asked before. Askable e only needs to be asked if the user says "yes" to c. Askable c doesn't need to be asked if the user previously replied "no" to e.

   This form of ask-the-user can ask a different set of questions than the top-down interpreter that asks questions when encountered. Give an example where they ask different questions (neither set of questions asked is a subset of the other).

Exercise 5.2 This algorithm runs in time $O(n^2)$, where n is the number of clauses, for a bounded number of elements in the body; each iteration goes through each of the clauses, and in the worst case, it will do an iteration for each clause. It is possible to implement this in time $O(n)$ time by creating an index that maps an atom to the set of clauses with that atom in the body. Implement this. What is its complexity as a function of n and b, the maximum number of atoms in the body of a clause?

Exercise 5.3 It is possible to be asymptitocally more efficient (in terms of b) than the method in the previous question by noticing that each element of the body of clause only needs to be checked once. For example, the clause $a \leftarrow b \wedge c \wedge d$, needs only be considered when b is added to *fp*. Once b is added to *fp*, if c is already in *pf*, we know that a can be added as soon as d is added. Implement this. What is its complexity as a function of n and b, the maximum number of atoms in the body of a clause?

5.3 Top-down Proofs

prove(*kb*, *goal*) is used to prove *goal* from a knowledge base, *kb*, where a *goal* is a list of atoms. It returns *True* if *kb* ⊢ *goal*. The *indent* is used when tracing the code (and doesn't need to have a non-default value).

```
_____logicTopDown.py — Top-down Proof Procedure for Definite Clauses _____
11 from logicProblem import yes
12
13 def prove(kb, ans_body, indent=""):
14     """returns True if kb |- ans_body
15     ans_body is a list of atoms to be proved
16     """
17     kb.display(2,indent,'yes <-',' & '.join(ans_body))
18     if ans_body:
19         selected = ans_body[0] # select first atom from ans_body
20         if selected in kb.askables:
21             return (yes(input("Is "+selected+" true? "))
22                     and prove(kb,ans_body[1:],indent+" "))
23         else:
24             return any(prove(kb,cl.body+ans_body[1:],indent+" ")
25                     for cl in kb.clauses_for_atom(selected))
26     else:
27         return True   # empty body is true
```

Testing:

```
_____logicTopDown.py — (continued) _____
29 from logicProblem import triv_KB
30 def test():
31     a1 = prove(triv_KB,['i_am'])
32     assert a1, "triv_KB proving i_am gave "+str(a1)
33     a2 = prove(triv_KB,['i_smell'])
34     assert not a2, "triv_KB proving i_smell gave "+str(a2it)
35     print("Passed unit tests")
36 if __name__ == "__main__":
37     test()
38 # try
39 from logicProblem import elect
40 # elect.max_display_level=3 # give detailed trace
41 # prove(elect,['live_w6'])
42 # prove(elect,['lit_l1'])
```

Exercise 5.4 This code can re-ask a question multiple times. Implement this code so that it only asks a question once and remembers the answer. Also implement a function to forget the answers.

Exercise 5.5 What search method is this using? Implement the search interface so that it can use A^* or other searching methods. Define an admissible heuristic that is not always 0.

5.4 Assumables

Atom *a* can be made assumable by including *Assumable*(*a*) in the knowledge base. A knowledge base that can include assumables is declared with *KBA*.

_____logicAssumables.py — Definite clauses with assumables _____
```python
11   from logicProblem import Clause, Askable, KB, yes
12
13   class Assumable(object):
14       """An askable atom"""
15
16       def __init__(self,atom):
17           """clause with atom head and lost of atoms body"""
18           self.atom = atom
19
20       def __str__(self):
21           """returns the string representation of a clause.
22           """
23           return "assumable " + self.atom + "."
24
25   class KBA(KB):
26       """A knowledge base that can include assumables"""
27       def __init__(self,statements):
28           self.assumables = [c.atom for c in statements if isinstance(c, Assumable)]
29           KB.__init__(self,statements)
```

The top-down Horn clause interpreter, *prove_all_ass* returns a list of the sets of assumables that imply *ans_body*. This list will contain all of the minimal sets of assumables, but can also find non-minimal sets, and repeated sets, if they can be generated with separate proofs. The set *assumed* is the set of assumables already assumed.

_____logicAssumables.py — (continued) _____
```python
31       def prove_all_ass(self, ans_body, assumed=set()):
32           """returns a list of sets of assumables that extends assumed
33           to imply ans_body from self.
34           ans_body is a list of atoms (it is the body of the answer clause).
35           assumed is a set of assumables already assumed
36           """
37           if ans_body:
38               selected = ans_body[0] # select first atom from ans_body
39               if selected in self.askables:
40                   if yes(input("Is "+selected+" true? ")):
41                       return self.prove_all_ass(ans_body[1:],assumed)
42                   else:
43                       return []  # no answers
44               elif selected in self.assumables:
45                   return self.prove_all_ass(ans_body[1:],assumed|{selected})
46               else:
47                   return [ass
48                           for cl in self.clauses_for_atom(selected)
```

```
49                          for ass in self.prove_all_ass(cl.body+ans_body[1:],assumed)
50                          ] # union of answers for each clause with head=selected
51          else:                    # empty body
52              return [assumed]  # one answer
53
54      def conflicts(self):
55          """returns a list of minimal conflicts"""
56          return minsets(self.prove_all_ass(['false']))
```

Given a list of sets, *minsets* returns a list of the minimal sets in the list. For example, *minsets*([{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}]) returns [{2, 3}, {2, 4, 5}].

———————————— logicAssumables.py — (continued) ————————————

```
58  def minsets(ls):
59      """ls is a list of sets
60      returns a list of minimal sets in ls
61      """
62      ans = []     # elements known to be minimal
63      for c in ls:
64          if not any(c1<c for c1 in ls) and not any(c1 <= c for c1 in ans):
65              ans.append(c)
66      return ans
67
68  # minsets([{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}])
```

Warning: *minsets* works for a list of sets or for a set of (frozen) sets, but it does not work for a generator of sets. For example, try to predict and then test:

```
minsets(e for e in [{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}])
```

The diagnoses can be constructed from the (minimal) conflicts as follows. This also works if there are non-minimal conflicts, but is not as efficient.

———————————— logicAssumables.py — (continued) ————————————

```
69  def diagnoses(cons):
70      """cons is a list of (minimal) conflicts.
71      returns a list of diagnoses."""
72      if cons == []:
73          return [set()]
74      else:
75          return minsets([({e}|d)              # | is set union
76                          for e in cons[0]
77                          for d in diagnoses(cons[1:])])
```

Test cases:

———————————— logicAssumables.py — (continued) ————————————

```
80  electa = KBA([
81      Clause('light_l1'),
82      Clause('light_l2'),
83      Assumable('ok_l1'),
84      Assumable('ok_l2'),
```

```
85   │       Assumable('ok_s1'),
86   │       Assumable('ok_s2'),
87   │       Assumable('ok_s3'),
88   │       Assumable('ok_cb1'),
89   │       Assumable('ok_cb2'),
90   │       Assumable('live_outside'),
91   │       Clause('live_l1', ['live_w0']),
92   │       Clause('live_w0', ['up_s2','ok_s2','live_w1']),
93   │       Clause('live_w0', ['down_s2','ok_s2','live_w2']),
94   │       Clause('live_w1', ['up_s1', 'ok_s1', 'live_w3']),
95   │       Clause('live_w2', ['down_s1', 'ok_s1','live_w3' ]),
96   │       Clause('live_l2', ['live_w4']),
97   │       Clause('live_w4', ['up_s3','ok_s3','live_w3' ]),
98   │       Clause('live_p_1', ['live_w3']),
99   │       Clause('live_w3', ['live_w5', 'ok_cb1']),
100  │       Clause('live_p_2', ['live_w6']),
101  │       Clause('live_w6', ['live_w5', 'ok_cb2']),
102  │       Clause('live_w5', ['live_outside']),
103  │       Clause('lit_l1', ['light_l1', 'live_l1', 'ok_l1']),
104  │       Clause('lit_l2', ['light_l2', 'live_l2', 'ok_l2']),
105  │       Askable('up_s1'),
106  │       Askable('down_s1'),
107  │       Askable('up_s2'),
108  │       Askable('down_s2'),
109  │       Askable('up_s3'),
110  │       Askable('down_s2'),
111  │       Askable('dark_l1'),
112  │       Askable('dark_l2'),
113  │       Clause('false', ['dark_l1', 'lit_l1']),
114  │       Clause('false', ['dark_l2', 'lit_l2'])
115  │       ])
116  │ # electa.prove_all_ass(['false'])
117  │ # cs=electa.conflicts()
118  │ # print(cs)
119  │ # diagnoses(cs)      # diagnoses from conflicts
```

Exercise 5.6 To implement a version of *conflicts* that never generates non-minimal conflicts, modify *prove_all_ass* to implement iterative deepening on the number of assumables used in a proof, and prune any set of assumables that is a superset of a conflict.

Exercise 5.7 Implement *explanations*(*self*, *body*), where *body* is a list of atoms, that returns the a list of the minimal explanations of the body. This does not require modification of *prove_all_ass*.

Exercise 5.8 Implement *explanations*, as in the previous question, so that it never generates non-minimal explanations. Hint: modify *prove_all_ass* to implement iterative deepening on the number of assumptions, generating conflicts and explanations together, and pruning as early as possible.

Chapter 6

Planning with Certainty

6.1 Representing Actions and Planning Problems

The STRIPS representation of an action consists of:

- preconditions: a dictionary of *feature:value* pairs that specifies that the feature must have this value for the action to be possible.

- effects: a dictionary of *feature:value* pairs that are made true by this action. In particular, a feature in the dictionary has the corresponding value (and not its previous value) after the action, and a feature not in the dictionary keeps its old value.

_____stripsProblem.py — STRIPS Representations of Actions _____
```
11  class Strips(object):
12      def __init__(self, preconditions, effects, cost=1):
13          """
14          defines the STRIPS represtation for an action:
15          * preconditions is feature:value dictionary that must hold
16          for the action to be carried out
17          * effects is a feature:value map that this action makes
18          true. The action changes the value of any feature specified
19          here, and leaves other properties unchanged.
20          * cost is the cost of the action
21          """
22          self.preconditions = preconditions
23          self.effects = effects
24          self.cost = cost
```

A STRIPS domain consists of:

- A set of actions.

- A dictionary that maps each feature into a set of possible values for the feature.

- A dictionary that maps each action into a STRIPS representation of the action.

_____ stripsProblem.py — (continued) _____

```
26   class STRIPS_domain(object):
27       def __init__(self, feats_vals, strips_map):
28           """Problem domain
29           feats_vals is a feature:domain dictionary,
30                   mapping each feature to its domain
31           strips_map is an action:strips dictionary,
32                   mapping each action to its Strips representation
33           """
34           self.actions = set(strips_map) # set of all actions
35           self.feats_vals = feats_vals
36           self.strips_map = strips_map
```

6.1.1 Robot Delivery Domain

The following specifies the robot delivery domain of Chapter 8.

_____ stripsProblem.py — (continued) _____

```
38   boolean = {True, False}
39   delivery_domain = STRIPS_domain(
40       {'RLoc':{'cs', 'off', 'lab', 'mr'}, 'RHC':boolean, 'SWC':boolean,
41        'MW':boolean, 'RHM':boolean},        #feaures:values dictionary
42       {'mc_cs': Strips({'RLoc':'cs'}, {'RLoc':'off'}),
43        'mc_off': Strips({'RLoc':'off'}, {'RLoc':'lab'}),
44        'mc_lab': Strips({'RLoc':'lab'}, {'RLoc':'mr'}),
45        'mc_mr': Strips({'RLoc':'mr'}, {'RLoc':'cs'}),
46        'mcc_cs': Strips({'RLoc':'cs'}, {'RLoc':'mr'}),
47        'mcc_off': Strips({'RLoc':'off'}, {'RLoc':'cs'}),
48        'mcc_lab': Strips({'RLoc':'lab'}, {'RLoc':'off'}),
49        'mcc_mr': Strips({'RLoc':'mr'}, {'RLoc':'lab'}),
50        'puc': Strips({'RLoc':'cs', 'RHC':False}, {'RHC':True}),
51        'dc': Strips({'RLoc':'off', 'RHC':True}, {'RHC':False, 'SWC':False}),
52        'pum': Strips({'RLoc':'mr','MW':True}, {'RHM':True,'MW':False}),
53        'dm': Strips({'RLoc':'off', 'RHM':True}, {'RHM':False})
54       } )
```

A planning problem consists of a planning domain, an initial state, and a goal. The goal does not need to fully specify the final state.

_____ stripsProblem.py — (continued) _____

```
56   class Planning_problem(object):
```
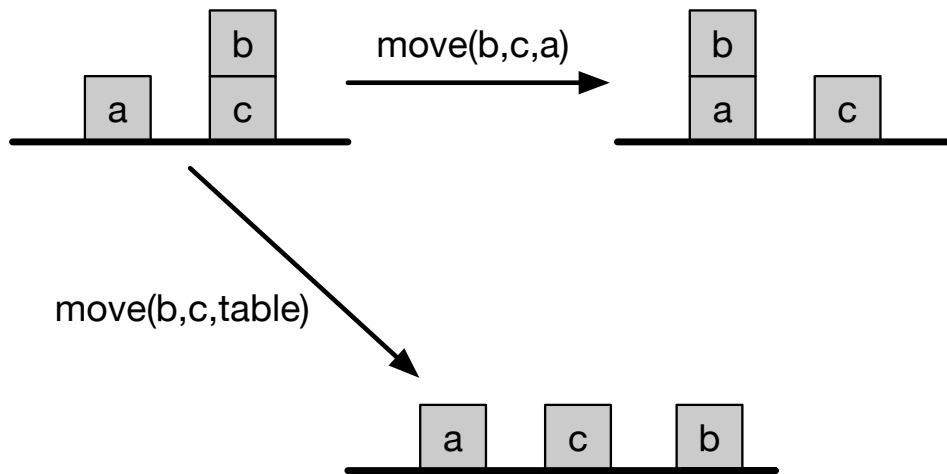
Figure 6.1: Blocks world with two actions

```
57      def __init__(self, prob_domain, initial_state, goal):
58          """
59          a planning problem consists of
60          * a planning domain
61          * the initial state
62          * a goal
63          """
64          self.prob_domain = prob_domain
65          self.initial_state = initial_state
66          self.goal = goal
67
68  problem0 = Planning_problem(delivery_domain,
69                          {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,
70                           'RHM':False},
71                          {'RLoc':'off'})
72  problem1 = Planning_problem(delivery_domain,
73                          {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,
74                           'RHM':False},
75                          {'SWC':False})
76  problem2 = Planning_problem(delivery_domain,
77                          {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,
78                           'RHM':False},
79                          {'SWC':False, 'MW':False, 'RHM':False})
```

6.1.2 Blocks World

The blocks world consist of blocks and a table. Each block can be on the table or on another block. A block can only have one other block on top of it. Figure 6.1 shows 3 states with some of the actions between them. The following

represents the blocks world. Note that the actions and the conditions are all
strings.

```
────────────────── stripsProblem.py — (continued) ──────────────────
81  ### blocks world
82  def move(x,y,z):
83      """string for the 'move' action"""
84      return 'move_'+x+'_from_'+y+'_to_'+z
85  def on(x,y):
86      """string for the 'on' feature"""
87      return x+'_on_'+y
88  def clear(x):
89      """string for the 'clear' feature"""
90      return 'clear_'+x
91  def create_blocks_world(blocks = ['a','b','c','d']):
92      blocks_and_table = blocks+['table']
93      stmap = {move(x,y,z):Strips({on(x,y):True, clear(x):True, clear(z):True},
94                              {on(x,z):True, on(x,y):False, clear(y):True, clear(z):False})
95                  for x in blocks
96                  for y in blocks_and_table
97                  for z in blocks
98                  if x!=y and y!=z and z!=x}
99      stmap.update({move(x,y,'table'):Strips({on(x,y):True, clear(x):True},
100                             {on(x,'table'):True, on(x,y):False, clear(y):True})
101                 for x in blocks
102                 for y in blocks
103                 if x!=y})
104     feats_vals = {on(x,y):boolean for x in blocks for y in blocks_and_table}
105     feats_vals.update({clear(x):boolean for x in blocks_and_table})
106     return STRIPS_domain(feats_vals, stmap)
```

This is a classic example, with 3 blocks, and the goal consists of two conditions.

```
────────────────── stripsProblem.py — (continued) ──────────────────
108  blocks1dom = create_blocks_world(['a','b','c'])
109  blocks1 = Planning_problem(blocks1dom,
110      {on('a','table'):True,clear('a'):True, clear('b'):True,on('b','c'):True,
111       on('c','table'):True,clear('c'):False}, # initial state
112      {on('a','b'):True, on('c','a'):True}) #goal
```

This is a problem of inverting a tower of size 4.

```
────────────────── stripsProblem.py — (continued) ──────────────────
114  blocks2dom = create_blocks_world(['a','b','c','d'])
115  tower4 = {clear('a'):True, on('a','b'):True, clear('b'):False,
116      on('b','c'):True, clear('c'):False, on('c','d'):True,
117      clear('b'):False, on('d','table'):True}
118  blocks2 = Planning_problem(blocks2dom,
119      tower4, # initial state
120      {on('d','c'):True,on('c','b'):True,on('b','a'):True}) #goal
```

Moving bottom block to top of a tower of size 4.

```
122  blocks3 = Planning_problem(blocks2dom,
123      tower4, # initial state
124      {on('d','a'):True, on('a','b'):True, on('b','c'):True}) #goal
```

Exercise 6.1 Represent the problem of given a tower of 4 blocks (*a* on *b* on *c* on *d* on table), the goal is to have a tower with the previous top block on the bottom (*b* on *c* on *d* on *a*). Do not include the table in your goal (the goal does not care whether *a* is on the table). [Before you run the program, estimate how many steps it will take to solve this.] How many steps does an optimal planner take?

Exercise 6.2 The representation of the state does not include negative *on* facts. Does it need to? Why or why not? (Note that this may depend on the planner; write your answer with respect to particular planners.)

Exercise 6.3 It is possible to write the representation of the problem without using *clear*, where *clear*(*x*) means nothing is on *x*. Change the definition of the blocks world so that it does not use *clear* but uses *on* being false instead. Does this work better for any of the planners? (Does this change an answer to the previous question?)

6.2 Forward Planning

> To run the demo, in folder "aipython", load "stripsForwardPlanner.py", and copy and paste the commented-out example queries at the bottom of that file.

In a forward planner, a node is a state. A state consists of an assignment, which is a variable:value dictionary. In order to be able to do multiple-path pruning, we need to define a hash function, and equality between states.

```
11  from searchProblem import Arc, Search_problem
12  from stripsProblem import Strips, STRIPS_domain
13
14  class State(object):
15      def __init__(self,assignment):
16          self.assignment = assignment
17          self.hash_value = None
18      def __hash__(self):
19          if self.hash_value is None:
20              self.hash_value = hash(frozenset(self.assignment.items()))
21          return self.hash_value
22      def __eq__(self,st):
23          return self.assignment == st.assignment
24      def __str__(self):
25          return str(self.assignment)
```

In order to define a search problem (page 31), we need to define the goal condition, the start nodes, the neighbours, and (optionally) a heuristic function. Here *zero* is the default heuristic function.

```
_____ stripsForwardPlanner.py — (continued) _____
27  def zero(*args,**nargs):
28      """always returns 0"""
29      return 0
30
31  class Forward_STRIPS(Search_problem):
32      """A search problem from a planning problem where:
33      * a node is a state object.
34      * the dynamics are specified by the STRIPS representation of actions
35      """
36      def __init__(self, planning_problem, heur=zero):
37          """creates a forward seach space from a planning problem.
38          heur(state,goal) is a heuristic function,
39             an underestimate of the cost from state to goal, where
40             both state and goals are feature:value dictionaries.
41          """
42          self.prob_domain = planning_problem.prob_domain
43          self.initial_state = State(planning_problem.initial_state)
44          self.goal = planning_problem.goal
45          self.heur = heur
46
47      def is_goal(self, state):
48          """is True if node is a goal.
49
50          Every goal feature has the same value in the state and the goal."""
51          state_asst = state.assignment
52          return all(prop in state_asst and state_asst[prop]==self.goal[prop]
53                      for prop in self.goal)
54
55      def start_node(self):
56          """returns start node"""
57          return self.initial_state
58
59      def neighbors(self,state):
60          """returns neighbors of state in this problem"""
61          cost=1
62          state_asst = state.assignment
63          return [ Arc(state,self.effect(act,state_asst),cost,act)
64                  for act in self.prob_domain.actions
65                  if self.possible(act,state_asst)]
66
67      def possible(self,act,state_asst):
68          """True if act is possible in state.
69          act is possible if all of its preconditions have the same value in the state"""
70          preconds = self.prob_domain.strips_map[act].preconditions
71          return all(pre in state_asst and state_asst[pre]==preconds[pre]
```

```
72                        for pre in preconds)
73
74      def effect(self,act,state_asst):
75          """returns the state that is the effect of doing act given state_asst"""
76          new_state_asst = self.prob_domain.strips_map[act].effects.copy()
77          for prop in state_asst:
78              if prop not in new_state_asst:
79                  new_state_asst[prop]=state_asst[prop]
80          return State(new_state_asst)
81
82      def heuristic(self,state):
83          """in the forward planner a node is a state.
84          the heuristic is an (under)estimate of the cost
85          of going from the state to the top-level goal.
86          """
87          return self.heur(state.assignment, self.goal)
```

Here are some test cases to try.

_____ stripsForwardPlanner.py — (continued) _____

```
89  from searchBranchAndBound import DF_branch_and_bound
90  from searchGeneric import AStarSearcher
91  from searchMPP import SearcherMPP
92  from stripsProblem import problem0, problem1, problem2, blocks1, blocks2, blocks3
93
94  # AStarSearcher(Forward_STRIPS(problem1)).search() #A*
95  # SearcherMPP(Forward_STRIPS(problem1)).search() #A* with MPP
96  # DF_branch_and_bound(Forward_STRIPS(problem1),10).search() #B&B
97  # To find more than one plan:
98  # s1 = SearcherMPP(Forward_STRIPS(problem1)) #A*
99  # s1.search() #find another plan
```

6.2.1 Defining Heuristics for a Planner

Each planning domain requires its own heuristics. If you change the actions, you will need to reconsider the heuristic function, as there might then be a lower-cost path, which might make the heuristic non-admissible.

Here is an example of defining a (not very good) heuristic for the coffee delivery planning domain.

First we define the distance between two locations, which is used for the heuristics.

_____ stripsHeuristic.py — Planner with Heursitic Function _____

```
11  def dist(loc1, loc2):
12      """returns the distance from location loc1 to loc2
13      """
14      if loc1==loc2:
15          return 0
16      if {loc1,loc2} in [{'cs','lab'},{'mr','off'}]:
```

```
17          return 2
18      else:
19          return 1
```

Note that the current state is a complete description; there is a value for every feature. However the goal need not be complete; it does not need to define a value for every feature. Before checking the value for a feature in the goal, a heuristic needs to define whether the feature is defined in the goal.

──────── stripsHeuristic.py — (continued) ────────

```
21  def h1(state,goal):
22      """ the distance to the goal location, if there is one"""
23      if 'RLoc' in goal:
24          return dist(state['RLoc'], goal['RLoc'])
25      else:
26          return 0
27
28  def h2(state,goal):
29      """ the distance to the coffee shop plus getting coffee and delivering it
30      if the robot needs to get coffee
31      """
32      if ('SWC' in goal and goal['SWC']==False
33              and state['SWC']==True
34              and state['RHC']==False):
35          return dist(state['RLoc'],'cs')+3
36      else:
37          return 0
```

The maximum of the values of a set of admissible heuristics is also an admissible heuristic. The function maxh takes a number of heuristic functions as arguments, and returns a new heuristic function that takes the maximum of the values of the heuristics. For example, h1 and h2 are heuristic functions and so maxh(h1,h2) is also. maxh can take an arbitrary number of arguments.

──────── stripsHeuristic.py — (continued) ────────

```
39  def maxh(*heuristics):
40      """Returns a new heuristic function that is the maximum of the functions in heuristics.
41      heuristics is the list of arguments which must be heuristic functions.
42      """
43      return lambda state,goal: max(h(state,goal) for h in heuristics)
```

The following runs the example with and without the heuristic. (Also try using *AStarSearcher* instead of *SearcherMPP*.)

──────── stripsHeuristic.py — (continued) ────────

```
45  ##### Forward Planner #####
46  from searchGeneric import AStarSearcher
47  from searchMPP import SearcherMPP
48  from stripsForwardPlanner import Forward_STRIPS
49  from stripsProblem import problem0, problem1, problem2
50
```

```
51  def test_forward_heuristic(thisproblem=problem1):
52      print("\n***** FORWARD NO HEURISTIC")
53      print(SearcherMPP(Forward_STRIPS(thisproblem)).search())
54
55      print("\n***** FORWARD WITH HEURISTIC h1")
56      print(SearcherMPP(Forward_STRIPS(thisproblem,h1)).search())
57
58      print("\n***** FORWARD WITH HEURISTICs h1 and h2")
59      print(SearcherMPP(Forward_STRIPS(thisproblem,maxh(h1,h2))).search())
60
61  if __name__ == "__main__":
62      test_forward_heuristic()
```

Exercise 6.4 Try the forward planner with a heuristic function of just *h*1, with just *h*2 and with both. Explain how each one prunes or doesn't prune the search space.

Exercise 6.5 Create a better heuristic than *maxh*(*h*1, *h*2). Try it for a number of different problems.

Exercise 6.6 Create an admissible heuristic for the blocks world.

6.3 Regression Planning

> To run the demo, in folder "aipython", load "stripsRegressionPlanner.py", and copy and paste the commented-out example queries at the bottom of that file.

In a regression planner a node is a subgoal that need to be achieved.

A *Subgoal* object consists of an assignment, which is *variable:value* dictionary. We make it hashable so that multiple path pruning can work. The hash is only computed when necessary (and only once).

—————— stripsRegressionPlanner.py — Regression Planner with STRIPS actions ——————
```
11  from searchProblem import Arc, Search_problem
12
13  class Subgoal(object):
14      def __init__(self,assignment):
15          self.assignment = assignment
16          self.hash_value = None
17      def __hash__(self):
18          if self.hash_value is None:
19              self.hash_value = hash(frozenset(self.assignment.items()))
20          return self.hash_value
21      def __eq__(self,st):
22          return self.assignment == st.assignment
23      def __str__(self):
24          return str(self.assignment)
```

A regression search has subgoals as nodes. The initial node is the top-level goal of the planner. The goal for the search (when the search can stop) is a subgoal that holds in the initial state.

```
────────────────────── stripsRegressionPlanner.py — (continued) ──────────────────────
26  from stripsForwardPlanner import zero
27
28  class Regression_STRIPS(Search_problem):
29      """A search problem where:
30      * a node is a goal to be achieved, represented by a set of propositions.
31      * the dynamics are specified by the STRIPS representation of actions
32      """
33
34      def __init__(self, planning_problem, heur=zero):
35          """creates a regression seach space from a planning problem.
36          heur(state,goal) is a heuristic function;
37              an underestimate of the cost from state to goal, where
38              both state and goals are feature:value dictionaries
39          """
40          self.prob_domain = planning_problem.prob_domain
41          self.top_goal = Subgoal(planning_problem.goal)
42          self.initial_state = planning_problem.initial_state
43          self.heur = heur
44
45      def is_goal(self, subgoal):
46          """if subgoal is true in the initial state, a path has been found"""
47          goal_asst = subgoal.assignment
48          return all((g in self.initial_state) and (self.initial_state[g]==goal_asst[g])
49                      for g in goal_asst)
50
51      def start_node(self):
52          """the start node is the top-level goal"""
53          return self.top_goal
54
55      def neighbors(self,subgoal):
56          """returns a list of the arcs for the neighbors of subgoal in this problem"""
57          cost = 1
58          goal_asst = subgoal.assignment
59          return [ Arc(subgoal,self.weakest_precond(act,goal_asst),cost,act)
60                      for act in self.prob_domain.actions
61                      if self.possible(act,goal_asst)]
62
63      def possible(self,act,goal_asst):
64          """True if act is possible to achieve goal_asst.
65
66          the action achieves an element of the effects and
67          the action doesn't delete something that needs to be achieved and
68          the precoditions are consistent with other subgoals that need to be achieved
69          """
70          effects = self.prob_domain.strips_map[act].effects
```

```
71          preconds = self.prob_domain.strips_map[act].preconditions
72          return ( any(goal_asst[prop]==effects[prop]
73                      for prop in effects if prop in goal_asst)
74                 and all(goal_asst[prop]==effects[prop]
75                      for prop in effects if prop in goal_asst)
76                 and all(goal_asst[prop]==preconds[prop]
77                      for prop in preconds if prop not in effects and prop in goal_asst)
78                 )
79
80      def weakest_precond(self,act,goal_asst):
81          """returns the subgoal that must be true so goal_asst holds after act"""
82          new_asst = self.prob_domain.strips_map[act].preconditions.copy()
83          for g in goal_asst:
84              if g not in self.prob_domain.strips_map[act].effects:
85                  new_asst[g] = goal_asst[g]
86          return Subgoal(new_asst)
87
88      def heuristic(self,subgoal):
89          """in the regression planner a node is a subgoal.
90          the heuristic is an (under)estimate of the cost of going from the initial state to subgoal.
91          """
92          return self.heur(self.initial_state, subgoal.assignment)
```

_____ stripsRegressionPlanner.py — (continued) _____

```
94  from searchBranchAndBound import DF_branch_and_bound
95  from searchGeneric import AStarSearcher
96  from searchMPP import SearcherMPP
97  from stripsProblem import problem0, problem1, problem2
98
99  # AStarSearcher(Regression_STRIPS(problem1)).search() #A*
100 # SearcherMPP(Regression_STRIPS(problem1)).search() #A* with MPP
101 # DF_branch_and_bound(Regression_STRIPS(problem1),10).search() #B&B
```

Exercise 6.7 Multiple path pruning could be used to prune more than the current code. In particular, if the current node contains more conditions than a previously visited node, it can be pruned. For example, if {*a* : *True*, *b* : *False*} has been visited, then any node that is a superset, e.g., {*a* : *True*, *b* : *False*, *d* : *True*}, need not be expanded. If the simpler subgoal does not lead to a solution, the more complicated one wont either. Implement this more severe pruning. (Hint: This may require modifications to the searcher.)

Exercise 6.8 It is possible that, as knowledge of the domain, that some assignment of values to variables can never be achieved. For example, the robot cannot be holding mail when there is mail waiting (assuming it isn't holding mail initially). An assignment of values to (some of the) variables is incompatible if no possible (reachable) state can include that assignment. For example, {$'MW'$: *True*,$' RHM'$: *True*} is an incompatible assignment. This information may be useful information for a planner; there is no point in trying to achieve these together. Define a subclass of *STRIPS_domain* that can accept a list of incompatible

assignments. Modify the regression planner code to use such a list of incompatible assignments. Give an example where the search space is smaller.

Exercise 6.9 After completing the previous exercise, design incompatible assignments for the blocks world. (This should result in dramatic search improvements.)

6.3.1 Defining Heuristics for a Regression Planner

The regression planner can use the same heuristic function as the forward planner. However, just because a heuristic is useful for a forward planner does not mean it is useful for a regression planner, and vice versa. you should experiment with whether the same heuristic works well for both a a regression planner and a forward planner.

The following runs the same example as the forward planner with and without the heuristic defined for the forward planner:

_____ stripsHeuristic.py — (continued) _____

```
64  ##### Regression Planner
65  from stripsRegressionPlanner import Regression_STRIPS
66
67  def test_regression_heuristic(thisproblem=problem1):
68      print("\n***** REGRESSION NO HEURISTIC")
69      print(SearcherMPP(Regression_STRIPS(thisproblem)).search())
70
71      print("\n***** REGRESSION WITH HEURISTICs h1 and h2")
72      print(SearcherMPP(Regression_STRIPS(thisproblem,maxh(h1,h2))).search())
73
74  if __name__ == "__main__":
75      test_regression_heuristic()
```

Exercise 6.10 Try the regression planner with a heuristic function of just $h1$ and with just $h2$ (defined in Section 6.2.1). Explain how each one prunes or doesn't prune the search space.

Exercise 6.11 Create a better heuristic than *heuristic_fun* defined in Section 6.2.1.

6.4   Planning as a CSP

> To run the demo, in folder "aipython", load "stripsCSPPlanner.py", and copy and paste the commented-out example queries at the bottom of that file. This assumes Python 3.

Here we implement the CSP planner assuming there is a single action at each step. This creates a CSP that can use any of the CSP algorithms to solve (e.g., stochastic local search or arc consistency with domain splitting).

This assumes the same action representation as before; we do not consider factored actions (action features), nor do we implement state constraints.

```
                    stripsCSPPlanner.py — CSP planner where actions are represented using STRIPS
11  from cspProblem import CSP, Constraint
12
13  class CSP_from_STRIPS(CSP):
14      """A CSP where:
15      * a CSP variable is constructed by st(var,stage).
16      * the dynamics are specified by the STRIPS representation of actions
17      """
18
19      def __init__(self, planning_problem, number_stages=2):
20          prob_domain = planning_problem.prob_domain
21          initial_state = planning_problem.initial_state
22          goal = planning_problem.goal
23          self.act_vars = [st('action',stage) for stage in range(number_stages)]
24          domains = {av:prob_domain.actions for av in self.act_vars}
25          domains.update({ st(var,stage):dom
26                          for (var,dom) in prob_domain.feats_vals.items()
27                          for stage in range(number_stages+1)})
28          # intial state constraints:
29          constraints = [Constraint((st(var,0),), is_(val))
30                          for (var,val) in initial_state.items()]
31          # goal constraints on the final state:
32          constraints += [Constraint((st(var,number_stages),),
33                                      is_(val))
34                          for (var,val) in goal.items()]
35          # precondition constraints:
36          constraints += [Constraint((st(var,stage), st('action',stage)),
37                                  if_(val,act)) # st(var,stage)==val if st('action',stage)=act
38                          for act,strps in prob_domain.strips_map.items()
39                          for var,val in strps.preconditions.items()
40                          for stage in range(number_stages)]
41          # effect constraints:
42          constraints += [Constraint((st(var,stage+1), st('action',stage)),
43                                  if_(val,act)) # st(var,stage+1)==val if st('action',stage)==act
44                          for act,strps in prob_domain.strips_map.items()
45                          for var,val in strps.effects.items()
46                          for stage in range(number_stages)]
47          # frame constraints:
48          constraints += [Constraint((st(var,stage), st('action',stage), st(var,stage+1)),
49                                  eq_if_not_in_({act for act in prob_domain.actions
50                                              if var in prob_domain.strips_map[act].effects}))
51                          for var in prob_domain.feats_vals
52                          for stage in range(number_stages) ]
53          CSP.__init__(self, domains, constraints)
54
55      def extract_plan(self,soln):
56          return [soln[a] for a in self.act_vars]
57
58  def st(var,stage):
59      """returns a string for the var-stage pair that can be used as a variable"""
```

```
60 |     return str(var)+"_"+str(stage)
```

The following methods return methods which can be applied to the particular environment.

For example, *is_*(3) returns a function that when applied to 3, returns True and when aplied to any other value returns False. So *is_*(3)(3) trurns *True* and *is_*(3)(7) returns *False*.

Note that the underscore ('_') is part of the name; here we use it as the convention that it is a function that returns a function. This uses two different styles to define *is_* and *if_*; returning a function defined by *lambda* is equivaent to returning the embedded function, except that the embedded function has a name. The embedded function can also be given a docstring.

_____ stripsCSPPlanner.py — (continued) _____

```
62 | def is_(val):
63 |     """returns a function that is true when it is it applied to val.
64 |     """
65 |     return lambda x: x == val
66 |
67 | def if_(v1,v2):
68 |     """if the second argument is v2, the first argument must be v1"""
69 |     #return lambda x1,x2: x1==v1 if x2==v2 else True
70 |     def if_fun(x1,x2):
71 |         return x1==v1 if x2==v2 else True
72 |     if_fun.__doc__ = "if x2 is "+str(v2)+" then x1 is "+str(v1)
73 |     return if_fun
74 |
75 | def eq_if_not_in_(actset):
76 |     """first and third arguments are equal if action is not in actset"""
77 |     return lambda x1, a, x2: x1==x2 if a not in actset else True
```

Putting it together, this returns a list of actions that solves the problem *prob* for a given horizon. If you want to do more than just return the list of actions, you might want to get it to return the solution. Or even enumerate the solutions (by using *Search_with_AC_from_CSP*).

_____ stripsCSPPlanner.py — (continued) _____

```
79 | def con_plan(prob,horizon):
80 |     """finds a plan for problem prob given horizon.
81 |     """
82 |     csp = CSP_from_STRIPS(prob, horizon)
83 |     sol = Con_solver(csp).solve_one()
84 |     return csp.extract_plan(sol) if sol else sol
```

The following are some example queries.

_____ stripsCSPPlanner.py — (continued) _____

```
86 | from searchGeneric import Searcher
87 | from stripsProblem import delivery_domain
88 | from cspConsistency import Search_with_AC_from_CSP, Con_solver
89 | from stripsProblem import Planning_problem, problem0, problem1, problem2
```

```
90
91  # Problem 0
92  # con_plan(problem0,1) # should it succeed?
93  # con_plan(problem0,2) # should it succeed?
94  # con_plan(problem0,3) # should it succeed?
95  # To use search to enumerate solutions
96  #searcher0a = Searcher(Search_with_AC_from_CSP(CSP_from_STRIPS(problem0, 1)))
97  #print(searcher0a.search())
98
99  ## Problem 1
100 # con_plan(problem1,5) # should it succeed?
101 # con_plan(problem1,4) # should it succeed?
102 ## To use search to enumerate solutions:
103 #searcher15a = Searcher(Search_with_AC_from_CSP(CSP_from_STRIPS(problem1, 5)))
104 #print(searcher15a.search())
105
106 ## Problem 2
107 #con_plan(problem2, 6) # should fail??
108 #con_plan(problem2, 7) # should succeed???
109
110 ## Example 6.13
111 problem3 = Planning_problem(delivery_domain,
112                         {'SWC':True, 'RHC':False}, {'SWC':False})
113 #con_plan(problem3,2) # Horizon of 2
114 #con_plan(problem3,3) # Horizon of 3
115
116 problem4 = Planning_problem(delivery_domain,{'SWC':True},
117                         {'SWC':False, 'MW':False, 'RHM':False})
118
119 # For the stochastic local search:
120 #from cspSLS import SLSearcher, Runtime_distribution
121 # cspplanning15 = CSP_from_STRIPS(problem1, 5) # should succeed
122 #se0 = SLSearcher(cspplanning15); print(se0.search(100000,0.5))
123 #p = Runtime_distribution(cspplanning15)
124 #p.plot_run(1000,1000,0.7) # warning will take a few minutes
```

6.5   Partial-Order Planning

> To run the demo, in folder "aipython", load "stripsPOP.py", and copy
> and paste the commented-out example queries at the bottom of that
> file.

A partial order planner maintains a partial order of action instances. An
action instance consists of a name and an index. We need action instances
because the same action could be carried out at different times.

───────────── stripsPOP.py — Partial-order Planner using STRIPS representation ─────────────

```
11  from searchProblem import Arc, Search_problem
```

```
12  import random
13
14  class Action_instance(object):
15      next_index = 0
16      def __init__(self,action,index=None):
17          if index is None:
18              index = Action_instance.next_index
19              Action_instance.next_index += 1
20          self.action = action
21          self.index = index
22
23      def __str__(self):
24          return str(self.action)+"#"+str(self.index)
25
26      __repr__ = __str__ # __repr__ function is the same as the __str__ function
```

A node (as in the abstraction of search space) in a partial-order planner
consists of:

- *actions*: a set of action instances.

- *constraints*: a set of (a_1, a_2) pairs, where a_1 and a_2 are action instances,
  which represents that a_1 must come before a_2 in the partial order. There
  are a number of ways that this could be represented. Here we represent
  the set of pairs that are in transitive closure of the *before* relation. This lets
  us quickly determine whether some before relation is consistent with the
  current constraints.

- *agenda*: a list of (s, a) pairs, where s is a (var, val) pair and a is an action
  instance. This means that variable *var* must have value *val* before a can
  occur.

- *causal_links*: a set of $(a0, g, a1)$ triples, where a_1 and a_2 are action instances
  and g is a (var, val) pair. This holds when action a_0 makes g true for action
  a_1.

_____ stripsPOP.py — (continued) _____

```
28  class POP_node(object):
29      """a (partial) partial-order plan. This is a node in the search space."""
30      def __init__(self, actions, constraints, agenda, causal_links):
31          """
32          * actions is a set of action instances
33          * constraints a set of (a0,a1) pairs, representing a0<a1,
34            closed under transitivity
35          * agenda list of (subgoal,action) pairs to be achieved, where
36            subgoal is a (variable,value) pair
37          * causal_links is a set of (a0,g,a1) triples,
38            where ai are action instances, and g is a (variable,value) pair
39          """
40          self.actions = actions # a set of action instances
```

```
41        self.constraints = constraints # a set of (a0,a1) pairs
42        self.agenda = agenda # list of (subgoal,action) pairs to be achieved
43        self.causal_links = causal_links # set of (a0,g,a1) triples
44
45    def __str__(self):
46        return ("actions: "+str({str(a) for a in self.actions})+
47                "\nconstraints: "+
48                str({(str(a1),str(a2)) for (a1,a2) in self.constraints})+
49                "\nagenda: "+
50                str([(str(s),str(a)) for (s,a) in self.agenda])+
51                "\ncausal_links:"+
52                str({(str(a0),str(g),str(a2)) for (a0,g,a2) in self.causal_links})
)
```

extract_plan constructs a total order of action instances that is consistent with the partial order.

_____ stripsPOP.py — (continued) _____

```
54    def extract_plan(self):
55        """returns a total ordering of the action instances consistent
56        with the constraints.
57        raises IndexError if there is no choice.
58        """
59        sorted_acts = []
60        other_acts = set(self.actions)
61        while other_acts:
62            a = random.choice([a for a in other_acts if
63                    all(((a1,a) not in self.constraints) for a1 in other_acts)])
64            sorted_acts.append(a)
65            other_acts.remove(a)
66        return sorted_acts
```

POP_search_from_STRIPS is an instance of a search problem. As such, we need to define the start nodes, the goal, and the neighbors of a node.

_____ stripsPOP.py — (continued) _____

```
68  from utilities import Displayable
69
70  class POP_search_from_STRIPS(Search_problem, Displayable):
71      def __init__(self,planning_problem):
72          Search_problem.__init__(self)
73          self.planning_problem = planning_problem
74          self.start = Action_instance("start")
75          self.finish = Action_instance("finish")
76
77      def is_goal(self, node):
78          return node.agenda == []
79
80      def start_node(self):
81          constraints = {(self.start, self.finish)}
82          agenda = [(g, self.finish) for g in self.planning_problem.goal.items()]
83          return POP_node([self.start,self.finish], constraints, agenda, [] )
```

The *neighbors* method is a coroutine that enumerates the neighbors of a given node.

```
_____ stripsPOP.py — (continued) _____
```
```python
85      def neighbors(self, node):
86          """enumerates the neighbors of node"""
87          self.display(3,"finding neighbors of\n",node)
88          if node.agenda:
89              subgoal,act1 = node.agenda[0]
90              self.display(2,"selecting",subgoal,"for",act1)
91              new_agenda = node.agenda[1:]
92              for act0 in node.actions:
93                  if (self.achieves(act0, subgoal) and
94                      self.possible((act0,act1),node.constraints)):
95                      self.display(2," reusing",act0)
96                      consts1 = self.add_constraint((act0,act1),node.constraints)
97                      new_clink = (act0,subgoal,act1)
98                      new_cls = node.causal_links + [new_clink]
99                      for consts2 in self.protect_cl_for_actions(node.actions,consts1,new_clink):
100                         yield Arc(node,
101                                 POP_node(node.actions,consts2,new_agenda,new_cls),
102                                 cost=0)
103             for a0 in self.planning_problem.prob_domain.strips_map: #a0 is an action
104                 if self.achieves(a0, subgoal):
105                     #a0 acheieves subgoal
106                     new_a = Action_instance(a0)
107                     self.display(2," using new action",new_a)
108                     new_actions = node.actions + [new_a]
109                     consts1 = self.add_constraint((self.start,new_a),node.constraints)
110                     consts2 = self.add_constraint((new_a,act1),consts1)
111                     preconds = self.planning_problem.prob_domain.strips_map[a0].preconditions
112                     new_agenda = new_agenda + [(pre,new_a) for pre in preconds.items()]
113                     new_clink = (new_a,subgoal,act1)
114                     new_cls = node.causal_links + [new_clink]
115                     for consts3 in self.protect_all_cls(node.causal_links,new_a,consts2):
116                         for consts4 in self.protect_cl_for_actions(node.actions,consts3,new_clink):
117                             yield Arc(node,
118                                     POP_node(new_actions,consts4,new_agenda,new_cls),
119                                     cost=1)
```

Given a casual link (*a*0, *subgoal*, *a*1), the following method protects the causal link from each action in *actions*. Whenever an action deletes *subgoal*, the action needs to be before *a*0 or after *a*1. This method enumerates all constraints that result from protecting the causal link from all actions.

```
_____ stripsPOP.py — (continued) _____
```
```python
121     def protect_cl_for_actions(self, actions, constrs, clink):
122         """yields constriants that extend constrs and
123         protect causal link (a0, subgoal, a1)
124         for each action in actions
125         """
```

```
126        if actions:
127            a = actions[0]
128            rem_actions = actions[1:]
129            a0, subgoal, a1 = clink
130            if a != a0 and a != a1 and self.deletes(a,subgoal):
131                if self.possible((a,a0),constrs):
132                    new_const = self.add_constraint((a,a0),constrs)
133                    for e in self.protect_cl_for_actions(rem_actions,new_const,clink): yield e
# could be "yield from"
134                if self.possible((a1,a),constrs):
135                    new_const = self.add_constraint((a1,a),constrs)
136                    for e in self.protect_cl_for_actions(rem_actions,new_const,clink): yield e
137            else:
138                for e in self.protect_cl_for_actions(rem_actions,constrs,clink): yield e
139        else:
140            yield constrs
```

Given an action *act*, the following method protects all the causal links in *clinks* from *act*. Whenever *act* deletes *subgoal* from some causal link $(a0, subgoal, a1)$, the action *act* needs to be before *a0* or after *a1*. This method enumerates all constraints that result from protecting the causal links from *act*.

_____stripsPOP.py — (continued) _____

```
142    def protect_all_cls(self, clinks, act, constrs):
143        """yields constraints that protect all causal links from act"""
144        if clinks:
145            (a0,cond,a1) = clinks[0] # select a causal link
146            rem_clinks = clinks[1:] # remaining causal links
147            if act != a0 and act != a1 and self.deletes(act,cond):
148                if self.possible((act,a0),constrs):
149                    new_const = self.add_constraint((act,a0),constrs)
150                    for e in self.protect_all_cls(rem_clinks,act,new_const): yield e
151                if self.possible((a1,act),constrs):
152                    new_const = self.add_constraint((a1,act),constrs)
153                    for e in self.protect_all_cls(rem_clinks,act,new_const): yield e
154            else:
155                for e in self.protect_all_cls(rem_clinks,act,constrs): yield e
156        else:
157            yield constrs
```

The following methods check whether an action (or action instance) achives or deletes some subgoal.

_____stripsPOP.py — (continued) _____

```
159    def achieves(self,action,subgoal):
160        var,val = subgoal
161        return var in self.effects(action) and self.effects(action)[var] == val
162
163    def deletes(self,action,subgoal):
164        var,val = subgoal
165        return var in self.effects(action) and self.effects(action)[var] != val
```

```
166
167      def effects(self,action):
168          """returns the variable:value dictionary of the effects of action.
169          works for both actions and action instances"""
170          if isinstance(action, Action_instance):
171              action = action.action
172          if action == "start":
173              return self.planning_problem.initial_state
174          elif action == "finish":
175              return {}
176          else:
177              return self.planning_problem.prob_domain.strips_map[action].effects
```

The constriants are represented as a set of pairs closed under transitivity.
Thus if (a, b) and (b, c) are the list, then (a, c) must also be in the list. This means
that adding a new constraint means adding the implied pairs, but querying
whether some order is consistent is quick.

_____ stripsPOP.py — (continued) _____
```
179      def add_constraint(self, pair, const):
180          if pair in const:
181              return const
182          todo = [pair]
183          newconst = const.copy()
184          while todo:
185              x0,x1 = todo.pop()
186              newconst.add((x0,x1))
187              for x,y in newconst:
188                  if x==x1 and (x0,y) not in newconst:
189                      todo.append((x0,y))
190                  if y==x0 and (x,x1) not in newconst:
191                      todo.append((x,x1))
192          return newconst
193
194      def possible(self,pair,constraint):
195          (x,y) = pair
196          return (y,x) not in constraint
```

Some code for testing:

_____ stripsPOP.py — (continued) _____
```
198  from searchBranchAndBound import DF_branch_and_bound
199  from searchGeneric import AStarSearcher
200  from searchMPP import SearcherMPP
201  from stripsProblem import problem0, problem1, problem2
202
203  rplanning0 = POP_search_from_STRIPS(problem0)
204  rplanning1 = POP_search_from_STRIPS(problem1)
205  rplanning2 = POP_search_from_STRIPS(problem2)
206  searcher0 = DF_branch_and_bound(rplanning0,5)
207  searcher0a = AStarSearcher(rplanning0)
```

```
208  searcher1 = DF_branch_and_bound(rplanning1,10)
209  searcher1a = AStarSearcher(rplanning1)
210  searcher2 = DF_branch_and_bound(rplanning2,10)
211  searcher2a = AStarSearcher(rplanning2)
212  # Try one of the following searchers
213  # a = searcher0.search()
214  # a = searcher0a.search()
215  # a.end().extract_plan() # print a plan found
216  # a.end().constraints   # print the constraints
217  # AStarSearcher.max_display_level = 0 # less detailed display
218  # DF_branch_and_bound.max_display_level = 0 # less detailed display
219  # a = searcher1.search()
220  # a = searcher1a.search()
221  # a = searcher2.search()
222  # a = searcher2a.search()
```

Chapter 7

Supervised Machine Learning

A good source of datasets is the UCI machine Learning Repository [**?**]; the SPECT and car datasets are from this repository.

7.1 Representations of Data and Predictions

A **data set** is an enumeration of examples. Each **example** is a list (or tuple) of feature values. The feature values can be numbers or strings. A **feature** is a function from the examples into the range of the feature. We assume each feature has a variable `frange` that gives the range of the feature. A **Boolean feature** is a function from the examples into {*False, True*}. So $f(e)$ is either *True* or *False*, where f is a feature and e is an example. The `__doc__` variable contains the docstring, a string description of the function.

_____learnProblem.py — A Learning Problem_____
```
11  import math, random
12  import csv
13  from utilities import Displayable
14
15  boolean = [False, True]
```

When creating a data set, we partition the data into a training set (*train*) and a test set (*test*). The target feature is the feature that we are making a prediction of.

_____learnProblem.py — (continued)_____
```
17  class Data_set(Displayable):
18      """ A data set consists of a list of training data and a list of test data.
19      """
20      seed = None #123456 # make it None for a different test set each time
21
```

```
22    def __init__(self, train, test=None, prob_test=0.30, target_index=0, header=None):
23        """A dataset for learning.
24        train is a list of tuples representing the training examples
25        test is the list of tuples representing the test examples
26        if test is None, a test set is created by selecting each
27            example with probability prob_test
28        target_index is the index of the target. If negative, it counts from right.
29            If target_index is larger than the number of properties,
30            there is no target (for unsupervised learning)
31        header is a list of names for the features
32        """
33        if test is None:
34            train,test = partition_data(train, prob_test, seed=self.seed)
35        self.train = train
36        self.test = test
37        self.display(1,"Tuples read. \nTraining set", len(train),
38                    "examples. Number of columns:",{len(e) for e in train},
39                    "\nTest set", len(test),
40                    "examples. Number of columns:",{len(e) for e in test}
41                    )
42        self.prob_test = prob_test
43        self.num_properties = len(self.train[0])
44        if target_index < 0: #allows for -1, -2, etc.
45            target_index = self.num_properties + target_index
46        self.target_index = target_index
47        self.header = header
48        self.create_features()
49        self.display(1,"There are",len(self.input_features),"input features")
```

Initially we assume that all of the properties can be mapped directly into features. If all values are 0 or 1 they can be used as Boolean features. This will be overridden to allow for more general features.

―――――――――――――――――――― learnProblem.py — (continued) ――――――――――――――――――――

```
51    def create_features(self):
52        """create the input features and target feature.
53        This assumes that the features all have domain {0,1}.
54        This should be overridden if the features have a different domain.
55        """
56        self.input_features = []
57        for i in range(self.num_properties):
58            def feat(e,index=i):
59                return e[index]
60            if self.header:
61                feat.__doc__ = self.header[i]
62            else:
63                feat.__doc__ = "e["+str(i)+"]"
64            feat.frange = [0,1]
65            if i == self.target_index:
66                self.target = feat
67            else:
```

```
68                    self.input_features.append(feat)
```

7.1.1 Evaluating Predictions

A **predictor** is a function that takes an example and makes a prediction on the value of the target feature. A predictor can be judged according to a number of evaluation criteria. The function *evaluate_dataset* returns the average error for each example, where the error for each example depends on the evaluation criteria. Here we consider three evaluation criteria, the sum-of-squares, the sum of absolute errors and the logloss (the negative log-likelihood, which is the number of bits to describe the data using a code based on the prediction treated as a probability).

```
———————————learnProblem.py — (continued) ———————————
70    evaluation_criteria = ["sum-of-squares","sum_absolute","logloss"]
71
72    def evaluate_dataset(self, data, predictor, evaluation_criterion):
73        """Evaluates predictor on data according to the evaluation_criterion.
74        predictor is a function that takes an example and returns a
75                prediction for the target feature.
76        evaluation_criterion is one of the evaluation_criteria.
77        """
78        assert evaluation_criterion in self.evaluation_criteria,"given: "+str(evaluation_criterion)
79        if data:
80            try:
81                error = sum(error_example(predictor(example), self.target(example),
82                                          evaluation_criterion)
83                        for example in data)/len(data)
84            except ValueError:
85                return float("inf") # infinity
86            return error
```

error_example is used to evaluate a single example, based on the predicted value, the actual value and the evaluation criterion. Note that for logloss, the actual value must be 0 or 1.

```
———————————learnProblem.py — (continued) ———————————
88    def error_example(predicted, actual, evaluation_criterion):
89        """returns the error of the for the predicted value given the actual value
90        according to evaluation_criterion.
91        Throws ValueError if the error is infinite (log(0))
92        """
93        if evaluation_criterion=="sum-of-squares":
94            return (predicted-actual)**2
95        elif evaluation_criterion=="sum_absolute":
96            return abs(predicted-actual)
97        elif evaluation_criterion=="logloss":
98            assert actual in [0,1], "actual="+str(actual)
99            if actual==0:
```

```
100            return -math.log2(1-predicted)
101        else:
102            return -math.log2(predicted)
103    elif evaluation_criterion=="characteristic_ss":
104        return sum((1-predicted[i])**2 if actual==i else predicted[i]**2
105                   for i in range(len(predicted)))
106    else:
107        raise RuntimeError("Not evaluation criteria: "+str(evaluation_criterion))
```

7.1.2   Creating Test and Training Sets

The following method partitions the data into a training set and a test set. Note that this does not guarantee that the test set will contain exactly a proportion of the data equal to *prob_test*.

[An alternative is to use *random.sample*() which can guarantee that the test set will contain exactly a particular proportion of the data. However this would require knowing how many elements are in the data set, which we may not know, as *data* may just be a generator of the data (e.g., when reading the data from a file).]

_____ learnProblem.py — (continued) _____

```
109  def partition_data(data, prob_test=0.30, seed=None):
110      """partitions the data into a training set and a test set, where
111      prob_test is the probability of each example being in the test set.
112      """
113      train = []
114      test = []
115      if seed:     # given seed makes the partition consistent from run-to-run
116          random.seed(seed)
117      for example in data:
118          if random.random() < prob_test:
119              test.append(example)
120          else:
121              train.append(example)
122      return train, test
```

7.1.3   Importing Data From File

A data set is typically loaded from a file. The default here is that it loaded from a CSV (comma separated values) file, although the default separator can be changed. This assumes that all lines that contain the separator are valid data (so we only include those data items that contain more than one element). This allows for blank lines and comment lines that do not contain the separator. However, it means that this method is not suitable for cases where there is only one feature.

Note that *data_all* and *data_tuples* are generators. *data_all* is a generator of a list of list of strings. This version assumes that CSV files are simple. The

standard *csv* package, that allows quoted arguments, can be used by uncommenting the line for *data_aa* and commenting out the following line. *data_tuples* contains only those lines that contain the delimiter (others lines are assumed to be empty or comments), and tries to convert the elements to numbers whenever possible.

This allows for some of the columns to be included. Note that if include_only is specified, the target index is in the resulting

```
_____learnProblem.py — (continued) _____
124 │ class Data_from_file(Data_set):
125 │     def __init__(self, file_name, separator=',', num_train=None, prob_test=0.3,
126 │                  has_header=False, target_index=0, boolean_features=True,
127 │                  categorical=[], include_only=None):
128 │         """create a dataset from a file
129 │         separator is the character that separates the attributes
130 │         num_train is a number n specifying the first n tuples are training, or None
131 │         prob_test is the probability an example should in the test set (if num_train is None)
132 │         has_header is True if the first line of file is a header
133 │         target_index specifies which feature is the target
134 │         boolean_features specifies whether we want to create Boolean features
135 │             (if False, is uses the original features).
136 │         categorical is a set (or list) of features that should be treated as categorical
137 │         include_only is a list or set of indexes of columns to include
138 │         """
139 │         self.boolean_features = boolean_features
140 │         with open(file_name,'r',newline='') as csvfile:
141 │             # data_all = csv.reader(csvfile,delimiter=separator) # for more complicted CSV files
142 │             data_all = (line.strip().split(separator) for line in csvfile)
143 │             if include_only is not None:
144 │                 data_all = ([v for (i,v) in enumerate(line) if i in include_only] for line in data_
145 │             if has_header:
146 │                 header = next(data_all)
147 │             else:
148 │                 header = None
149 │             data_tuples = (make_num(d) for d in data_all if len(d)>1)
150 │             if num_train is not None:
151 │                 # training set is divided into training then text examples
152 │                 # the file is only read once, and the data is placed in appropriate list
153 │                 train = []
154 │                 for i in range(num_train):  # will give an error if insufficient examples
155 │                     train.append(next(data_tuples))
156 │                 test = list(data_tuples)
157 │                 Data_set.__init__(self,train, test=test, target_index=target_index,header=header)
158 │             else:    # randomly assign training and test examples
159 │                 Data_set.__init__(self,data_tuples, prob_test=prob_test,
160 │                                   target_index=target_index, header=header)
161 │
162 │     def __str__(self):
163 │         if self.train and len(self.train)>0:
164 │             return ("Data: "+str(len(self.train))+" training examples, "
```

```
165                     +str(len(self.test))+" test examples, "
166                     +str(len(self.train[0]))+" features.")
167           else:
168               return ("Data: "+str(len(self.train))+" training examples, "
169                     +str(len(self.test))+" test examples.")
```

7.1.4 Creating Binary Features

Some of the algorithms require Boolean features or features with domain $\{0, 1\}$. In order to be able to use these on datasets that allow for arbitrary ranges of input variables, we construct binary features from the attributes. This method overrides the method in *Data_set*.

There are 3 cases:

- When the attribute only has two values, we designate one to be the "true" value.

- When the values are all numeric, we assume they are ordered (as opposed to just being some classes that happen to be labelled with numbers, but where the numbers have no meaning) and construct Boolean features for splits of the data. That is, the feature is $e[ind] < cut$ for some value *cut*. We choose a number of *cut* values, up to a maximum number of cuts, given by *max_num_cuts*.

- When the values are not all numeric, we assume they are unordered, and create an indicator function for each value. An indicator function for a value returns true when that value is given and false otherwise. Note that we can't create an indicator function for values that appear in the test set but not in the training set because we haven't seen the test set. For the examples in the test set with that value, the indicator functions return false.

_____ learnProblem.py — (continued) _____

```
171      def create_features(self, max_num_cuts=8):
172          """creates boolean features from input features.
173          max_num_cuts is the maximum number of binary variables
174             to split a numerical feature into.
175          """
176          ranges = [set() for i in range(self.num_properties)]
177          for example in self.train:
178              for ind,val in enumerate(example):
179                  ranges[ind].add(val)
180          if self.target_index <= self.num_properties:
181              def target(e,index=self.target_index):
182                  return e[index]
183              if self.header:
184                  target.__doc__ = self.header[ind]
```

```
185                    else:
186                        target.__doc__ = "e["+str(ind)+"]"
187                target.frange = ranges[self.target_index]
188                self.target = target
189        if self.boolean_features:
190            self.input_features = []
191            for ind,frange in enumerate(ranges):
192                if ind != self.target_index and len(frange)>1:
193                    if len(frange) == 2:
194                        # two values, the feature is equality to one of them.
195                        true_val = list(frange)[1] # choose one as true
196                        def feat(e, i=ind, tv=true_val):
197                            return e[i]==tv
198                        if self.header:
199                            feat.__doc__ = self.header[ind]+"=="+str(true_val)
200                        else:
201                            feat.__doc__ = "e["+str(ind)+"]=="+str(true_val)
202                        feat.frange = boolean
203                        self.input_features.append(feat)
204                    elif all(isinstance(val,(int,float)) for val in frange):
205                        # all numeric, create cuts of the data
206                        sorted_frange = sorted(frange)
207                        num_cuts = min(max_num_cuts,len(frange))
208                        cut_positions = [len(frange)*i//num_cuts for i in range(1,num_cuts)]
209                        for cut in cut_positions:
210                            cutat = sorted_frange[cut]
211                            def feat(e, ind_=ind, cutat=cutat):
212                                return e[ind_] < cutat
213
214                            if self.header:
215                                feat.__doc__ = self.header[ind]+"<"+str(cutat)
216                            else:
217                                feat.__doc__ = "e["+str(ind)+"]<"+str(cutat)
218                            feat.frange = boolean
219                            self.input_features.append(feat)
220                    else:
221                        # create an indicator function for every value
222                        for val in frange:
223                            def feat(e, ind_=ind, val_=val):
224                                return e[ind_] == val_
225                            if self.header:
226                                feat.__doc__ = self.header[ind]+"=="+str(val)
227                            else:
228                                feat.__doc__= "e["+str(ind)+"]=="+str(val)
229                            feat.frange = boolean
230                            self.input_features.append(feat)
231        else: # boolean_features is off
232            self.input_features = []
233            for i in range(self.num_properties):
234                def feat(e,index=i):
```

```
235                    return e[index]
236                 if self.header:
237                    feat.__doc__ = self.header[i]
238                 else:
239                     feat.__doc__ = "e["+str(i)+"]"
240                 feat.frange = ranges[i]
241                 if i == self.target_index:
242                    self.target = feat
243                 else:
244                    self.input_features.append(feat)
```

Exercise 7.1 Change the code so that it splits using $e[ind] \leq cut$ instead of $e[ind] <$ *cut*. Check boundary cases, such as 3 elements with 2 cuts, and if there are 30 elements (integers from 100 to 129), and you want 2 cuts, the resulting Boolean features should be $e[ind] \leq 109$ and $e[ind] \leq 119$ to make sure that each of the resulting ranges is equal size.

Exercise 7.2 This splits on whether the feature is less than one of the values in the training set. Sam suggested it might be better to split between the values in the training set, and suggested using

$$cutat = (sorted\_frange[cut] + sorted\_frange[cut-1])/2$$

Why might Sam have suggested this? Does this work better? (Try it on a few data sets).

When reading from a file all of the values are strings. This next method tries to convert each values into a number (an int or a float), if it is possible.

_____learnProblem.py — (continued) _____

```
245  def make_num(str_list):
246      """make the elements of string list str_list numerical if possible.
247      Otherwise remove initial and trailing spaces.
248      """
249      res = []
250      for e in str_list:
251          try:
252              res.append(int(e))
253          except ValueError:
254              try:
255                  res.append(float(e))
256              except ValueError:
257                  res.append(e.strip())
258      return res
```

7.1.5 Augmented Features

Sometimes we want to augment the features with new features computed from the old features (eg. the product of features). Here we allow the creation of a new dataset from an old dataset but with new features.

A feature is a function of examples. A unary feature constructor takes a feature and returns a new feature. A binary feature combiner takes two features and returns a new feature.

_____learnProblem.py — (continued)_____

```
260  class Data_set_augmented(Data_set):
261      def __init__(self, dataset, unary_functions=[], binary_functions=[], include_orig=True):
262          """creates a dataset like dataset but with new features
263          unary_function is a list of unary feature constructors
264          binary_functions is a list of binary feature combiners.
265          include_orig specifies whether the original features should be included
266          """
267          self.orig_dataset = dataset
268          self.unary_functions = unary_functions
269          self.binary_functions = binary_functions
270          self.include_orig = include_orig
271          self.target = dataset.target
272          Data_set.__init__(self,dataset.train, test=dataset.test,
273                          target_index = dataset.target_index)
274
275      def create_features(self):
276          if self.include_orig:
277              self.input_features = self.orig_dataset.input_features.copy()
278          else:
279              self.input_features = []
280          for u in self.unary_functions:
281              for f in self.orig_dataset.input_features:
282                  self.input_features.append(u(f))
283          for b in self.binary_functions:
284              for f1 in self.orig_dataset.input_features:
285                  for f2 in self.orig_dataset.input_features:
286                      if f1 != f2:
287                          self.input_features.append(b(f1,f2))
```

The following are useful unary feature constructors and binary feature combiner.

_____learnProblem.py — (continued)_____

```
289  def square(f):
290      """a unary feature constructor to construct the square of a feature
291      """
292      def sq(e):
293          return f(e)**2
294      sq.__doc__ = f.__doc__+"**2"
295      return sq
296
297  def power_feat(n):
298      """given n returns a unary feature constructor to construct the nth power of a feature.
299      e.g., power_feat(2) is the same as square
300      """
301      def fn(f,n=n):
```

```
302            def pow(e,n=n):
303                return f(e)**n
304            pow.__doc__ = f.__doc__+"**"+str(n)
305            return pow
306        return fn
307
308  def prod_feat(f1,f2):
309      """a new feature that is the product of features f1 and f2
310      """
311      def feat(e):
312          return f1(e)*f2(e)
313      feat.__doc__ = f1.__doc__+"*"+f2.__doc__
314      return feat
315
316  def eq_feat(f1,f2):
317      """a new feature that is 1 if f1 and f2 give same value
318      """
319      def feat(e):
320          return 1 if f1(e)==f2(e) else 0
321      feat.__doc__ = f1.__doc__+"=="+f2.__doc__
322      return feat
323
324  def xor_feat(f1,f2):
325      """a new feature that is 1 if f1 and f2 give different values
326      """
327      def feat(e):
328          return 1 if f1(e)!=f2(e) else 0
329      feat.__doc__ = f1.__doc__+"!="+f2.__doc__
330      return feat
```

Example:

_____ learnProblem.py — (continued) _____

```
332  # from learnProblem import Data_set_augmented,prod_feat
333  # data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
334  ## data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
335  # dataplus = Data_set_augmented(data,[],[prod_feat])
336  # dataplus = Data_set_augmented(data,[],[prod_feat,xor_feat])
```

Exercise 7.3 For symmetric properties, such as product, we don't need both $f1 * f2$ as well as $f2 * f1$ as extra properties. Allow the user to be able to declare feature constructors as symmetric (by associating a Boolean feature with them). Change *construct_features* so that it does not create both versions for symmetric combiners.

7.1.6   Learner

A learner takes a dataset (and possible other arguments specific to the method). To get it to learn, we call the *learn*() method. This implements *Displayable* so

that we can display traces at multiple levels of detail (and perhaps with a GUI).

_____learnProblem.py — (continued) _____

```
337  from utilities import Displayable
338
339  class Learner(Displayable):
340      def __init__(self, dataset):
341          raise NotImplementedError("Learner.__init__") # abstract method
342
343      def learn(self):
344          """returns a predictor, a function from a tuple to a value for the target feature
345          """
346          raise NotImplementedError("learn") # abstract method
```

7.2 Learning With No Input Features

If we make the same prediction for each example, what prediction should we make?

There are a few alternatives as to what could be allowed in a prediction:

- a point prediction, where we are only allowed to predict one of the values of the feature. For example, if the values of the feature are $\{0, 1\}$ we are only allowed to predict 0 or 1 or of the values are ratings in $\{1, 2, 3, 4, 5\}$, we can only predict one of these integers.

- a point prediction, where we are allowed to predict any value. For example, if the values of the feature are $\{0, 1\}$ we may be allowed to predict 0.3, 1, or even 1.7. For all of the criteria we can imagine, there is no point in predicting a value greater than 1 or less that zero (but that doesn't mean we can't), but it is often useful to predict a value between 0 and 1. If the values are ratings in $\{1, 2, 3, 4, 5\}$, we may want to predict 3.4.

- a probability distribution over the values of the feature. For each value v, we predict a non-negative number p_v, such that the sum over all predictions is 1.

The following code assumes the second of these, where we can make a point prediction of any value (although median will only predict one of the actual values for the feature).

The *point_prediction* function takes in a target feature (which is assumed to be numeric), some training data, and a section of what to return, and returns a function that takes in an example, and makes a prediction of a value for the target variable, but makes same prediction for all examples.

This method uses *selection*, whose value should be "median", "proportion", or "Laplace" determine what prediction should be made.

```python
from learnProblem import Learner, Data_set
import math, random

selections = ["median", "mean", "Laplace"]

def point_prediction(target, training_data,
                     selection="mean" ):
    """makes a point prediction for a set of training data.
    target provides the target
    training_data provides the training data to use (often a subset of train).
    selection specifies what statistic of the data to use as the evaluation.
    to_optimize provides a criteria to optimize (used to guess selection)
    """
    assert len(training_data)>0
    if selection == "median":
        counts,total = target_counts(target,training_data)
        middle = total/2
        cumulative = 0
        for val,num in sorted(counts.items()):
            cumulative += num
            if cumulative > middle:
                break # exit loop with val as the median
    elif selection == "mean":
        val = mean((target(e) for e in training_data))
    elif selection == "Laplace":
        val = mean((target(e) for e in training_data),len(target.frange),1)
    elif selection == "mode":
        raise NotImplementedError("mode")
    else:
        raise RuntimeError("Not valid selection: "+str(selection))
    fun = lambda x: val
    fun.__doc__ = str(val)
    return fun

def mean(enum,count=0,sum=0):
    """returns the mean of enumeration enum,
       count and sum are initial counts and the initial sum.
       This works for enumerations, even where len() is not defined"""
    for e in enum:
        count += 1
        sum += e
    return sum/count

def target_counts(target, data_subset):
    """returns a value:count dictionary of the count of the number of
    times target has this value in data_subset, and the number of examples.
    """
    counts = {val:0 for val in target.frange}
    total = 0
```

```
60      for instance in data_subset:
61          total += 1
62          counts[target(instance)] += 1
63      return counts, total
```

7.2.1 Testing

To test the point prediction, we will first generate some data from a simple (Bernoulli) distribution, where there are two possible values, 0 and 1 for the target feature. Given *prob*, a number in the range $[0,1]$, this generate some training and test data where *prob* is the probability of each example being 1.

———————— learnNoInputs.py — (continued) ————————

```
65  class Data_set_random(Data_set):
66      """A data set of a {0,1} feature generated randomly given a probability"""
67      def __init__(self, prob, train_size, test_size=100):
68          """a data set of with train_size training examples,
69          test_size test examples
70          where each examples in generated where prob i the probability of 1
71          """
72          self.max_display_level = 0
73          train = [[1] if random.random()<prob else [0] for i in range(train_size)]
74          test = [[1] if random.random()<prob else [0] for i in range(test_size)]
75          Data_set.__init__(self, train, test, target_index=0)
```

Let's try to evaluate the predictions of the possible selections according to the different evaluation criteria, for various training sizes.

———————— learnNoInputs.py — (continued) ————————

```
77  def test_no_inputs():
78      num_samples = 1000 #number of runs to average over
79      test_size = 100   # number of test examples for each prediction
80      for train_size in [1,2,3,4,5,10,20,100,1000]:
81          total_error = {(select,crit):0
82                          for select in selections
83                          for crit in Data_set.evaluation_criteria}
84          for sample in range(num_samples): # average over num_samples
85              p = random.random()
86              data = Data_set_random(p, train_size, test_size)
87              for select in selections:
88                  prediction = point_prediction(data.target, data.train, selection=select)
89                  for ecrit in Data_set.evaluation_criteria:
90                      test_error = data.evaluate_dataset(data.test,prediction,ecrit)
91                      total_error[(select,ecrit)] += test_error
92          print("For training size",train_size,":")
93          for ecrit in Data_set.evaluation_criteria:
94              print("   Evaluated according to",ecrit,":")
95              for select in selections:
96                  print("       Average error of",select,"is",
97                      total_error[(select,ecrit)]/num_samples)
```

```
98
99  if __name__ == "__main__":
100     test_no_inputs()
```

7.3   Decision Tree Learning

> To run the decision tree learning demo, in folder "aipython", load
> "learnDT.py", using e.g., `ipython -i learnDT.py`, and it prints some
> test results. To try more examples, copy and paste the commented-
> out commands at the bottom of that file. This requires Python 3 with
> `matplotlib`.

The decision tree algorithm does binary splits, and assumes that all input
features are binary functions of the examples. It stops splitting if there are
no input features, the number of examples is less than a specified number of
examples or all of the examples agree on the target feature.

_____ learnDT.py — Learning a binary decision tree _____

```
11  from learnProblem import Learner, error_example
12  from learnNoInputs import point_prediction, target_counts, selections
13  import math
14
15  class DT_learner(Learner):
16      def __init__(self,
17                   dataset,
18                   to_optimize="sum-of-squares",
19                   leaf_selection="mean", # what to use for point prediction at leaves
20                   train=None,            # used for cross validation
21                   min_number_examples=10):
22          self.dataset = dataset
23          self.target = dataset.target
24          self.to_optimize = to_optimize
25          self.leaf_selection = leaf_selection
26          self.min_number_examples = min_number_examples
27          if train is None:
28              self.train = self.dataset.train
29          else:
30              self.train = train
31
32      def learn(self):
33          return self.learn_tree(self.dataset.input_features, self.train)
```

The main recursive algorithm, takes in a set of input features and a set of
training data. It first decides whether to split. If it doesn't split, it makes a point
prediction, ignoring the input features.

It doesn't split if there are no more input features, if there are fewer exam-
ples than *min_number_examples*, if all the examples agree on the value of the
target or if the best split makes all examples in the same partition

If it decides to split, it selects the best split and returns the condition to split on (in the variable *split*) and the corresponding partition of the examples.

```
_____learnDT.py — (continued)_____
35   def learn_tree(self, input_features, data_subset):
36       """returns a decision tree
37       for input_features is a set of possible conditions
38       data_subset is a subset of the data used to build this (sub)tree
39
40       where a decision tree is a function that takes an example and
41       makes a prediction on the target feature
42       """
43       if (input_features and len(data_subset) >= self.min_number_examples):
44           first_target_val = self.target(data_subset[0])
45           allagree = all(self.target(inst)==first_target_val for inst in data_subset)
46           if not allagree:
47               split, partn = self.select_split(input_features, data_subset)
48               if split: # the split succeeded in splitting the data
49                   false_examples, true_examples = partn
50                   rem_features = [fe for fe in input_features if fe != split]
51                   self.display(2,"Splitting on",split.__doc__,"with examples split",
52                               len(true_examples),":",len(false_examples))
53                   true_tree = self.learn_tree(rem_features,true_examples)
54                   false_tree = self.learn_tree(rem_features,false_examples)
55                   def fun(e):
56                       if split(e):
57                           return true_tree(e)
58                       else:
59                           return false_tree(e)
60                   #fun = lambda e: true_tree(e) if split(e) else false_tree(e)
61                   fun.__doc__ = ("if "+split.__doc__+" then ("+true_tree.__doc__+
62                                 ") else ("+false_tree.__doc__+")")
63                   return fun
64       # don't expand the trees but return a point prediction
65       return point_prediction(self.target, data_subset, selection=self.leaf_selection)
```

```
_____learnDT.py — (continued)_____
67   def select_split(self, input_features, data_subset):
68       """finds best feature to split on.
69
70       input_features is a non-empty list of features.
71       returns feature, partition
72       where feature is an input feature with the smallest error as
73             judged by to_optimize or
74             feature==None if there are no splits that improve the error
75       partition is a pair (false_examples, true_examples) if feature is not None
76       """
77       best_feat = None # best feature
78       # best_error = float("inf") # infinity - more than any error
79       best_error = training_error(self.dataset, data_subset, self.to_optimize)
```

```
80              best_partition = None
81           for feat in input_features:
82               false_examples, true_examples = partition(data_subset,feat)
83               if false_examples and true_examples: #both partitons are non-empty
84                   err = (training_error(self.dataset,false_examples,self.to_optimize)
85                          + training_error(self.dataset,true_examples,self.to_optimize))
86                   self.display(3," split on",feat.__doc__,"has err=",err,
87                           "splits into",len(true_examples),":",len(false_examples))
88                   if err < best_error:
89                       best_feat = feat
90                       best_error=err
91                       best_partition = false_examples, true_examples
92           self.display(3,"best split is on",best_feat.__doc__,
93                              "with err=",best_error)
94           return best_feat, best_partition
95
96  def partition(data_subset,feature):
97      """partitions the data_subset by the feature"""
98      true_examples = []
99      false_examples = []
100     for example in data_subset:
101         if feature(example):
102             true_examples.append(example)
103         else:
104             false_examples.append(example)
105     return false_examples, true_examples
106
107
108 def training_error(dataset, data_subset, to_optimize):
109     """returns training error for dataset on to_optimize.
110     This assumes that we choose the best value for the optimization
111     criteria for dataset according to point_prediction
112     """
113     select_dict = {"sum-of-squares":"mean", "sum_absolute":"median",
114                     "logloss":"Laplace"} # arbitrary mapping. Perhaps wrong.
115     selection = select_dict[to_optimize]
116     predictor = point_prediction(dataset.target, data_subset, selection=selection)
117     error = sum(error_example(predictor(example),
118                         dataset.target(example),
119                         to_optimize)
120             for example in data_subset)
121     return error
```

Test cases:

```
————————————————————————————learnDT.py — (continued) ————————————————————————————
123 from learnProblem import Data_set, Data_from_file
124
125 def test(data):
126     """Prints errors and the trees for various evaluation criteria and ways to select leaves.
127     """
```

```
128    for crit in Data_set.evaluation_criteria:
129        for leaf in selections:
130            tree = DT_learner(data, to_optimize=crit, leaf_selection=leaf).learn()
131            print("For",crit,"using",leaf,"at leaves, tree built is:",tree.__doc__)
132            if data.test:
133                for ecrit in Data_set.evaluation_criteria:
134                    test_error = data.evaluate_dataset(data.test, tree, ecrit)
135                    print("   Average error for", ecrit,"using",leaf, "at leaves is", test_error)
136
137 if __name__ == "__main__":
138    #print("carbool.csv"); test(data = Data_from_file('data/carbool.csv', target_index=-1))
139    # print("SPECT.csv"); test(data = Data_from_file('data/SPECT.csv', target_index=0))
140    print("mail_reading.csv"); test(data = Data_from_file('data/mail_reading.csv', target_index=-1
141    # print("holiday.csv"); test(data = Data_from_file('data/holiday.csv', num_train=19, target_in
```

Exercise 7.4 The current algorithm does not have a very sophisticated stopping criterion. What is the current stopping criterion? (Hint: you need to look at both *learn_tree* and *select_split*.)

Exercise 7.5 Extend the current algorithm to include in the stopping criterion

(a) A minimum child size; don't use a split if one of the children has fewer elements that this.

(b) A depth-bound on the depth of the tree.

(c) An improvement bound such that a split is only carried out if error with the split is better than the error without the split by at least the improvement bound.

Which values for these parameters make the prediction errors on the test set the smallest? Try it on more than one dataset.

Exercise 7.6 Without any input features, it is often better to include a pseudo-count that is added to the counts from the training data. Modify the code so that it includes a pseudo-count for the predictions. When evaluating a split, including pseudo counts can make the split worse than no split. Does pruning with an improvement bound and pseudo-counts make the algorithm work better than with an improvement bound by itself?

Exercise 7.7 Some people have suggested using information gain (which is equivalent to greedy optimization of logloss) as the measure of improvement when building the tree, even in they want to have non-probabilistic predictions in the final tree. Does this work better than myopically choosing the split that is best for the evaluation criteria we will use to judge the final prediction?

7.4   Cross Validation and Parameter Tuning

> To run the cross validation demo, in folder "aipython",
> load "learnCrossValidation.py", using e.g., ipython -i
> learnCrossValidation.py.   Run plot_fig_7_15() to produce a
> graph like Figure 7.15. Note that different runs will produce different
> graphs, so your graph will not look like the one in the textbook. To try
> more examples, copy and paste the commented-out commands at the
> bottom of that file. This requires Python 3 with matplotlib.

The above decision tree overfits the data. One way to determine whether
the prediction is overfitting is by cross validation. The code below implements
k-fold cross validation, which can be used to choose the value of parameters
to best fit the training data. If we want to use parameter tuning to improve
predictions on a particular data set, we can only use the training data (and not
the test data) to tune the parameter.

In *k*-fold cross validation, we partition the training set into *k* approximately
equal-sized folds (each fold is an enumeration of examples). For each fold, we
train on the other examples, and determine the error of the prediction on that
fold. For example, if there are 10 folds, we train on 90% of the data, and then
test on remaining 10% of the data. We do this 10 times, so that each example
gets used as a test set once, and in the training set 9 times.

The code below creates one copy of the data, and multiple views of the data.
For each fold, *fold* enumerates the examples in the fold, and *fold_complement*
enumerates the examples not in the fold.

───────────── learnCrossValidation.py — Cross Validation for Parameter Tuning ─────────────

```
11  from learnProblem import Data_set, Data_from_file, error_example
12  from learnDT import DT_learner
13  import matplotlib.pyplot as plt
14  import random
15
16  class K_fold_dataset(object):
17      def __init__(self, training_set, num_folds):
18          self.data = training_set.train.copy()
19          self.target = training_set.target
20          self.input_features = training_set.input_features
21          self.num_folds = num_folds
22          random.shuffle(self.data)
23          self.fold_boundaries = [(len(self.data)*i)//num_folds
24                                  for i in range(0,num_folds+1)]
25
26      def fold(self, fold_num):
27          for i in range(self.fold_boundaries[fold_num],
28                         self.fold_boundaries[fold_num+1]):
29              yield self.data[i]
30
```

```
31      def fold_complement(self, fold_num):
32          for i in range(0,self.fold_boundaries[fold_num]):
33              yield self.data[i]
34          for i in range(self.fold_boundaries[fold_num+1],len(self.data)):
35              yield self.data[i]
```

The validation error is the average error for each example, where we test on each fold, and learn on the other folds.

──────────────── learnCrossValidation.py — (continued) ────────────────

```
37      def validation_error(self, learner, criterion, **other_params):
38          error = 0
39          try:
40              for i in range(self.num_folds):
41                  predictor = learner(self, train=list(self.fold_complement(i)),
42                                      **other_params).learn()
43                  error += sum( error_example(predictor(example),
44                                              self.target(example),
45                                              criterion)
46                          for example in self.fold(i))
47          except ValueError:
48              return float("inf") #infinity
49          return error/len(self.data)
```

The *plot_error* method plots the average error as a function of a the minimun number of examples in decision-tree search, both for the validation set and for the test set. The error on the validation set can be used to tune the parameter — choose the value of the parameter that minimizes the error. The error on the test set cannot be used to tune the parameters; if is were to be used this way then it cannot be used to test.

──────────────── learnCrossValidation.py — (continued) ────────────────

```
51  def plot_error(data,criterion="sum-of-squares", num_folds=5, xscale='log'):
52      """Plots the error on the validation set and the test set
53      with respect to settings of the minimum number of examples.
54      xscale should be 'log' or 'linear'
55      """
56      plt.ion()
57      plt.xscale('linear') # change between log and linear scale
58      plt.xlabel("minimum number of examples")
59      plt.ylabel("average "+criterion+" error")
60      folded_data = K_fold_dataset(data, num_folds)
61      verrors = [] # validation errors
62      terrors = [] # test set errors
63      for mne in range(1,len(data.train)+2):
64          verrors.append(folded_data.validation_error(DT_learner,criterion,
65                                              min_number_examples=mne))
66          tree = DT_learner(data, criterion, min_number_examples=mne).learn()
67          terrors.append(data.evaluate_dataset(data.test,tree,criterion))
68      plt.plot(range(1,len(data.train)+2), verrors, ls='-',color='k', label="validation for "+criter
69      plt.plot(range(1,len(data.train)+2), terrors, ls='--',color='k', label="test set for "+criteri
```

```
70 │     plt.legend()
71 │     plt.draw()
72 │
73 │ # Try
74 │ # data = Data_from_file('data/mail_reading.csv', target_index=-1)
75 │ # data = Data_from_file('data/SPECT.csv',target_index=0)
76 │ # data = Data_from_file('data/carbool.csv', target_index=-1)
77 │ # plot_error(data)  # warning, may take a long time depending on the dataset
78 │
79 │ def plot_fig_7_15():  # different runs produce different plots
80 │     data = Data_from_file('data/SPECT.csv',target_index=0)
81 │     # data = Data_from_file('data/carbool.csv', target_index=-1)
82 │     plot_error(data)
83 │ # plot_fig_7_15() # warning takes a long time!
```

7.5 Linear Regression and Classification

Here we give a gradient descent searcher for linear regression and classification.

_____ learnLinear.py — Linear Regression and Classification _____
```
11 │ from learnProblem import Learner
12 │ import random, math
13 │
14 │ class Linear_learner(Learner):
15 │     def __init__(self, dataset, train=None,
16 │                 learning_rate=0.1, max_init = 0.2,
17 │                 squashed=True):
18 │         """Creates a gradient descent searcher for a linear classifier.
19 │         The main learning is carried out by learn()
20 │
21 │         dataset provides the target and the input features
22 │         train provides a subset of the training data to use
23 │         number_iterations is the default number of steps of gradient descent
24 │         learning_rate is the gradient descent step size
25 │         max_init is the maximum absolute value of the initial weights
26 │         squashed specifies whether the output is a squashed linear function
27 │         """
28 │         self.dataset = dataset
29 │         self.target = dataset.target
30 │         if train==None:
31 │             self.train = self.dataset.train
32 │         else:
33 │             self.train = train
34 │         self.learning_rate = learning_rate
35 │         self.squashed = squashed
36 │         self.input_features = dataset.input_features+[one] # one is defined below
37 │         self.weights = {feat:random.uniform(-max_init,max_init)
38 │                     for feat in self.input_features}
```

predictor predicts the value of an example from the current parameter settings. *predictor_string* gives a string representation of the predictor.

_____learnLinear.py — (continued) _____

```
40
41   def predictor(self,e):
42       """returns the prediction of the learner on example e"""
43       linpred = sum(w*f(e) for f,w in self.weights.items())
44       if self.squashed:
45           return sigmoid(linpred)
46       else:
47           return linpred
48
49   def predictor_string(self, sig_dig=3):
50       """returns the doc string for the current prediction function
51       sig_dig is the number of significant digits in the numbers"""
52       doc = "+".join(str(round(val,sig_dig))+"*"+feat.__doc__
53                       for feat,val in self.weights.items())
54       if self.squashed:
55           return "sigmoid("+ doc+")"
56       else:
57           return doc
```

learn is the main algorithm of the learner. It does *num_iter* steps of gradient descent. The other parameters it gets from the class.

_____learnLinear.py — (continued) _____

```
59   def learn(self,num_iter=100):
60       for it in range(num_iter):
61           self.display(2,"prediction=",self.predictor_string())
62           for e in self.train:
63               predicted = self.predictor(e)
64               error = self.target(e) - predicted
65               update = self.learning_rate*error
66               for feat in self.weights:
67                   self.weights[feat] += update*feat(e)
68       #self.predictor.__doc__ = self.predictor_string()
69       #return self.predictor
```

one is a function that always returns 1. This is used for one of the input properties.

_____learnLinear.py — (continued) _____

```
71 def one(e):
72     "1"
73     return 1
```

sigmoid(x) is the function

$$\frac{1}{1+e^{-x}}$$

_____learnLinear.py — (continued) _____

```
75  def sigmoid(x):
76      return 1/(1+math.exp(-x))
```

The following tests the learner on a data sets. Uncomment the other data sets for different examples.

_____learnLinear.py — (continued) _____

```
78  from learnProblem import Data_set, Data_from_file
79  import matplotlib.pyplot as plt
80  def test(**args):
81      data = Data_from_file('data/SPECT.csv', target_index=0)
82      # data = Data_from_file('data/mail_reading.csv', target_index=-1)
83      # data = Data_from_file('data/carbool.csv', target_index=-1)
84      learner = Linear_learner(data,**args)
85      learner.learn()
86      print("function learned is", learner.predictor_string())
87      for ecrit in Data_set.evaluation_criteria:
88          test_error = data.evaluate_dataset(data.test, learner.predictor, ecrit)
89          print("   Average", ecrit, "error is", test_error)
```

The following plots the errors on the training and test sets as a function of the number of steps of gradient descent.

_____learnLinear.py — (continued) _____

```
91  def plot_steps(learner=None,
92                 data = None,
93                 criterion="sum-of-squares",
94                 step=1,
95                 num_steps=1000,
96                 log_scale=True,
97                 label=""):
98      """
99      plots the training and test error for a learner.
100     data is the
101     learner_class is the class of the learning algorithm
102     criterion gives the evaluation criterion plotted on the y-axis
103     step specifies how many steps are run for each point on the plot
104     num_steps is the number of points to plot
105
106     """
107     plt.ion()
108     plt.xlabel("step")
109     plt.ylabel("Average "+criterion+" error")
110     if log_scale:
111         plt.xscale('log') #plt.semilogx() #Makes a log scale
112     else:
113         plt.xscale('linear')
114     if data is None:
115         data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
116         #data = Data_from_file('data/SPECT.csv', target_index=0)
```

```
117         # data = Data_from_file('data/mail_reading.csv', target_index=-1)
118         # data = Data_from_file('data/carbool.csv', target_index=-1)
119     random.seed(None)  # reset seed
120     if learner is None:
121         learner = Linear_learner(data)
122     train_errors = []
123     test_errors = []
124     for i in range(1,num_steps+1,step):
125         test_errors.append(data.evaluate_dataset(data.test, learner.predictor, criterion))
126         train_errors.append(data.evaluate_dataset(data.train, learner.predictor, criterion))
127         learner.display(2, "Train error:",train_errors[-1],
128                         "Test error:",test_errors[-1])
129         learner.learn(num_iter=step)
130     plt.plot(range(1,num_steps+1,step),train_errors,ls='-',c='k',label="training errors")
131     plt.plot(range(1,num_steps+1,step),test_errors,ls='--',c='k',label="test errors")
132     plt.legend()
133     plt.draw()
134     learner.display(1, "Train error:",train_errors[-1],
135                     "Test error:",test_errors[-1])
136
137 if __name__ == "__main__":
138     test()
139
140 # This generates the figure
141 # from learnProblem import Data_set_augmented,prod_feat
142 # data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
143 # dataplus = Data_set_augmented(data,[],[prod_feat])
144 # plot_steps(data=data,num_steps=10000)
145 # plot_steps(data=dataplus,num_steps=10000) # warning very slow
```

Exercise 7.8 The squashed learner only makes predictions in the range $(0, 1)$. If the output values are $\{1, 2, 3, 4\}$ there is no use prediction less than 1 or greater than 4. Change the squashed learner so that it can learn values in the range $(1, 4)$. Test it on the file `'data/car.csv'`.

The following plots the prediction as a function of the function of the number of steps of gradient descent. We first define a version of *range* that allows for real numbers (integers and floats).

_____learnLinear.py — (continued)_____
```
146 def arange(start,stop,step):
147     """returns enumeration of values in the range [start,stop) separated by step.
148     like the built-in range(start,stop,step) but allows for integers and floats.
149     Note that rounding errors are expected with real numbers.
150     """
151     while start<stop:
152         yield start
153         start += step
154
155 def plot_prediction(learner=None,
156             data = None,
```

```
157              minx = 0,
158              maxx = 5,
159              step_size = 0.01,  # for plotting
160              label="function"):
161     plt.ion()
162     plt.xlabel("x")
163     plt.ylabel("y")
164     if data is None:
165         data = Data_from_file('data/simp_regr.csv', prob_test=0,
166                          boolean_features=False, target_index=-1)
167     if learner is None:
168         learner = Linear_learner(data,squashed=False)
169     learner.learning_rate=0.001
170     learner.learn(100)
171     learner.learning_rate=0.0001
172     learner.learn(1000)
173     learner.learning_rate=0.00001
174     learner.learn(10000)
175     learner.display(1,"function learned is", learner.predictor_string(),
176             "error=",data.evaluate_dataset(data.train, learner.predictor, "sum-of-squares"))
177     plt.plot([e[0] for e in data.train],[e[-1] for e in data.train],"bo",label="data")
178     plt.plot(list(arange(minx,maxx,step_size)),[learner.predictor([x])
179                                     for x in arange(minx,maxx,step_size)],
180                                  label=label)
181     plt.legend()
182     plt.draw()
```

_____learnLinear.py — (continued)_____

```
184  from learnProblem import Data_set_augmented, power_feat
185  def plot_polynomials(data=None,
186              learner_class = Linear_learner,
187              max_degree=5,
188              minx = 0,
189              maxx = 5,
190              num_iter = 100000,
191              learning_rate = 0.0001,
192              step_size = 0.01, # for plotting
193              ):
194     plt.ion()
195     plt.xlabel("x")
196     plt.ylabel("y")
197     if data is None:
198         data = Data_from_file('data/simp_regr.csv', prob_test=0,
199                          boolean_features=False, target_index=-1)
200     plt.plot([e[0] for e in data.train],[e[-1] for e in data.train],"ko",label="data")
201     x_values = list(arange(minx,maxx,step_size))
202     line_styles = ['-','--','-.',':']
203     colors = ['0.5','k','k','k','k']
204     for degree in range(max_degree):
205         data_aug = Data_set_augmented(data,[power_feat(n) for n in range(1,degree+1)],
```

```
206                                    include_orig=False)
207            learner = learner_class(data_aug,squashed=False)
208            learner.learning_rate=learning_rate
209            learner.learn(num_iter)
210            learner.display(1,"For degree",degree,
211                    "function learned is", learner.predictor_string(),
212                    "error=",data.evaluate_dataset(data.train, learner.predictor, "sum-of-squares")
213            ls = line_styles[degree % len(line_styles)]
214            col = colors[degree % len(colors)]
215            plt.plot(x_values,[learner.predictor([x]) for x in x_values], linestyle=ls, color=col,
216                        label="degree="+str(degree))
217            plt.legend(loc='upper left')
218            plt.draw()
219
220  # Try:
221  # plot_prediction()
222  # plot_polynomials()
223  #data = Data_from_file('data/mail_reading.csv', target_index=-1)
224  #plot_prediction(data=data)
```

7.5.1   Batched Stochastic Gradient Descent

This implements batched stochastic gradient descent. If the batch size is 1, it can be simplified by not storing the differences in *d*, but applying them directly; this would the be equivalent to the original code!

This overrides the learner *Linear_learner*. Note that the comparison with regular gradient descent is unfair as the number of updates per step is not the same. (How could it me made more fair?)

———————————learnLinearBSGD.py — Linear Learner with Batched Stochastic Gradient Descent ———————————
```
11  from learnLinear import Linear_learner
12  import random, math
13
14  class Linear_learner_bsgd(Linear_learner):
15      def __init__(self, *args, batch_size=10, **kargs):
16          Linear_learner.__init__(self, *args, **kargs)
17          self.batch_size = batch_size
18
19      def learn(self,num_iter=None):
20          if num_iter is None:
21              num_iter = self.number_iterations
22          batch_size = min(self.batch_size, len(self.train))
23          d = {feat:0 for feat in self.weights}
24          for it in range(num_iter):
25              self.display(2,"prediction=",self.predictor_string())
26              for e in random.sample(self.train, batch_size):
27                  predicted = self.predictor(e)
28                  error = self.target(e) - predicted
29                  update = self.learning_rate*error
```

```
30                    for feat in self.weights:
31                        d[feat] += update*feat(e)
32                for feat in self.weights:
33                    self.weights[feat] += d[feat]
34                    d[feat]=0
35
36  # from learnLinear import plot_steps
37  # from learnProblem import Data_from_file
38  # data = Data_from_file('data/holiday.csv', target_index=-1)
39  # learner = Linear_learner_bsgd(data)
40  # plot_steps(learner = learner, data=data)
41
42  # to plot polynomials with batching (compare to SGD)
43  # from learnLinear import plot_polynomials
44  # plot_polynomials(learner_class = Linear_learner_bsgd)
```

7.6   Deep Neural Network Learning

This provides a modular implementation that implements the layers modularly. Layers can easily be configured in many configurations. A layer needs to implement a function to compute the output values from the inputs and a way to back-propagate the error.

_____learnNN.py — Neural Network Learning _____
```
11  from learnProblem import Learner, Data_set, Data_from_file
12  from learnLinear import sigmoid, one
13  import random, math
14
15  class Layer(object):
16      def __init__(self,nn,num_outputs=None):
17          """Given a list of inputs, outputs will produce a list of length num_outputs.
18          nn is the neural network this is part of
19          num outputs is the number of outputs for this layer.
20          """
21          self.nn = nn
22          self.num_inputs = nn.num_outputs # output of nn is the input to this layer
23          if num_outputs:
24              self.num_outputs = num_outputs
25          else:
26              self.num_outputs = nn.num_outputs # same as the inputs
27
28      def output_values(self,input_values):
29          """Return the outputs for this layer for the given input values.
30          input_values is a list of the inputs to this layer (of length num_inputs)
31          returns a list of length self.num_outputs
32          """
33          raise NotImplementedError("output_values") # abstract method
34
35      def backprop(self,errors):
```

```
36          """Backpropagate the errors on the outputs, return the errors on the inputs.
37          errors is a list of errors for the outputs (of length self.num_outputs).
38          Return the errors for the inputs to this layer (of length self.num_inputs).
39          You can assume that this is only called after corresponding output_values,
40             and it can remember information information required for the backpropagation.
41          """
42          raise NotImplementedError("backprop") # abstract method
```

A linear layer maintains an array of weights. *self.weights*[*o*][*i*] is the weight between input *i* and output *o*. A 1 is added to the inputs.

────────────── learnNN.py — (continued) ──────────────

```
44  class Linear_complete_layer(Layer):
45      """a completely connected layer"""
46      def __init__(self, nn, num_outputs, max_init=0.2):
47          """A completely connected linear layer.
48          nn is a neural network that the inputs come from
49          num_outputs is the number of outputs
50          max_init is the maximum value for random initialization of parameters
51          """
52          Layer.__init__(self, nn, num_outputs)
53          # self.weights[o][i] is the weight between input i and output o
54          self.weights = [[random.uniform(-max_init, max_init)
55                          for inf in range(self.num_inputs+1)]
56                         for outf in range(self.num_outputs)]
57
58      def output_values(self,input_values):
59          """Returns the outputs for the input values.
60          It remembers the values for the backprop.
61
62          Note in self.weights there is a weight list for every output,
63          so wts in self.weights effectively loops over the outputs.
64          """
65          self.inputs = input_values + [1]
66          return [sum(w*val for (w,val) in zip(wts,self.inputs))
67                     for wts in self.weights]
68
69      def backprop(self,errors):
70          """Backpropagate the errors, updating the weights and returning the error in its inputs.
71          """
72          input_errors = [0]*(self.num_inputs+1)
73          for out in range(self.num_outputs):
74              for inp in range(self.num_inputs+1):
75                  input_errors[inp] += self.weights[out][inp] * errors[out]
76                  self.weights[out][inp] += self.nn.learning_rate * self.inputs[inp] * errors[out]
77          return input_errors[:-1] # remove the error for the "1"
```

────────────── learnNN.py — (continued) ──────────────

```
79  class Sigmoid_layer(Layer):
80      """sigmoids of the inputs.
81      The number of outputs is equal to the number of inputs.
```

```
82        Each output is the sigmoid of its corresponding input.
83        """
84    def __init__(self, nn):
85        Layer.__init__(self, nn)
86
87    def output_values(self,input_values):
88        """Returns the outputs for the input values.
89        It remembers the output values for the backprop.
90        """
91        self.outputs= [sigmoid(inp) for inp in input_values]
92        return self.outputs
93
94    def backprop(self,errors):
95        """Returns the derivative of the errors"""
96        return [e*out*(1-out) for e,out in zip(errors, self.outputs)]
```

————————————————————learnNN.py — (continued) ————————————————————

```
98  class ReLU_layer(Layer):
99      """Rectified linear unit (ReLU) f(z) = max(0, z).
100     The number of outputs is equal to the number of inputs.
101     """
102     def __init__(self, nn):
103         Layer.__init__(self, nn)
104
105     def output_values(self,input_values):
106         """Returns the outputs for the input values.
107         It remembers the input values for the backprop.
108         """
109         self.input_values = input_values
110         self.outputs= [max(0,inp) for inp in input_values]
111         return self.outputs
112
113     def backprop(self,errors):
114         """Returns the derivative of the errors"""
115         return [e if inp>0 else 0 for e,inp in zip(errors, self.input_values)]
```

————————————————————learnNN.py — (continued) ————————————————————

```
117 class NN(Learner):
118     def __init__(self, dataset, learning_rate=0.1):
119         self.dataset = dataset
120         self.learning_rate = learning_rate
121         self.input_features = dataset.input_features
122         self.num_outputs = len(self.input_features)
123         self.layers = []
124
125     def add_layer(self,layer):
126         """add a layer to the network.
127         Each layer gets values from the previous layer.
128         """
129         self.layers.append(layer)
```

```
130          self.num_outputs = layer.num_outputs
131
132      def predictor(self,ex):
133          """Predicts the value of the first output feature for example ex.
134          """
135          values = [f(ex) for f in self.input_features]
136          for layer in self.layers:
137              values = layer.output_values(values)
138          return values[0]
139
140      def predictor_string(self):
141          return "not implemented"
```

The *test* method learns a network and evaluates it according to various criteria.

────────────── learnNN.py — (continued) ──────────────

```
143
144      def learn(self,num_iter):
145          """Learns parameters for a neural network using stochastic gradient decent.
146          num_iter is the number of iterations
147          """
148          for i in range(num_iter):
149              for e in random.sample(self.dataset.train,len(self.dataset.train)):
150                  # compute all outputs
151                  values = [f(e) for f in self.input_features]
152                  for layer in self.layers:
153                      values = layer.output_values(values)
154                  # backpropagate
155                  errors = self.sum_squares_error([self.dataset.target(e)],values)
156                  for layer in reversed(self.layers):
157                      errors = layer.backprop(errors)
158
159      def sum_squares_error(self,observed,predicted):
160          """Returns the errors for each of the target features.
161          """
162          return [obsd-pred for obsd,pred in zip(observed,predicted)]
```

This constructs a neural network consisting of neural network with one hidden layer. The hidden using used a ReLU activation function. The output layer used a sigmoid.

────────────── learnNN.py — (continued) ──────────────

```
165 data = Data_from_file('data/mail_reading.csv', target_index=-1)
166 #data = Data_from_file('data/mail_reading_consis.csv', target_index=-1)
167 #data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
168 #data = Data_from_file('data/holiday.csv', target_index=-1) #, num_train=19)
169 nn1 = NN(data)
170 nn1.add_layer(Linear_complete_layer(nn1,3))
171 nn1.add_layer(Sigmoid_layer(nn1)) # comment this or the next
172 # nn1.add_layer(ReLU_layer(nn1))
173 nn1.add_layer(Linear_complete_layer(nn1,1))
```

```
174  nn1.add_layer(Sigmoid_layer(nn1))
175  nn1.learning_rate=0.1
176  #nn1.learn(100)
177
178  from learnLinear import plot_steps
179  import time
180  start_time = time.perf_counter()
181  plot_steps(learner = nn1, data = data, num_steps=10000)
182  for eg in data.train:
183      print(eg,nn1.predictor(eg))
184  end_time = time.perf_counter()
185  print("Time:", end_time - start_time)
```

Exercise 7.9 In the definition of *nn*1 above, for each of the following, first hypothesize what will happen, then test your hypothesis, then explain whether you testing confirms your hypothesis or not. Test it for more than one data set, and use more than one run for each data set.

(a) Which fits the data better, having a sigmoid layer or a ReLU layer after the first linear layer?

(b) Which is faster, having a sigmoid layer or a ReLU layer after the first linear layer?

(c) What happens if you have both the sigmoid layer and then a ReLU layer after the first linear layer and before the second linear layer?

(d) What happens if you have neither the sigmoid layer nor a ReLU layer after the first linear layer?

(e) What happens if you have a ReLU layer then a sigmoid layer after the first linear layer and before the second linear layer?

Exercise 7.10 Do some

It is even possible to define a perceptron layer. Warning: you may need to change the learning rate to make this work. Should I add it into the code? It doesn't follow the official line.

```
class PerceptronLayer(Layer):
    def __init__(self, nn):
        Layer.__init__(self, nn)

    def output_values(self,input_values):
        """Returns the outputs for the input values.
        """
        self.outputs= [1 if inp>0 else -1 for inp in input_values]
        return self.outputs

    def backprop(self,errors):
        """Pass the errors through"""
        return errors
```

7.7 Boosting

The following code implements functional gradient boosting for regression.

A Boosted dataset is created from a base dataset by subtracting the prediction of the offset function from each example. This does not save the new dataset, but generates it as needed. The amount of space used is constant, independent on the size of the data set.

—————————————learnBoosting.py — Functional Gradient Boosting—————————————
```python
11  from learnProblem import Data_set, Learner
12
13  class Boosted_dataset(Data_set):
14      def __init__(self, base_dataset, offset_fun):
15          """new dataset which is like base_dataset,
16            but offset_fun(e) is subtracted from the target of each example e
17          """
18          self.base_dataset = base_dataset
19          self.offset_fun = offset_fun
20          Data_set.__init__(self, base_dataset.train, base_dataset.test,
21                          base_dataset.prob_test, base_dataset.target_index)
22
23      def create_features(self):
24          self.input_features = self.base_dataset.input_features
25          def newout(e):
26              return self.base_dataset.target(e) - self.offset_fun(e)
27          newout.frange = self.base_dataset.target.frange
28          self.target = newout
```

A boosting learner takes in a dataset and a base learner, and returns a new predictor. The base learner, takes a dataset, and returns a Learner object.

————————————————————————learnBoosting.py — (continued)————————————————————————
```python
30  class Boosting_learner(Learner):
31      def __init__(self, dataset, base_learner_class):
32          self.dataset = dataset
33          self.base_learner_class = base_learner_class
34          mean = sum(self.dataset.target(e)
35                    for e in self.dataset.train)/len(self.dataset.train)
36          self.predictor = lambda e:mean # function that returns mean for each example
37          self.predictor.__doc__ = "lambda e:"+str(mean)
38          self.offsets = [self.predictor]
39          self.errors = [data.evaluate_dataset(data.test, self.predictor, "sum-of-squares")]
40          self.display(1,"Predict mean test set error=", self.errors[0] )
41
42
43      def learn(self, num_ensemble=10):
44          """adds num_ensemble learners to the ensemble.
45          returns a new predictor.
46          """
47          for i in range(num_ensemble):
48              train_subset = Boosted_dataset(self.dataset, self.predictor)
```

```
49              learner = self.base_learner_class(train_subset)
50              new_offset = learner.learn()
51              self.offsets.append(new_offset)
52              def new_pred(e, old_pred=self.predictor, off=new_offset):
53                  return old_pred(e)+off(e)
54              self.predictor = new_pred
55              self.errors.append(data.evaluate_dataset(data.test, self.predictor,"sum-of-squares"))
56              self.display(1,"After Iteration",len(self.offsets)-1,"test set error=", self.errors[-1])
57          return self.predictor
```

For testing, *sp_DT_learner* returns a function that constructs a decision tree learner
where the minimum number of examples is a proportion of the number of
training examples. The value of 0.9 tends to have one split, and a value of 0.5
tends to have two splits (but test it). Thus this can be used to construct small
decision trees that can be used as weak learners.

_____learnBoosting.py — (continued) _____

```
59  # Testing
60
61  from learnDT import DT_learner
62  from learnProblem import Data_set, Data_from_file
63
64  def sp_DT_learner(min_prop=0.9):
65      def make_learner(dataset):
66          mne = len(dataset.train)*min_prop
67          return DT_learner(dataset,min_number_examples=mne)
68      return make_learner
69
70  data = Data_from_file('data/carbool.csv', target_index=-1)
71  #data = Data_from_file('data/SPECT.csv', target_index=0)
72  #data = Data_from_file('data/mail_reading.csv', target_index=-1)
73  #data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
74  learner9 = Boosting_learner(data, sp_DT_learner(0.9))
75  #learner7 = Boosting_learner(data, sp_DT_learner(0.7))
76  #learner5 = Boosting_learner(data, sp_DT_learner(0.5))
77  predictor9 =learner9.learn(10)
78  for i in learner9.offsets: print(i.__doc__)
79  import matplotlib.pyplot as plt
80
81  def plot_boosting(data,steps=10, thresholds=[0.5,0.1,0.01,0.001], markers=['-','--','-.',':'] ):
82      learners = [Boosting_learner(data, sp_DT_learner(th)) for th in thresholds]
83      predictors = [learner.learn(steps) for learner in learners]
84      plt.ion()
85      plt.xscale('linear') # change between log and linear scale
86      plt.xlabel("number of trees")
87      plt.ylabel(" error")
88      for (learner,(threshold,marker)) in zip(learners,zip(thresholds,markers)):
89          plt.plot(range(len(learner.errors)), learner.errors, ls=marker,c='k',
90                      label=str(round(threshold*100))+"% min example threshold")
91      plt.legend()
92      plt.draw()
```

```
93
94  # plot_boosting(data)
```

Chapter 8

Reasoning Under Uncertainty

8.1 Representing Probabilistic Models

In the implementation of probabilistic models we will assume that the variables
are objects, rather than the strings we used for CSPs. (Note that in the CSP code
variables could be anything; we just used strings for the examples.) We use a
class here because it is more amenable to extend to richer models, such as when
we introduce time.

A variable consists of a name and a domain. The domain of a variable is
a list or a tuple, as the ordering will matter in the representation of factors.
The code below internally uses the index of each value. We define a function
val_to_index that maps from the value to the index.

probVariables.py — Probabilistic Variables

```python
11 class Variable(object):
12     """A random variable.
13     name (string) - name of the variable
14     domain (list) - a list of the values for the variable.
15     Variables are ordered according to their name.
16     """
17
18     def __init__(self,name,domain):
19         self.name = name
20         self.size = len(domain)
21         self.domain = domain
22         self.val_to_index = {} # map from domain to index
23         for i,val in enumerate(domain):
24             self.val_to_index[val]=i
25
26     def __str__(self):
27         return self.name
```

A	B	C	Value
0	a	s	v_0
0	a	t	v_1
0	b	s	v_2
0	b	t	v_3
0	c	s	v_4
0	c	t	v_5
1	a	s	v_6
1	a	t	v_7
1	b	s	v_8
1	b	t	v_9
1	c	s	v_{10}
1	c	t	v_{11}

Figure 8.1: A representation for a factor for the variable ordering A, B, C

```
28
29       def __repr__(self):
30           return "Variable('"+self.name+"')"
```

8.2 Factors

Factors are functions from variables into values. The main problem with variable elimination is the amount of space used, because it saves the intermediate factors. (If instead it recomputed factors rather than saving the factors, it would be effectively enumerating the worlds, and so would be exponential in the number of variables). We only want to store the list of numbers, with as little bookkeeping as possible.

A total ordering of the variables, and a total ordering of the values in the domains of the variables induces a total ordering of the values of the factor according to the lexicographic ordering. E.g., suppose the domain of A is $[0, 1]$, domain of B is $['a', 'b', 'c']$, and the domain of C is $['s', 't']$, the ordering $[A, B, C]$ of variables induces an ordering on the values of the factor, as in Figure 8.1. We just need to store the list of variables and the v_is. For any assignment to A, B and C, we can compute the index of the value for that assignment. $A = a, B = b, C = c$ is stored at location $a' * 6 + b' * 2 + c'$, where a' is $A.val\_to\_index[a]$, and similarly for b' and c'.

```
_____probFactors.py — Factor manipulation for graphical models _____
11   from functools import reduce
12   #from probVariables import Variable
13
14   class Factor(object):
15       nextid=0 # each factor has a unique identifier; for printing
```

```
16
17      def __init__(self,variables):
18          """variables is the ordered list of variables
19          """
20          self.variables = variables # ordered list of variables
21          # Compute the size and the offsets for the variables
22          self.var_offsets = {}
23          self.size = 1
24          for i in range(len(variables)-1,-1,-1):
25              self.var_offsets[variables[i]]=self.size
26              self.size *= variables[i].size
27          self.id = Factor.nextid
28          self.name = "f"+str(self.id)
29          Factor.nextid += 1
```

For each factor, *get_value* returns the value of the factor for an assignment. An
assignment is a variable:value dictionary. The assignment must include all of
the variables involved in the factor, and can include variables not in the factor.
This needs to be defined for every subclass.

_____probFactors.py — (continued) _____

```
31      def get_value(self,assignment):
32          raise NotImplementedError("get_value") # abstract method
```

The methods *str* and *brief* return string representations of the factor, as a table
or just as a name with the variables it is a factor on.

_____probFactors.py — (continued) _____

```
34      def __str__(self, variables=None):
35          """returns a string representation of the factor.
36          Allows for an arbitrary variable ordering.
37          variables is a list of the variables in the factor
38          (can contain other variables)"""
39          if variables==None:
40              variables = self.variables
41          else:
42              variables = [v for v in variables if v in self.variables]
43          res = ""
44          for v in variables:
45              res += str(v) + "\t"
46          res += self.name+"\n"
47          for i in range(self.size):
48              asst = self.index_to_assignment(i)
49              for v in variables:
50                  res += str(asst[v])+"\t"
51              res += str(self.get_value(asst))
52              res += "\n"
53          return res
54
55      def brief(self):
56          """returns a string representing a summary of the factor"""
```

```
57          res = self.name+"("
58          for i in range(0,len(self.variables)-1):
59              res += str(self.variables[i])+","
60          if len(self.variables)>0:
61              res += str(self.variables[len(self.variables)-1])
62          res += ")"
63          return res
```

The methods *assignment_to_index* and *index_to_assignment* map between the assignments of values to variables and the index of where that assignment would be stored.

_____probFactors.py — (continued) _____

```
65      def assignment_to_index(self,assignment):
66          """returns the index where the variable:value assignment is stored"""
67          index = 0
68          for var in self.variables:
69              index += var.val_to_index[assignment[var]]*self.var_offsets[var]
70          return index
71
72      def index_to_assignment(self,index):
73          """gives a dict representation of the variable assignment for index
74          """
75          asst = {}
76          for i in range(len(self.variables)-1,-1,-1):
77              asst[self.variables[i]] = self.variables[i].domain[index % self.variables[i].size]
78              index = index // self.variables[i].size
79          return asst
```

A *Factor_stored* is a factor that has the values stored in a list.

_____probFactors.py — (continued) _____

```
81  class Factor_stored(Factor):
82      def __init__(self,variables,values):
83          Factor.__init__(self, variables)
84          self.values = values
85
86      def get_value(self,assignment):
87          return self.values[self.assignment_to_index(assignment)]
```

A *Factor_observed* is a factor that is the result of some observations on another factor. We don't store the values in a list; we just look them up as needed. The observations can include variables that are not in the list, but should have some intersection with the variables in the factor.

_____probFactors.py — (continued) _____

```
89  class Factor_observed(Factor):
90      def __init__(self,factor,obs):
91          Factor.__init__(self, [v for v in factor.variables if v not in obs])
92          self.observed = obs
93          self.orig_factor = factor
```

```
94
95      def get_value(self,assignment):
96          ass = assignment.copy()
97          for ob in self.observed:
98              ass[ob]=self.observed[ob]
99          return self.orig_factor.get_value(ass)
```

A *Factor_sum* is a factor that is the result of summing out a variable from the product of other factors. Ie., it constructs a representation of:

$$\sum_{var} \prod_{f \in factors} f.$$

We store the values in a list in a lazy manner; if they are already computed, we used the stored values. If they are not already computed we can compute and store them.

_____ probFactors.py — (continued) _____
```
101  class Factor_sum(Factor_stored):
102      def __init__(self,var,factors):
103          self.var_summed_out = var
104          self.factors = factors
105          vars = []
106          for fac in factors:
107              for v in fac.variables:
108                  if v is not var and v not in vars:
109                      vars.append(v)
110          Factor_stored.__init__(self,vars,None)
111          self.values = [None]*self.size
112
113      def get_value(self,assignment):
114          """lazy implementation: if not saved, compute it. Return saved value"""
115          index = self.assignment_to_index(assignment)
116          if self.values[index]:
117              return self.values[index]
118          else:
119              total = 0
120              new_asst = assignment.copy()
121              for val in self.var_summed_out.domain:
122                  new_asst[self.var_summed_out] = val
123                  prod = 1
124                  for fac in self.factors:
125                      prod *= fac.get_value(new_asst)
126                  total += prod
127              self.values[index] = total
128              return total
```

The method *factor_times* multiples a set of factors that are all factors on the same variable (or on no variables). This is the last step in variable elimination before normalizing. It returns an array giving the product for each value of *variable*.

---
probFactors.py — (continued)

```python
130  def factor_times(variable,factors):
131      """when factors are factors just on variable (or on no variables)"""
132      prods= []
133      facs = [f for f in factors if variable in f.variables]
134      for val in variable.domain:
135          prod = 1
136          ast = {variable:val}
137          for f in facs:
138              prod *= f.get_value(ast)
139          prods.append(prod)
140      return prods
```

Prob is a factor that represents a conditional probability.

---
probFactors.py — (continued)

```python
142  class Prob(Factor_stored):
143      """A factor defined by a conditional probability table"""
144      def __init__(self,var,pars,cpt):
145          """Creates a factor from a conditional probability table, cptf.
146          The cpt values are assumed to be for the ordering par+[var]
147          """
148          Factor_stored.__init__(self,pars+[var],cpt)
149          self.child = var
150          self.parents = pars
151          assert self.size==len(cpt),"Table size incorrect "+str(self)
```

cond_dist returns the probability distribution of the child given values from the
parent. This code is based on *assignment_to_index*. Similarly, *cont_prob* returns
the probability that the child has a particular value given an assignment of
values to the parents. In both of these *par_assignment* is a dict that assigns all of
the parents (and can also assign other variables, but these are ignored).

---
probFactors.py — (continued)

```python
153      def cond_dist(self,par_assignment):
154          """returns the distribution (a val:prob dictionary) over the child given
155          assignment to the parents
156
157          par_assignment is a variable:value dictionary that assigns values to parents
158          """
159          index = 0
160          for var in self.parents:
161              index += var.val_to_index[par_assignment[var]]*self.var_offsets[var]
162          # index is the position where the disgribution starts
163          return {self.child.domain[i]:self.values[index+i] for i in range(len(self.child.domain))}
164
165      def cond_prob(self,par_assignment,child_value):
166          """returns the probability child has child_value given
167          assignment to the parents
168
169          par_assignment is a variable:value dictionary that assigns values to parents
```

```
170            child_value is a value to the child
171            """
172            index = self.child.val_to_index[child_value]
173            for var in self.parents:
174                index += var.val_to_index[par_assignment[var]]*self.var_offsets[var]
175            return self.values[index]
```

A *Factor_rename* is a factor that is the result renaming the variables in the factor. It takes a factor, *fac*, and a *new* : *old* dictionary, where *new* is the name of a variable in the resulting factor and *old* is the corresponding name in *fac*. This assumes that the all variables are renamed.

─────────────────── probFactors.py — (continued) ───────────────────

```
177  class Factor_rename(Factor):
178      def __init__(self,fac,renaming):
179          Factor.__init__(self,list(renaming.keys()))
180          self.orig_fac = fac
181          self.renaming = renaming
182
183      def get_value(self,assignment):
184          return self.orig_fac.get_value({self.renaming[var]:val
185                                          for (var,val) in assignment.items()
186                                          if var in self.variables})
```

8.3 Graphical Models

A graphical model consists of a set of variables and a set of factors. A belief network is a graphical model where all of the factors represent conditional probabilities. There are some operations (such as pruning variables) which are applicable to belief networks, but are not applicable to more general models. At the moment, we will treat them as the same.

─────────── probGraphicalModels.py — Graphical Models and Belief Networks ───────────

```
11  class Graphical_model(object):
12      """The class of graphical models.
13      A graphical model consists of a set of variables and a set of factors.
14
15      List vars is a list of variables
16      List factors is a list of factors
17      """
18      def __init__(self,vars=None,factors=None):
19          self.variables = vars
20          self.factors = factors
```

A belief network is a graphical model where all of the factors are conditional probabilities, and every variable has a conditional probability. This only checks the first condition:

─────────────────── probGraphicalModels.py — (continued) ───────────────────

```
22  class Belief_network(Graphical_model):
23      """The class of belief networks."""
24
25      def __init__(self,vars=None,factors=None):
26          """vars is a list of variables
27          factors is a list of factors. Here we assume that all of the factors are instances of Prob.
28          """
29          Graphical_model.__init__(self,vars,factors)
30          assert all(isinstance(f,Prob) for f in factors) if factors else True
```

Each of the inference methods implements the query method that computes the posterior probability of a variable given a dictionary of variable:value observations. These are all Displayable because they implement the *display* method which is currently text-based.

—————————————— probGraphicalModels.py — (continued) ——————————————

```
32  from utilities import Displayable
33
34  class Inference_method(Displayable):
35      """The abstract class of graphical model inference methods"""
36      def query(self,qvar,obs={}):
37          raise NotImplementedError("Inference_method query") # abstract method
```

The first example belief network is a simple chain $A \longrightarrow B \longrightarrow C$.

—————————————— probGraphicalModels.py — (continued) ——————————————

```
39  from probVariables import Variable
40  from probFactors import Prob
41
42  boolean = [False, True]
43  A = Variable("A", boolean)
44  B = Variable("B", boolean)
45  C = Variable("C", boolean)
46
47  f_a = Prob(A,[],[0.4,0.6])
48  f_b = Prob(B,[A],[0.9,0.1,0.2,0.8])
49  f_c = Prob(C,[B],[0.5,0.5,0.3,0.7])
50
51  bn1 = Belief_network([A,B,C],[f_a,f_b,f_c])
```

The second Bayesian network is the report-of-leaving example from Poole and Mackworth, Artificial Intelligence, 2010 `http://artint.info`. This is Example 6.10 (page 236) shown in Figure 6.1.

—————————————— probGraphicalModels.py — (continued) ——————————————

```
53  # Bayesian network report of leaving example from
54  # Poole and Mackworth, Artificial Intelligence, 2010 http://artint.info
55  # This is Example 6.10 (page 236) shown in Figure 6.1
56
57  Al = Variable("Alarm", boolean)
58  Fi = Variable("Fire", boolean)
```

```
59  Le = Variable("Leaving", boolean)
60  Re = Variable("Report", boolean)
61  Sm = Variable("Smoke", boolean)
62  Ta = Variable("Tamper", boolean)
63
64  f_ta = Prob(Ta,[],[0.98,0.02])
65  f_fi = Prob(Fi,[],[0.99,0.01])
66  f_sm = Prob(Sm,[Fi],[0.99,0.01,0.1,0.9])
67  f_al = Prob(Al,[Fi,Ta],[0.9999, 0.0001, 0.15, 0.85, 0.01, 0.99, 0.5, 0.5])
68  f_lv = Prob(Le,[Al],[0.999, 0.001, 0.12, 0.88])
69  f_re = Prob(Re,[Le],[0.99, 0.01, 0.25, 0.75])
70
71  bn2 = Belief_network([Al,Fi,Le,Re,Sm,Ta],[f_ta,f_fi,f_sm,f_al,f_lv,f_re])
```

The third Bayesian network is the sprinkler example from Pearl.

_____ probGraphicalModels.py — (continued) _____

```
73
74  Season = Variable("Season",["summer","winter"])
75  Sprinkler = Variable("Sprinkler",["on","off"])
76  Rained = Variable("Rained",boolean)
77  Grass_wet = Variable("Grass wet",boolean)
78  Grass_shiny = Variable("Grass shiny",boolean)
79  Shoes_wet = Variable("Shoes wet",boolean)
80
81  f_season = Prob(Season,[],[0.5,0.5])
82  f_sprinkler = Prob(Sprinkler,[Season],[0.9,0.1,0.05,0.95])
83  f_rained = Prob(Rained,[Season],[0.7,0.3,0.2,0.8])
84  f_wet = Prob(Grass_wet,[Sprinkler,Rained], [1,0,0.1,0.9,0.2,0.8,0.02,0.98])
85  f_shiny = Prob(Grass_shiny, [Grass_wet], [0.95,0.05,0.3,0.7])
86  f_shoes = Prob(Shoes_wet, [Grass_wet], [0.92,0.08,0.35,0.65])
87
88  bn3 = Belief_network([Season, Sprinkler, Rained, Grass_wet, Grass_shiny, Shoes_wet],
89              [f_season, f_sprinkler, f_rained, f_wet, f_shiny, f_shoes])
```

8.4 Variable Elimination

An instance of a *VE* object takes in a graphical model. The query method uses variable elimination to compute the probability of a variable given observations on some other variables.

_____ probVE.py — Variable Elimination for Graphical Models _____

```
11  from probFactors import Factor, Factor_observed, Factor_sum, factor_times
12  from probGraphicalModels import Graphical_model, Inference_method
13
14  class VE(Inference_method):
15      """The class that queries Graphical Models using variable elimination.
16
17      gm is graphical model to query
```

```
18          """
19      def __init__(self,gm=None):
20          self.gm = gm
21
22      def query(self,var,obs={},elim_order=None):
23          """computes P(var|obs) where
24          var is a variable
25          obs is a variable:value dictionary"""
26          if var in obs:
27              return [1 if val == obs[var] else 0 for val in var.domain]
28          else:
29              if elim_order == None:
30                  elim_order = self.gm.variables
31              projFactors = [self.project_observations(fact,obs)
32                             for fact in self.gm.factors]
33              for v in elim_order:
34                  if v != var and v not in obs:
35                      projFactors = self.eliminate_var(projFactors,v)
36              unnorm = factor_times(var,projFactors)
37              p_obs=sum(unnorm)
38              self.display(1,"Unnormalized probs:",unnorm,"Prob obs:",p_obs)
39              return {val:pr/p_obs for val,pr in zip(var.domain, unnorm)}
```

To project observations onto a factor, for each variable that is observed in the factor, we construct a new factor that is the factor projected onto that variable. *Factor_observed* creates a new factor that is the result is assigning a value to a single variable.

---
probVE.py — (continued)

```
41      def project_observations(self,factor,obs):
42          """Returns the resulting factor after observing obs
43
44          obs is a dictionary of variable:value pairs.
45          """
46          if any((var in obs) for var in factor.variables):
47              # a variable in factor is observed
48              return Factor_observed(factor,obs)
49          else:
50              return factor
51
52      def eliminate_var(self,factors,var):
53          """Eliminate a variable var from a list of factors.
54          Returns a new set of factors that has var summed out.
55          """
56          self.display(2,"eliminating ",str(var))
57          contains_var = []
58          not_contains_var = []
59          for fac in factors:
60              if var in fac.variables:
61                  contains_var.append(fac)
62              else:
```

```
63                    not_contains_var.append(fac)
64           if contains_var == []:
65               return factors
66           else:
67               newFactor = Factor_sum(var,contains_var)
68               self.display(2,"Multiplying:",[f.brief() for f in contains_var])
69               self.display(2,"Creating factor:", newFactor.brief())
70               self.display(3, newFactor) # factor in detail
71               not_contains_var.append(newFactor)
72               return not_contains_var
73
74  from probGraphicalModels import bn1, A,B,C
75  bn1v = VE(bn1)
76  ## bn1v.query(A,{})
77  ## bn1v.query(C,{})
78  ## Inference_method.max_display_level = 3 # show more detail in displaying
79  ## Inference_method.max_display_level = 1 # show less detail in displaying
80  ## bn1v.query(A,{C:True})
81  ## bn1v.query(B,{A:True,C:False})
82
83  from probGraphicalModels import bn2,Al,Fi,Le,Re,Sm,Ta
84  bn2v = VE(bn2)   # answers queries using variable elimination
85  ## bn2v.query(Ta,{})
86  ## Inference_method.max_display_level = 0 # show no detail in displaying
87  ## bn2v.query(Le,{})
88  ## bn2v.query(Ta,{},elim_order=[Sm,Re,Le,Al,Fi])
89  ## bn2v.query(Ta,{Re:True})
90  ## bn2v.query(Ta,{Re:True,Sm:False})
91
92  from probGraphicalModels import bn3, Season, Sprinkler, Rained, Grass_wet, Grass_shiny, Shoes_wet
93  bn3v = VE(bn3)
94  ## bn3v.query(Shoes_wet,{})
95  ## bn3v.query(Shoes_wet,{Rained:True})
96  ## bn3v.query(Shoes_wet,{Grass_shiny:True})
97  ## bn3v.query(Shoes_wet,{Grass_shiny:False,Rained:True})
```

8.5   Stochastic Simulation

8.5.1   Sampling from a discrete distribution

The method *sample_one* generates a single sample from a (possible unnormalized) distribution. *dist* is a *value* : *weight* dictionary, where *weight* $\geq$ 0. This returns a value with probability in proportion to its weight.

─────────probStochSim.py — Probabilistic inference using stochastic simulation ─────────
```
11  import random
12  from probGraphicalModels import Inference_method
13
14  def sample_one(dist):
```

```
15        """returns the index of a single sample from normalized distribution dist."""
16        rand = random.random()*sum(dist.values())
17        cum = 0    # cumulative weights
18        for v in dist:
19            cum += dist[v]
20            if cum > rand:
21                return v
```

If we want to generate multiple samples, repeatedly calling *sample_one* may not be efficient. If we want to generate n samples, and the distribution is over m values, *sample_one* takes time $O(mn)$. If m and n are of the same order of magnitude, we can do better.

The method *sample_multiple* generates multiple samples from a distribution defined by *dist*, where *dist* is a *value* : *weight* dictionary, where *weight* ≥ 0 and the weights cannot all be zero. This returns a list of values, of length *num_samples*, where each sample is selected with a probability proportional to its weight.

The method generates all of the random numbers, sorts them, and then goes through the distribution once, saving the selected samples.

——————————— probStochSim.py — (continued) ———————————

```
23  def sample_multiple(dist, num_samples):
24      """returns a list of num_samples values selected using distribution dist.
25      dist is a value:weight dictionary that does not need to be normalized
26      """
27      total = sum(dist.values())
28      rands = sorted(random.random()*total for i in range(num_samples))
29      result = []
30      dist_items = list(dist.items())
31      cum = dist_items[0][1] # cumulative sum
32      index = 0
33      for r in rands:
34          while r>cum:
35              index += 1
36              cum += dist_items[index][1]
37          result.append(dist_items[index][0])
38      return result
```

Exercise 8.1

What is the time and space complexity the following 4 methods to generate n samples, where m is the length of *dist*:

(a) n calls to *sample_one*

(b) *sample_multiple*

(c) Create the cumulative distribution (choose how this is represented) and, for each random number, do a binary search to determine the sample associated with the random number.

(d) Choose a random number in the range $[i/n, (i+1)/n)$ for each $i \in range(n)$, where n is the number of samples. Use these as the random numbers to select the particles. (Does this give random samples?)

For each method suggest when it might be the best method.

The *test_sampling* method can be used to generate the statistics from a number of samples. It is useful to see the variability as a function of the number of samples. Try it for few samples and also for many samples.

_____ probStochSim.py — (continued) _____

```
40  def test_sampling(dist, num_samples):
41      """Given a distribution, dist, draw num_samples samples
42      and return the resulting counts
43      """
44      result = {v:0 for v in dist}
45      for v in sample_multiple(dist, num_samples):
46          result[v] += 1
47      return result
48
49  # try the following queries a number of times each:
50  # test_sampling({1:1,2:2,3:3,4:4}, 100)
51  # test_sampling({1:1,2:2,3:3,4:4}, 100000)
```

8.5.2   Sampling Methods for Belief Network Inference

A *Sampling_inference_method* is an *Inference_method*, but the query method also takes arguments for the number of samples and the sample-order (which is an ordering of factors). The first methods assume a Bayesian network (and not an undirected graphical model).

_____ probStochSim.py — (continued) _____

```
53  class Sampling_inference_method(Inference_method):
54      """The abstract class of sampling-based belief network inference methods"""
55      def query(self,qvar,obs={},number_samples=1000,sample_order=None):
56          raise NotImplementedError("Sampling_inference_method query") # abstract
```

Some of the sampling methods require a sample order of factors representing conditional probabilities, where the parents of a node must come before the node in the sample order. The following method computes such a sample ordering, and is used when the *sample_order* argument is *None*.

_____ probStochSim.py — (continued) _____

```
58  def select_sample_ordering(bn):
59      """creates a sample ordering of factors such that the parents of a node
60      are before the node.
61      raises StopIteration if there is no such ordering. This would occur in next(.).
62      """
63      sample_order=[]
64      defined = set() # set of variables whose probability is defined
65      factors_to_sample = bn.factors.copy()
66      while factors_to_sample:
67          fac = next(f for f in factors_to_sample
68                  if all(par in defined for par in f.parents))
```

```
69              factors_to_sample.remove(fac)
70              sample_order.append(fac)
71              defined.add(fac.child)
72          return sample_order
```

8.5.3   Rejection Sampling

—————————————————————probStochSim.py — (continued) —————————————————————
```
74  class Rejection_sampling(Sampling_inference_method):
75      """The class that queries Graphical Models using Rejection Sampling.
76
77      bn is a belief network to query
78      """
79      def __init__(self,bn=None):
80          self.bn = bn
81          self.label = "Rejection Sampling"
82
83      def query(self,qvar,obs={},number_samples=1000,sample_order=None):
84          """computes P(qvar|obs) where
85          qvar is a variable.
86          obs is a variable:value dictionary.
87          sample_order is a list of factors where factors defining the parents
88            come before the factors for the child.
89          """
90          if sample_order is None:
91              sample_order = select_sample_ordering(self.bn)
92          self.display(2,*[f.child for f in sample_order],sep="\t")
93          counts = {val:0 for val in qvar.domain}
94          for i in range(number_samples):
95              rejected = False
96              sample = {}
97              for fac in sample_order:
98                  nvar = fac.child    #next variable
99                  val = sample_one(fac.cond_dist(sample))
100                 self.display(2,val,end="\t")
101                 if nvar in obs and obs[nvar] != val:
102                     rejected = True
103                     self.display(2,"Rejected")
104                     break
105                 sample[nvar] = val
106             if not rejected:
107                 counts[sample[qvar]] += 1
108                 self.display(2,"Accepted")
109         tot = sum(counts.values())
110         return counts, {c:divide(v,tot) for (c,v) in counts.items()}
```

It is possible that all samples get rejected. In that case, Python would give as a arithmetic error. Instead, we implement the convention that $0/0 = 1$. You need to be careful is using these numbers as probabilities.

probStochSim.py — (continued)

```
112  def divide(num,denom):
113      """returns num/denom without divide-by-zero errors.
114      defines 0/0 to be 1."""
115      if denom == 0:
116          return 1.0
117      else:
118          return num/denom
```

8.5.4 Likelihood Weighting

Likelihood weighting includes a weight for each sample. Instead of rejecting samples based on observations, likelihood weighting changes the weights of the sample in proportion with the probability of the observation. The weight then becomes the probability that the variable would have been rejected.

probStochSim.py — (continued)

```
120  class Likelihood_weighting(Sampling_inference_method):
121      """The class that queries Graphical Models using Likelihood weighting.
122
123      bn is a belief network to query
124      """
125      def __init__(self,bn=None):
126          self.bn = bn
127          self.label = "Likelihood weighting"
128
129      def query(self,qvar,obs={},number_samples=1000,sample_order=None):
130          """computes P(qvar|obs) where
131          qvar is a variable.
132          obs is a variable:value dictionary.
133          sample_order is a list of factors where factors defining the parents
134            come before the factors for the child.
135          """
136          if sample_order is None:
137              sample_order = select_sample_ordering(self.bn)
138          self.display(2,*[f.child for f in sample_order
139                          if f.child not in obs],sep="\t")
140          counts = [0 for val in qvar.domain]
141          for i in range(number_samples):
142              sample = {}
143              weight = 1.0
144              for fac in sample_order:
145                  nvar = fac.child # next variable sampled
146                  if nvar in obs:
147                      sample[nvar] = obs[nvar]
148                      weight *= fac.get_value(sample)
149                  else:
150                      val = sample_one(fac.cond_dist(sample))
151                      self.display(2,val,end="\t")
152                      sample[nvar] = val
153              counts[sample[qvar]] += weight
```

```
154                 self.display(2,weight)
155             tot = sum(counts)
156             return counts, {c:v/tot for (c,v) in counts.items()}
```

Exercise 8.2 Change this algorithm so that it does **importance sampling** using a proposal distribution. It needs *sample_one* using a different distribution and then update the weight of the current sample. For testing, use a proposal distribution that only specifies probabilities for some of the variables (and the algorithm uses the probabilities for the network in other cases).

8.5.5   Particle Filtering

In this implementation, a particle is a *variable* : *value* dictionary. Because adding a new value to dictionary involves a side effect, the dictionaries need to be copied during resampling.

_____probStochSim.py — (continued) _____

```
158  class Particle_filtering(Sampling_inference_method):
159      """The class that queries Graphical Models using Particle Filtering.
160
161      bn is a belief network to query
162      """
163      def __init__(self,bn=None):
164          self.bn = bn
165          self.label = "Particle Filtering"
166
167      def query(self, qvar, obs={}, number_samples=1000, sample_order=None):
168          """computes P(qvar|obs) where
169          qvar is a variable.
170          obs is a variable:value dictionary.
171          sample_order is a list of factors where factors defining the parents
172            come before the factors for the child.
173          """
174          if sample_order is None:
175              sample_order = select_sample_ordering(self.bn)
176          self.display(2,*[f.child for f in sample_order
177                          if f.child not in obs],sep="\t")
178          particles = [{} for i in range(number_samples)]
179          for fac in sample_order:
180              nvar = fac.child # the variable sampled
181              if nvar in obs:
182                  weights = {part:fac.cond_prob(part,obs[nvar]) for part in particles}
183                  particles = [p.copy for p in resample(particles, weights, number_samples)]
184              else:
185                  for part in particles:
186                      part[nvar] = sample_one(fac.cond_dist(part))
187                  self.display(2,part[nvar],end="\t")
188          counts = [0 for val in qvar.domain]
189          for part in particles:
190              counts[part[qvar]] += 1
```

```
191          self.display(2,weight)
192          return counts
```

Resampling

Resample is based on *sample_multiple* but works with an array of particles. (Aside: Python doesn't let us use *sample_multiple* directly as it uses a dictionary, and particles, represented as dictionaries can't be the key of dictionaries).

—————————————— probStochSim.py — (continued) ——————————————

```
194  def resample(particles, weights, num_samples):
195      """returns num_samples copies of particles resampled according to weights.
196      particles is a list of particles
197      weights is a list of positive numbers, of same length as particles
198      num_samples is n integer
199      """
200      total = sum(weights)
201      rands = sorted(random.random()*total for i in range(num_samples))
202      result = []
203      cum = weights[0]   # cumulative sum
204      index = 0
205      for r in rands:
206          while r>cum:
207              index += 1
208              cum += weights[index]
209          result.append(particles[index])
210      return result
```

8.5.6 Examples

—————————————— probStochSim.py — (continued) ——————————————

```
212  from probGraphicalModels import bn1, A,B,C
213  bn1r = Rejection_sampling(bn1)
214  bn1L = Likelihood_weighting(bn1)
215  ## Inference_method.max_display_level = 2 # detailed tracing for all inference methods
216  ## bn1r.query(A,{})
217  ## bn1r.query(C,{})
218  ## bn1r.query(A,{C:True})
219  ## bn1r.query(B,{A:True,C:False})
220
221  from probGraphicalModels import bn2,Al,Fi,Le,Re,Sm,Ta
222  bn2r = Rejection_sampling(bn2) # answers queries using rejection sampling
223  bn2L = Likelihood_weighting(bn2) # answers queries using rejection sampling
224  bn2p = Particle_filtering(bn2) # answers queries using particle filtering
225  ## bn2r.query(Ta,{})
226  ## bn2r.query(Ta,{})
227  ## bn2r.query(Ta,{Re:True})
228  ## Inference_method.max_display_level = 0 # no detailed tracing for all inference methods
```

```
229  ## bn2r.query(Ta,{Re:True},number_samples=100000)
230  ## bn2r.query(Ta,{Re:True,Sm:False})
231  ## bn2r.query(Ta,{Re:True,Sm:False},number_samples=100)
232
233  ## bn2L.query(Ta,{Re:True,Sm:False},number_samples=100)
234  ## bn2L.query(Ta,{Re:True,Sm:False},number_samples=100)
235
236
237  from probGraphicalModels import bn3,Season, Sprinkler
238  from probGraphicalModels import Rained, Grass_wet, Grass_shiny, Shoes_wet
239  bn3r = Rejection_sampling(bn3) # answers queries using rejection sampling
240  bn3L = Likelihood_weighting(bn3) # answers queries using rejection sampling
241  bn3p = Particle_filtering(bn3) # answers queries using particle filtering
242  #bn3r.query(Shoes_wet,{Grass_shiny:True,Rained:True})
243  #bn3L.query(Shoes_wet,{Grass_shiny:True,Rained:True})
244  #bn3p.query(Shoes_wet,{Grass_shiny:True,Rained:True})
```

Exercise 8.3 This code keeps regenerating the distribution of a variable given its parents. Implement one or both of the following, and compare them to the original. Make *cond_dist* return a slice that corresponds to the distribution, and then use the slice instead of the dictionary (a list slice does not generate new data structures). Make *cond_dist* remember values it has already computed, and only return these.

8.5.7 Plotting Behaviour of Stochastic Simulators

The stochastic simulation runs can give different answers each time they are run. For the algorithms that give the same answer in the limit as the number of samples approaches infinity (as do all of these algorithms), the algorithms can be compared by comparing the accuracy for multiple runs. Summary statistics like the variance may provide some information, but the assumptions behind the variance being appropriate (namely that the distribution is approximately Gaussian) may not hold for cases where the predictions are bounded and often skewed.

It is more appropriate to plot the distribution of predictions over multiple runs. The *plot_stats* method plots the prediction of a particular variable (or for the partition function) for a number of runs of the same algorithm. On the *x*-axis, is the prediction of the algorithm. On the *y*-axis is the number of runs with prediction less than or equal to the *x* value. Thus this is like a cumulative distribution over the predictions, but with counts on the *y*-axis.

Note that for runs where there are no samples that are consistent with the observations (as can happen with rejection sampling), the prediction of probability is 1.0 (as a convention for 0/0).

That variable *what* contains the query variable, or *what* is "*prob_ev*", the probability of evidence.

───────────────── probStochSim.py — (continued) ─────────────────

```
246  import matplotlib.pyplot as plt
```

```
247
248 def plot_stats(method, what, qvar, obs, number_samples=100, number_runs=1000):
249     """Plots a cumulative distribution of the prediction of the model.
250     method is a Sampling_inference_method (that implements appropriate query(.))
251     what is either "prob_ev" or the value of qvar to plot
252     qvar is the query variable
253     obs is the variable:value dictionary representing the observations
254     number_samples is the number of samples for each run
255     number_iterations is the number of runs that are plotted
256     """
257     plt.ion()
258     plt.xlabel("value")
259     plt.ylabel("Cumulative Number")
260     Inference_method.max_display_level, prev_max_display_level = 0, Inference_method.max_display_l
261     answers = [method.query(qvar,obs,number_samples=number_samples)
262                 for i in range(number_runs)]
263     if what == "prob_ev":
264         values = [sum(ans)/number_samples for ans in answers]
265         label = method.label+"(prob of evidence)"
266     else:
267         values = [divide(ans[qvar.val_to_index[what]],sum(ans)) for ans in answers]
268         label = method.label+" ("+str(qvar)+"="+str(what)+")"
269     values.sort()
270     plt.plot(values,range(number_runs),label=label)
271     plt.legend(loc="upper left")
272     plt.draw()
273     Inference_method.max_display_level = prev_max_display_level # restore display level
274
275
276 # plot_stats(bn2r,False,Ta,{Re:True,Sm:False},number_samples=1000, number_runs=1000)
277 # plot_stats(bn2L,False,Ta,{Re:True,Sm:False},number_samples=1000, number_runs=1000)
278 # plot_stats(bn2r,False,Ta,{Re:True,Sm:False},number_samples=100, number_runs=1000)
279 # plot_stats(bn2L,False,Ta,{Re:True,Sm:False},number_samples=100, number_runs=1000)
280 # plot_stats(bn3r,True,Shoes_wet,{Grass_shiny:True,Rained:True},number_samples=1000)
281 # plot_stats(bn3L,True,Shoes_wet,{Grass_shiny:True,Rained:True},number_samples=1000)
282 # plot_stats(bn2r,"prob_ev",Ta,{Re:True,Sm:False},number_samples=1000, number_runs=1000)
283 # plot_stats(bn2L,"prob_ev",Ta,{Re:True,Sm:False},number_samples=1000, number_runs=1000)
```

8.6   Markov Chain Monte Carlo

The following implements **Gibbs sampling**, a form of **Markov Chain Monte Carlo** MCMC.

_____probMCMC.py — Markov Chain Monte Carlo (Gibbs sampling) _____

```
11 import random
12 from probGraphicalModels import Inference_method
13
14 from probStochSim import sample_one, Sampling_inference_method
15
```

```
16  class Gibbs_sampling(Sampling_inference_method):
17      """The class that queries Graphical Models using Gibbs Sampling.
18
19      bn is a graphical model (e.g., a belief network) to query
20      """
21      def __init__(self,bn=None):
22          self.bn = bn
23          self.label = "Gibbs Sampling"
24
25      def query(self, qvar, obs={}, number_samples=1000, burn_in=100, sample_order=None):
26          """computes P(qvar|obs) where
27          qvar is a variable.
28          obs is a variable:value dictionary.
29          sample_order is a list of non-observed variables in order.
30          """
31          counts = {val:0 for val in qvar.domain}
32          if sample_order is not None:
33              variables = sample_order
34          else:
35              variables = [v for v in self.bn.variables if v not in obs]
36          var_to_factors = {v:set() for v in self.bn.variables}
37          for fac in self.bn.factors:
38              for var in fac.variables:
39                  var_to_factors[var].add(fac)
40          sample = {var:random.choice(var.domain) for var in variables}
41          self.display(2,"Sample:",sample)
42          sample.update(obs)
43          for i in range(burn_in + number_samples):
44              if sample_order == None:
45                  random.shuffle(variables)
46              for var in variables:
47                  # get probability distribution of var given its neighbours
48                  vardist = {val:1 for val in var.domain}
49                  for val in var.domain:
50                      sample[var] = val
51                      for fac in var_to_factors[var]: # Markov blanket
52                          vardist[val] *= fac.get_value(sample)
53                  sample[var] = sample_one(vardist)
54              if i >= burn_in:
55                  counts[sample[qvar]] +=1
56          tot = sum(counts.values())
57          return counts, {c:v/tot for (c,v) in counts.items()}
58
59  from probGraphicalModels import bn1, A,B,C
60  bn1g = Gibbs_sampling(bn1)
61  ## Inference_method.max_display_level = 2 # detailed tracing for all inference methods
62  bn1g.query(A,{})
63  ## bn1g.query(C,{})
64  ## bn1g.query(A,{C:True})
65  ## bn1g.query(B,{A:True,C:False})
```

```
66
67  from probGraphicalModels import bn2,Al,Fi,Le,Re,Sm,Ta
68  bn2g = Gibbs_sampling(bn2)
69  ## bn2g.query(Ta,{Re:True},number_samples=100000)
```

Exercise 8.4 Change the code so that it can have multiple query variables. Make the list of query variable be an input to the algorithm, so that the default value is the list of all non-observed variables.

Exercise 8.5 In this algorithm, explain where it computes the probability of a variable given its Markov blanket. Instead of returning the average of the samples for the query variable, it is possible to return the average estimate of the probability of the query variable given its Markov blanket. Does this converge to the same answer as the given code? Does it converge faster, slower, or the same?

8.7 Hidden Markov Models

This code for hidden Markov models is independent of the graphical models code, to keep it simple. Section 8.8 gives code that models hidden Markov models, and more generally, dynamic belief networks, using the graphical models code.

This HMM code assumes there are multiple Boolean observation variables that depend on the current state and are independent of each other given the state.

```
_____ probHMM.py — Hidden Markov Model _____
11  import random
12  from probStochSim import sample_one, sample_multiple
13
14  class HMM(object):
15      def __init__(self, states, obsvars,pobs,trans,indist):
16          """A hidden Markov model.
17          states - set of states
18          obsvars - set of observation variables
19          pobs - probability of observations, pobs[i][s] is P(Obs_i=True | State=s)
20          trans - transition probability - trans[i][j] gives P(State=j | State=i)
21          indist - initial distribution - indist[s] is P(State_0 = s)
22          """
23          self.states = states
24          self.obsvars = obsvars
25          self.pobs = pobs
26          self.trans = trans
27          self.indist = indist
```

Consider the following example. Suppose you want to unobtrusively keep track of an animal in a triangular enclosure using sound. Suppose you have 3 microphones that provide unreliable (noisy) binary information at each time step. The animal is either close to one of the 3 points of the triangle or in the middle of the triangle.

```
_____ probHMM.py — (continued) _____
29  # state
30  #        0=middle, 1,2,3 are corners
31  states1 = {'middle', 'c1', 'c2', 'c3'} # states
32  obs1 = {'m1','m2','m3'} # microphones
```

The observation model is as follows. If the animal is in a corner, it will be detected by the microphone at that corner with probability 0.6, and will be independently detected by each of the other microphones with a probability of 0.1. If the animal is in the middle, it will be detected by each microphone with a probability of 0.4.

```
_____ probHMM.py — (continued) _____
34  # pobs gives the observation model:
35  #pobs[mi][state] is P(mi=on | state)
36  closeMic=0.6; farMic=0.1; midMic=0.4
37  pobs1 = {'m1':{'middle':midMic, 'c1':closeMic, 'c2':farMic, 'c3':farMic}, # mic 1
38          'm2':{'middle':midMic, 'c1':farMic, 'c2':closeMic, 'c3':farMic}, # mic 2
39          'm3':{'middle':midMic, 'c1':farMic, 'c2':farMic, 'c3':closeMic}} # mic 3
```

The transition model is as follows: If the animal is in a corner it stays in the same corner with probability 0.80, goes to the middle with probability 0.1 or goes to one of the other corners with probability 0.05 each. If it is in the middle, it stays in the middle with probability 0.7, otherwise it moves to one the corners, each with probability 0.1.

```
_____ probHMM.py — (continued) _____
41  # trans specifies the dynamics
42  # trans[i] is the distribution over states resulting from state i
43  # trans[i][j] gives P(S=j | S=i)
44  sm=0.7; mmc=0.1              # transition probabilities when in middle
45  sc=0.8; mcm=0.1; mcc=0.05   # transition probabilities when in a corner
46  trans1 = {'middle':{'middle':sm, 'c1':mmc, 'c2':mmc, 'c3':mmc}, # was in middle
47          'c1':{'middle':mcm, 'c1':sc, 'c2':mcc, 'c3':mcc}, # was in corner 1
48          'c2':{'middle':mcm, 'c1':mcc, 'c2':sc, 'c3':mcc}, # was in corner 2
49          'c3':{'middle':mcm, 'c1':mcc, 'c2':mcc, 'c3':sc}} # was in corner 3
```

Initially the animal is in one of the four states, with equal probability.

```
_____ probHMM.py — (continued) _____
51  # initially we have a uniform distribution over the animal's state
52  indist1 = {st:1.0/len(states1) for st in states1}
53
54  hmm1 = HMM(states1, obs1, pobs1, trans1, indist1)
```

8.7.1   Exact Filtering for HMMs

A *HMM_VE_filter* has a current state distribution which can be updated by observing or by advancing to the next time.

_____ probHMM.py — (continued) _____

```
56  from utilities import Displayable
57
58  class HMM_VE_filter(Displayable):
59      def __init__(self,hmm):
60          self.hmm = hmm
61          self.state_dist = hmm.indist
62
63      def filter(self, obsseq):
64          """updates and returns the state distribution following the sequence of
65          observations in obsseq using variable elimination.
66
67          Note that it first advances time.
68          This is what is required if it is called sequentially.
69          If that is not what is wanted initially, do an observe first.
70          """
71          for obs in obsseq:
72              self.advance()     # advance time
73              self.observe(obs) # observe
74          return self.state_dist
75
76      def observe(self, obs):
77          """updates state conditioned on observations.
78          obs is a list of values for each observation variable"""
79          for i in self.hmm.obsvars:
80              self.state_dist = {st:self.state_dist[st]*(self.hmm.pobs[i][st]
81                                              if obs[i] else (1-self.hmm.pobs[i][st]))
82                          for st in self.hmm.states}
83          norm = sum(self.state_dist.values()) # normalizing constant
84          self.state_dist = {st:self.state_dist[st]/norm for st in self.hmm.states}
85          self.display(2,"After observing",obs,"state distribution:",self.state_dist)
86
87      def advance(self):
88          """advance to the next time"""
89          nextstate = {st:0.0 for st in self.hmm.states} # distribution over next states
90          for j in self.hmm.states:      # j ranges over next states
91              for i in self.hmm.states: # i ranges over previous states
92                  nextstate[j] += self.hmm.trans[i][j]*self.state_dist[i]
93          self.state_dist = nextstate
```

The following are some queries for *hmm*1.

_____ probHMM.py — (continued) _____

```
95  hmm1f1 = HMM_VE_filter(hmm1)
96  # hmm1f1.filter([{'m1':0, 'm2':1, 'm3':1}, {'m1':1, 'm2':0, 'm3':1}])
97  ## HMM_VE_filter.max_display_level = 2 # show more detail in displaying
98  # hmm1f2 = HMM_VE_filter(hmm1)
99  # hmm1f2.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':1, 'm3':0}, {'m1':1, 'm2':0, 'm3':0},
100 #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0},
101 #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':1}, {'m1':0, 'm2':0, 'm3':1},
102 #                {'m1':0, 'm2':0, 'm3':1}])
```

```
103   # hmm1f3 = HMM_VE_filter(hmm1)
104   # hmm1f3.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':1, 'm2':0, 'm3':0}, {'m1':
105
106   # How do the following differ in the resulting state distribution?
107   # Note they start the same, but have different initial observations.
108   ## HMM_VE_filter.max_display_level = 1 # show less detail in displaying
109   # for i in range(100): hmm1f1.advance()
110   # hmm1f1.state_dist
111   # for i in range(100): hmm1f3.advance()
112   # hmm1f3.state_dist
```

Exercise 8.6 The localization example in the book is a controlled HMM, where there is a given action at each time and the transition depends on the action. Change the code to allow for controlled HMMs. Hint: the action only influences the state transition.

Exercise 8.7 The representation assumes that there are a list of Boolean observations. Extend the representation so that the each observation variable can have multiple discrete values. You need to choose a representation for the model, and change the algorithm.

8.7.2   Particle Filtering for HMMs

In this implementation a particle is just a state. If you want to do some form of smooting, a particle should probably be a history of states. This maintains, *particles*, an array of states, *weights* an array of (non-negative) real numbers, such that *weights*[*i*] is the weight of *particles*[*i*].

```
                          probHMM.py — (continued)
113   from utilities import Displayable
114   from probStochSim import resample
115
116   class HMM_particle_filter(Displayable):
117       def __init__(self,hmm,number_particles=1000):
118           self.hmm = hmm
119           self.particles = [sample_one(hmm.indist)
120                             for i in range(number_particles)]
121           self.weights = [1 for i in range(number_particles)]
122
123       def filter(self, obsseq):
124           """returns the state distribution following the sequence of
125           observations in obsseq using particle filtering.
126
127           Note that it first advances time.
128           This is what is required if it is called after previous filtering.
129           If that is not what is wanted initially, do an observe first.
130           """
131           for obs in obsseq:
132               self.advance()   # advance time
133               self.observe(obs) # observe
```

```
134                  self.resample_particles()
135              self.display(2,"After observing", str(obs),
136                          "state distribution:", self.histogram(self.particles))
137          self.display(1,"Final state distribution:", self.histogram(self.particles))
138          return self.histogram(self.particles)
139
140      def advance(self):
141          """advance to the next time.
142          This assumes that all of the weights are 1."""
143          self.particles = [sample_one(self.hmm.trans[st])
144                          for st in self.particles]
145
146      def observe(self, obs):
147          """reweight the particles to incorporate observations obs"""
148          for i in range(len(self.particles)):
149              for obv in obs:
150                  if obs[obv]:
151                      self.weights[i] *= self.hmm.pobs[obv][self.particles[i]]
152                  else:
153                      self.weights[i] *= 1-self.hmm.pobs[obv][self.particles[i]]
154
155      def histogram(self, particles):
156          """returns list of the probability of each state as represented by
157          the particles"""
158          tot=0
159          hist = {st: 0.0 for st in self.hmm.states}
160          for (st,wt) in zip(self.particles,self.weights):
161              hist[st]+=wt
162              tot += wt
163          return {st:hist[st]/tot for st in hist}
164
165      def resample_particles(self):
166          """resamples to give a new set of particles."""
167          self.particles = resample(self.particles, self.weights, len(self.particles))
168          self.weights = [1] * len(self.particles)
```

The following are some queries for *hmm*1.

```
                          probHMM.py — (continued)
170  hmm1pf1 = HMM_particle_filter(hmm1)
171  # HMM_particle_filter.max_display_level = 2 # show each step
172  # hmm1pf1.filter([{'m1':0, 'm2':1, 'm3':1}, {'m1':1, 'm2':0, 'm3':1}])
173  # hmm1pf2 = HMM_particle_filter(hmm1)
174  # hmm1pf2.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':1, 'm3':0}, {'m1':1, 'm2':0, 'm3':0},
175  #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0},
176  #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':1}, {'m1':0, 'm2':0, 'm3':1},
177  #                {'m1':0, 'm2':0, 'm3':1}])
178  # hmm1pf3 = HMM_particle_filter(hmm1)
179  # hmm1pf3.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':1, 'm2':0, 'm3':0}, {'
```

Exercise 8.8 A form of importance sampling can be obtained by not resampling.

Is it better or worse than particle filtering? Hint: you need to think about how they can be compared. Is the comparison different if there are more states than particles?

Exercise 8.9 Extend the particle filtering code to continuous variables and observations. In particular, suppose the state transition is a linear function with Gaussian noise of the previous state, and the observations are linear functions with Gaussian noise of the state. You may need to research how to sample from a Gaussian distribution.

8.7.3 Generating Examples

The following code is useful for generating examples.

```
                          probHMM.py — (continued)
181  def simulate(hmm,horizon):
182      """returns a pair of (state sequence, observation sequence) of length horizon.
183      for each time t, the agent is in state_sequence[t] and
184      observes observation_sequence[t]
185      """
186      state = sample_one(hmm.indist)
187      obsseq=[]
188      stateseq=[]
189      for time in range(horizon):
190          stateseq.append(state)
191          newobs = {obs:sample_one({0:1-hmm.pobs[obs][state],1:hmm.pobs[obs][state]})
192                  for obs in hmm.obsvars}
193          obsseq.append(newobs)
194          state = sample_one(hmm.trans[state])
195      return stateseq,obsseq
196
197  def simobs(hmm,stateseq):
198      """returns observation sequence for the state sequence"""
199      obsseq=[]
200      for state in stateseq:
201          newobs = {obs:sample_one({0:1-hmm.pobs[obs][state],1:hmm.pobs[obs][state]})
202                  for obs in hmm.obsvars}
203          obsseq.append(newobs)
204      return obsseq
205
206  def create_eg(hmm,n):
207      """Create an annotated example for horizon n"""
208      seq,obs = simulate(hmm,n)
209      print("True state sequence:",seq)
210      print("Sequence of observations:\n",obs)
211      hmmfilter = HMM_VE_filter(hmm)
212      dist = hmmfilter.filter(obs)
213      print("Resulting distribution over states:\n",dist)
```

8.8 Dynamic Belief Networks

A dynamic belief network consists of:

- A set of features. A variable is a feature-time pair.

- An initial distribution over the features at time 0. This is a belief network with all variables being time 0 variables.

- A specification of the dynamics. Here we define the how the variables one time depend on variables at that time and the previous time, in such a way that the graph is acyclic.

There are a number of ways that reasoning can be carried out in a DBN, including:

- Rolling out the DBN for some time period, and using standard belief network inference. The latest time that needs to be in the rolled out network is the time of the latest observation or the time of a query (whichever is later). This allows us to observe any variables at any time and query any variables at any time. However, the unrolled Bayesian network may be very large. We also need to construct multiple copies of each feature.

- Just representing the variables "now". In this approach we can observe and query the current variables. We can them move to the next time. This does not allow for arbitrary historical queries (about the past or the future), but can be much simpler.

Here we will implement the second of these.

```
_____probDBN.py — Dynamic belief networks _____
11  from probVariables import Variable
12  from probGraphicalModels import Graphical_model
13  from probFactors import Prob, Factor_rename
14  from probVE import VE
15  from utilities import Displayable
16
17  class DBN_variable(Variable):
18      """A random variable that incorporates
19
20      A variable can have both a name and an index. The index defaults to 1.
21      Equality is true if they are both the name and the index are the same."""
22      def __init__(self,name,domain=[False,True],index=1):
23          Variable.__init__(self,name,domain)
24          self.index = index
25          self.previous = None
26
27      def __lt__(self,other):
28          if self.name != other.name:
29              return self.name<other.name
```

```
30              else:
31                  return self.index<other.index
32
33      def __gt__(self,other):
34          return other<self
35
36      def __str__(self):
37  #          if self.index==1:
38  #              return self.name
39  #          else:
40              return self.name+"_"+str(self.index)
41
42      __repr__ = __str__
43
44  def variable_pair(name,domain=[False,True]):
45      """returns a variable and its predecessor. This is used to define 2-stage DBNs
46
47      If the name is X, it returns the pair of variables X0,X"""
48      var = DBN_variable(name,domain)
49      var0 = DBN_variable(name,domain,index=0)
50      var.previous = var0
51      return var0, var
```

───────────── probDBN.py — (continued) ─────────────

```
53  class DBN(Displayable):
54      """The class of stationary Dynamic Bayesian networks.
55
56      * vars1 is a list of current variables (each must have
57      previous variable).
58      * transition_factors is a list of factors for P(X|parents) where X
59      is a current variable and parents is a list of current or previous variables.
60      * init_factors is a list of factors for P(X|parents) where X is a
61      current variable and parents can only include current variables
62      The graph of transition factors + init factors must be acyclic.
63
64      """
65      def __init__(self,vars1, transition_factors=None, init_factors=None):
66          self.vars1 = vars1
67          self.vars0 = [v.previous for v in vars1]
68          self.transition_factors = transition_factors
69          self.init_factors = init_factors
70          self.var_index = {}      # var_index[v] is the index of variable v
71          for i,v in enumerate(vars1):
72              self.var_index[v]=i
```

Here is a 3 variable DBN:

───────────── probDBN.py — (continued) ─────────────

```
74  A0,A1 = variable_pair("A")
75  B0,B1 = variable_pair("B")
76  C0,C1 = variable_pair("C")
```

```
77
78  # dynamics
79  pc = Prob(C1,[B1,C0],[0.03,0.97,0.38,0.62,0.23,0.77,0.78,0.22])
80  pb = Prob(B1,[A0,A1],[0.5,0.5,0.77,0.23,0.4,0.6,0.83,0.17])
81  pa = Prob(A1,[A0,B0],[0.1,0.9,0.65,0.35,0.3,0.7,0.8,0.2])
82
83  # initial distribution
84  pa0 = Prob(A1,[],[0.9,0.1])
85  pb0 = Prob(B1,[A1],[0.3,0.7,0.8,0.2])
86  pc0 = Prob(C1,[],[0.2,0.8])
87
88  dbn1 = DBN([A1,B1,C1],[pa,pb,pc],[pa0,pb0,pc0])
```

Here is the animal example

──────────────── probDBN.py — (continued) ────────────────
```
90   from probHMM import closeMic, farMic, midMic, sm, mmc, sc, mcm, mcc
91
92   Pos_0,Pos_1 = variable_pair("Position",domain=[0,1,2,3])
93   Mic1_0,Mic1_1 = variable_pair("Mic1")
94   Mic2_0,Mic2_1 = variable_pair("Mic2")
95   Mic3_0,Mic3_1 = variable_pair("Mic3")
96
97   # conditional probabilities - see hmm for the values of sm,mmc, etc
98   ppos = Prob(Pos_1, [Pos_0],
99             [sm, mmc, mmc, mmc, #was in middle
100             mcm, sc, mcc, mcc, #was in corner 1
101             mcm, mcc, sc, mcc, #was in corner 2
102             mcm, mcc, mcc, sc]) #was in corner 3
103  pm1 = Prob(Mic1_1, [Pos_1], [1-midMic, midMic, 1-closeMic, closeMic,
104                         1-farMic, farMic, 1-farMic, farMic])
105  pm2 = Prob(Mic2_1, [Pos_1], [1-midMic, midMic, 1-farMic, farMic,
106                         1-closeMic, closeMic, 1-farMic, farMic])
107  pm3 = Prob(Mic3_1, [Pos_1], [1-midMic, midMic, 1-farMic, farMic,
108                         1-farMic, farMic, 1-closeMic, closeMic])
109  ipos = Prob(Pos_1,[], [0.25, 0.25, 0.25, 0.25])
110  dbn_an =DBN([Pos_1,Mic1_1,Mic2_1,Mic3_1],
111             [ppos, pm1, pm2, pm3],
112             [ipos, pm1, pm2, pm3])
```

──────────────── probDBN.py — (continued) ────────────────
```
114  class DBN_VE_filter(VE):
115      def __init__(self,dbn):
116          self.dbn = dbn
117          self.current_factors = dbn.init_factors
118          self.current_obs = {}
119
120      def observe(self, obs):
121          """updates the current observations with obs.
122          obs is a variable:value dictionary where variable is a current
123          variable.
```

```python
124           """
125           assert all(self.current_obs[var]==obs[var] for var in obs
126                      if var in self.current_obs),"inconsistent current observations"
127           self.current_obs.update(obs)
128
129       def query(self,var):
130           """returns the posterior probability of current variable var"""
131           return VE(Graphical_model(self.dbn.vars1,self.current_factors)).query(var,self.current_obs)
132
133       def advance(self):
134           """advance to the next time"""
135           prev_factors = [self.make_previous(fac) for fac in self.current_factors]
136           prev_obs = {var.previous:val for var,val in self.current_obs.items()}
137           two_stage_factors = prev_factors + self.dbn.transition_factors
138           self.current_factors = self.elim_vars(two_stage_factors,self.dbn.vars0,prev_obs)
139           self.current_obs = {}
140
141       def make_previous(self,fac):
142           """Creates new factor from fac where the current variables in fac
143           are renamed to previous variables.
144           """
145           return Factor_rename(fac, {var.previous:var for var in fac.variables})
146
147       def elim_vars(self,factors, vars, obs):
148           for var in vars:
149               if var in obs:
150                   factors = [self.project_observations(fac,obs) for fac in factors]
151               else:
152                   factors = self.eliminate_var(factors, var)
153           return factors
```

Example queries:

```python
                              ___probDBN.py — (continued) ___
155  df = DBN_VE_filter(dbn1)
156  #df.observe({B1:True}); df.advance(); df.observe({C1:False})
157  #df.query(B1)
158  #df.advance()
159  #df.query(B1)
160  dfa = DBN_VE_filter(dbn_an)
161  # dfa.observe({Mic1_1:0, Mic2_1:1, Mic3_1:1})
162  # dfa.advance()
163  # dfa.observe({Mic1_1:1, Mic2_1:0, Mic3_1:1})
164  # dfa.query(Pos_1)
```

Chapter 9

Planning with Uncertainty

9.1   Decision Networks

The decision network code builds on the representation for belief networks of
Chapter 8.

   We first allow for factors that define the utility. Here the utility is a function
of the variables in *vars*, and the table is a list that enumerates the values as in
Section 8.2.

---
decnNetworks.py — Representations for Decision Networks

```
11  from probGraphicalModels import Graphical_model
12  from probFactors import Factor_stored
13  from probVariables import Variable
14  from probFactors import Prob
15
16  class Utility(Factor_stored):
17      """A factor defined by a utility"""
18      def __init__(self,vars,table):
19          """Creates a factor on vars from the table.
20          The table is ordered according to vars.
21          """
22          Factor_stored.__init__(self,vars,table)
23          assert self.size==len(table),"Table size incorrect "+str(self)
```

A decision variable is a like a random variable with a string name, and a do-
main, which is a list of possible values. The decision variable also includes the
parents, a list of the variables whose value will be known when the decision is
made.

---
decnNetworks.py — (continued)

```
25  class DecisionVariable(Variable):
26      def __init__(self,name,domain,parents):
```

```
27            Variable.__init__(self,name,domain)
28            self.parents = parents
29            self.all_vars = set(parents) | {self}
```

A decision network is a graphical model where the variables can be random variables or decision variables. In the factors we assume there is one utility factor.

```
_____decnNetworks.py — (continued)_____
31  class DecisionNetwork(Graphical_model):
32      def __init__(self,vars=None,factors=None):
33          """vars is a list of variables
34          factors is a list of factors (instances of Prob and Utility)
35          """
36          Graphical_model.__init__(self,vars,factors)
```

VE_DN is variable elimination for decision networks. The method *optimize* is used to optimize all the decisions. Note that *optimize* requires a legal emimination ordering of the random and decision variables, otherwise it will give an exception. (A decision node can only be maximized if the variables that are not its parents have already been eliminated.)

```
_____decnNetworks.py — (continued)_____
38  from probFactors import factor_times, Factor_stored
39  from probVE import VE
40
41  class VE_DN(VE):
42      """Variable Elimination for Decision Networks"""
43      def __init__(self,dn=None):
44          """dn is a decision network"""
45          VE.__init__(self,dn)
46          self.dn = dn
47
48      def optimize(self,elim_order=None,obs={}):
49          if elim_order == None:
50                  elim_order = self.gm.variables
51          policy = []
52          proj_factors = [self.project_observations(fact,obs)
53                          for fact in self.dn.factors]
54          for v in elim_order:
55              if isinstance(v,DecisionVariable):
56                  to_max = [fac for fac in proj_factors
57                          if v in fac.variables and set(fac.variables) <= v.all_vars]
58                  assert len(to_max)==1, "illegal variable order "+str(elim_order)+" at "+str(v)
59                  newFac = Factor_max(v, to_max[0])
60                  policy.append(newFac.decision_fun)
61                  proj_factors = [fac for fac in proj_factors if fac is not to_max[0]]+[newFac]
62                  self.display(2,"maximizing",v,"resulting factor",newFac.brief() )
63                  self.display(3,newFac)
64              else:
65                  proj_factors = self.eliminate_var(proj_factors, v)
```

```
66          assert len(proj_factors)==1,"Should there be only one element of proj_factors?"
67          value = proj_factors[0].get_value({})
68          return value,policy
```

---------------------------------- decnNetworks.py — (continued) ----------------------------------

```
70  class Factor_max(Factor_stored):
71      """A factor obtained by maximizing a variable in a factor.
72      Also builds a decision_function. This is based on Factor_sum.
73      """
74
75      def __init__(self, dvar, factor):
76          """dvar is a decision variable.
77          factor is a factor that contains dvar and only parents of dvar
78          """
79          self.dvar = dvar
80          self.factor = factor
81          vars = [v for v in factor.variables if v is not dvar]
82          Factor_stored.__init__(self,vars,None)
83          self.values = [None]*self.size
84          self.decision_fun = Factor_DF(dvar,vars,[None]*self.size)
85
86      def get_value(self,assignment):
87          """lazy implementation: if saved, return saved value, else compute it"""
88          index = self.assignment_to_index(assignment)
89          if self.values[index]:
90              return self.values[index]
91          else:
92              max_val = float("-inf") # -infinity
93              new_asst = assignment.copy()
94              for elt in self.dvar.domain:
95                  new_asst[self.dvar] = elt
96                  fac_val = self.factor.get_value(new_asst)
97                  if fac_val>max_val:
98                      max_val = fac_val
99                      best_elt = elt
100             self.values[index] = max_val
101             self.decision_fun.values[index] = best_elt
102             return max_val
```

A decision function is a stored factor.

---------------------------------- decnNetworks.py — (continued) ----------------------------------

```
104 class Factor_DF(Factor_stored):
105     """A decision function"""
106     def __init__(self,dvar, vars, values):
107         Factor_stored.__init__(self,vars,values)
108         self.dvar = dvar
109         self.name = str(dvar) # Used in printing
```

The fire decision network of Figure 9.1 is represented as:

Figure 9.1: Fire Decision Network

```
_____ decnNetworks.py — (continued) _____
111  boolean = [False, True]
112  Al = Variable("Alarm", boolean)
113  Fi = Variable("Fire", boolean)
114  Le = Variable("Leaving", boolean)
115  Re = Variable("Report", boolean)
116  Sm = Variable("Smoke", boolean)
117  Ta = Variable("Tamper", boolean)
118  SS = Variable("See Sm", boolean)
119  CS = DecisionVariable("Ch Sm", boolean,{Re})
120  Call = DecisionVariable("Call", boolean,{SS,CS,Re})
121
122  f_ta = Prob(Ta,[],[0.98,0.02])
123  f_fi = Prob(Fi,[],[0.99,0.01])
124  f_sm = Prob(Sm,[Fi],[0.99,0.01,0.1,0.9])
125  f_al = Prob(Al,[Fi,Ta],[0.9999, 0.0001, 0.15, 0.85, 0.01, 0.99, 0.5, 0.5])
126  f_lv = Prob(Le,[Al],[0.999, 0.001, 0.12, 0.88])
127  f_re = Prob(Re,[Le],[0.99, 0.01, 0.25, 0.75])
128  f_ss = Prob(SS,[CS,Sm],[1,0,1,0,1,0,0,1])
129
130  ut = Utility([CS,Fi,Call],[0,-200,-5000,-200,-20,-220,-5020,-220])
131
132  dnf = DecisionNetwork([Ta,Fi,Al,Le,Sm,Call,SS,CS,Re],[f_ta,f_fi,f_sm,f_al,f_lv,f_re,f_ss,ut])
133  # v,p = VE_DN(dnf).optimize()
134  # for df in p: print(df,"\n")
```

The following is the representation of the cheating decision of Figure 9.2. Note
that we keep the names of the variables short (less than 8 characters) so that
tables Python prints look good.

Figure 9.2: Cheating Decision Network

```
                          decnNetworks.py — (continued)
136  grades = ["A","B","C","F"]
137  Wa = Variable("Watched", boolean)
138  CC1 = Variable("Caught1", boolean)
139  CC2 = Variable("Caught2", boolean)
140  Pun = Variable("Punish",["None","Suspension","Recorded"])
141  Gr1 = Variable("Grade_1",grades)
142  Gr2 = Variable("Grade_2",grades)
143  GrF = Variable("Fin_Grd",grades)
144  Ch1 = DecisionVariable("Cheat_1", boolean,set()) #no parents
145  Ch2 = DecisionVariable("Cheat_2", boolean,{Ch1,CC1})
146
147  p_wa = Prob(Wa,[],[0.7, 0.3])
148  p_cc1 = Prob(CC1,[Wa,Ch1],[1.0, 0.0, 0.9, 0.1, 1.0, 0.0, 0.5, 0.5])
149  p_cc2 = Prob(CC2,[Wa,Ch2],[1.0, 0.0, 0.9, 0.1, 1.0, 0.0, 0.5, 0.5])
150  p_pun = Prob(Pun,[CC1,CC2],[1.0, 0.0, 0.0, 0.5, 0.4, 0.1, 0.6, 0.2, 0.2, 0.2, 0.5, 0.3])
151  p_gr1 = Prob(Gr1,[Ch1], [0.2, 0.3, 0.3, 0.2, 0.5, 0.3, 0.2, 0.0])
152  p_gr2 = Prob(Gr2,[Ch2], [0.2, 0.3, 0.3, 0.2, 0.5, 0.25, 0.25, 0.0])
153  p_fg = Prob(GrF,[Gr1,Gr2],
154          [1.0, 0.0, 0.0, 0.0, 0.5, 0.5, 0.0, 0.0, 0.25, 0.5, 0.25, 0.0, 0.25,
155          0.25, 0.25, 0.25, 0.5, 0.5, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.5,
156          0.5, 0.0, 0.0, 0.25, 0.5, 0.25, 0.25, 0.5, 0.25, 0.0, 0.0, 0.5, 0.5,
157          0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.25, 0.75, 0.25, 0.5, 0.25, 0.0,
158          0.0, 0.25, 0.5, 0.25, 0.0, 0.0, 0.25, 0.75, 0.0, 0.0, 0.0, 1.0])
159  utc = Utility([Pun,GrF],[100,90,70,50,40,20,10,0,70,60,40,20])
```

```
160
161   cheat_dn = DecisionNetwork([Pun,CC2,Wa,GrF,Gr2,Gr1,Ch2,CC1,Ch1],
162                             [p_wa, p_cc1, p_cc2, p_pun, p_gr1, p_gr2,p_fg,utc])
163
164   # VE_DN.max_display_level = 3 # if you want to show lots of detail
165   # v,p = VE_DN(cheat_dn).optimize(); print(v)
166   # for df in p: print(df,"\n") # print decision functions
```

9.2 Markov Decision Processes

We will represent a **Markov decision process** (**MDP**) directly, rather than using the variable elimination code, as we did for decision networks.

States and actions are represented as lists of strings. The data structures for transitions, rewards, q-values, etc., use the index of the state or the action. The names of the state with index i is in $states[i]$, and the name of action with index i is in $actions[i]$.

─────────── mdpProblem.py — Representations for Markov Decision Processes ───────────
```python
11   from utilities import argmax
12
13   class MDP(object):
14       def __init__(self, states, actions, trans, reward, discount):
15           """states is a list or tuple of states.
16           actions is a list or tuple of actions
17           trans[s][a][s'] represents P(s'|a,s)
18           reward[s][a] gives the expected reward of doing a in state s
19           discount is a real in the range [0,1]
20           """
21           self.states = states
22           self.actions = actions
23           self.trans = trans
24           self.reward = reward
25           self.discount = discount
26           self.v0 = [0 for s in states] # initial value function
```

2 state partying example:

─────────────────────── mdpExamples.py — MDP Examples ───────────────────────
```python
11   from mdpProblem import MDP
12   #### Partying Decision Example ####
13
14   # States: Healthy Sick
15   # Actions: Relax Party
16
17   # trans[s][a][s'] gives P(s'|a,s)
18   #             Relax        Party
19   trans2 = (((0.95,0.05), (0.7, 0.3)), # Healthy
20            ((0.5,0.5),  (0.1, 0.9))  # Sick
21            )
```

```
22
23  # reward[s][a] gives the expected reward of doing a in state s.
24  reward2 = ((7,10),(0,2))
25
26  healthy2 = MDP(['Healthy','Sick'], ['Relax','Party'], trans2, reward2, discount=0.8)
```

Tiny game from Example 11.7 and Figure 11.8 of Poole and Mackworth, 2010:

———————— mdpExamples.py — (continued) ————————

```
28  ## Tiny Game from Example 11.7 and Figure 11.8 of Poole and Mackworth, 2010 #
29
30  # actions        up              right         upC           left
31  transt = (((0.1,0.1,0.8,0,0,0), (0,1,0,0,0,0), (0,0,1,0,0,0), (1,0,0,0,0,0)), #s0
32          ((0.1,0.1,0,0.8,0,0), (0,1,0,0,0,0), (0,0,0,1,0,0), (1,0,0,0,0,0)),
    #s1
33          ((0,0,0.1,0.1,0.8,0), (0,0,0,1,0,0), (0,0,0,0,1,0), (0,0,1,0,0,0)),
    #s2
34          ((0,0,0.1,0.1,0,0.8), (0,0,0,1,0,0), (0,0,0,0,0,1), (0,0,1,0,0,0)),
    #s3
35          ((0.1,0,0,0,0.8,0.1), (0,0,0,0,0,1), (0,0,0,0,1,0), (1,0,0,0,0,0)),
    #s4
36          ((0,0,0,0,0.1,0.9), (0,0,0,0,0,1), (0,0,0,0,0,1), (0,0,0,0,1,0)) ) #s5
37
38  # actions    up  rt upC left
39  rewardt = ((-0.1,  0, -1,   -1),   #s0
40            (-0.1, -1, -2,    0),   #s1
41            (-10,   0, -1, -100),   #s2
42            (-0.1, -1, -1,    0),   #s3
43            (-1,    0, -2,   10),   #s4
44            (-1,   -1, -2,    0))   #s5
45
46  mdpt = MDP(['s0','s1','s2','s3','s4','s5'], # states
47            ['up', 'right', 'upC', 'left'], # actions
48            transt, rewardt, discount=0.9)
```

9.3   Value Iteration

This implements value iteration, storing V.

This uses indexes of the states and actions (not the names). A value function is list, v, such that $v[s]$ is the value for state with index s. Similarly a policy *pi* is represented as a list where $pi[s]$, where s is the index of a state, returns the index of the action.

———————— mdpProblem.py — (continued) ————————

```
28      def vi1(self,v):
29          """carry out one iteration of value iteration and
30          returns a value function (a list of a value for each state).
31          v is the previous value function.
```

```
32          """
33          return [max([self.reward[s][a]+self.discount*product(self.trans[s][a],v)
34                      for a in range(len(self.actions))])
35                  for s in range(len(self.states))]
36
37      def vi(self,v0,n):
38          """carries out n iterations of value iteration starting with value v0.
39
40          Returns a value function
41          """
42          val = self.v0
43          for i in range(n):
44              val= self.vi1(val)
45          return val
46
47      def policy(self,v):
48          """returns an optimal policy assuming the next value function is v
49             v is a list of values for each state
50             returns a list of the indexes of optimal actions for each state
51          """
52          return [argmax(enumerate([self.reward[s][a]+self.discount*product(self.trans[s][a],v)
53                             for a in range(len(self.actions))]))
54                  for s in range(len(self.states))]
55
56      def q(self,v):
57          """returns the one-step-lookahead q-value assuming the next value function is v
58          v is a list of values for each state
59          returns a list of q values for each state. so that q[s][a] represents Q(s,a)
60          """
61          return [[self.reward[s][a]+self.discount*product(self.trans[s][a],v)
62                   for a in range(len(self.actions))]
63                  for s in range(len(self.states))]
```

—————————————— mdpProblem.py — (continued) ——————————————

```
65  def product(l1,l2):
66      """returns the dot product of l1 and l2"""
67      return sum([i1*i2 for (i1,i2) in zip(l1,l2)])
```

The following gives a trace for the examples:

—————————————— mdpExamples.py — (continued) ——————————————

```
50  def trace(mdp,numiter):
51      print("Q values are shown as",[[st+"_"+ac for ac in mdp.actions] for st in mdp.states])
52      print("One step lookahead Q-values:")
53      print(mdp.q(mdp.v0))
54      print("Values are for the states:", mdp.states)
55      print("One step lookahead values:")
56      print(mdp.vi(mdp.v0,1))
57      print("Two step lookahead Q-values:")
58      print(mdp.q(mdp.vi(mdp.v0,1)))
59      print("Two step lookahead values:")
```

```
60      print(mdp.vi(mdp.v0,2))
61      vfin = mdp.vi(mdp.v0,numiter)
62      print("After",numiter,"iterations, values:")
63      print(vfin)
64      print("After",numiter,"iterations, Q-values:")
65      print(mdp.q(vfin))
66      print("After",numiter,"iterations, Policy:",
67          [st+"->"+mdp.actions[act] for (st,act) in zip(mdp.states ,mdp.policy(vfin))])
68
69  # Try the following:
70  # trace(healthy2,10)
```

Exercise 9.1 Implement value iteration that stores the *Q*-values rather than the *V*-values. Does it work better than storing *V*? (What might better mean?)

Exercise 9.2 Implement asynchronous value iteration. Try a number of different ways to choose the states and actions to update (e.g., sweeping through the state-action pairs, choosing them at random). Note that the best way may be to determine which states have had their Q-values change the most, and then update the previous ones, but that is not so straightforward to implement, because you need to find those previous states.

Chapter 10

Learning with Uncertainty

10.1   K-means

The k-means learner maintains two lists that suffice as sufficient statistics to classify examples, and to learn the classification:

- *class_counts* is a list such that *class_counts*[*c*] is the number of examples in the training set with *class* = *c*.

- *feature_sum* is a list such that *feature_sum*[*i*][*c*] is sum of the values for the *i*'th feature *i* for members of class *c*. The average value of the *i*th feature in class *i* is

$$\frac{feature\_sum[i][c]}{class\_counts[c]}$$

The class is initialized by randomly assigning examples to classes, and updating the statistics for *class_counts* and *feature_sum*.

```
_____ learnKMeans.py — k-means learning _____
11  from learnProblem import Data_set, Learner, Data_from_file
12  import random
13  import matplotlib.pyplot as plt
14
15  class K_means_learner(Learner):
16      def __init__(self,dataset, num_classes):
17          self.dataset = dataset
18          self.num_classes = num_classes
19          self.random_initialize()
20
21      def random_initialize(self):
```

```
22          # class_counts[c] is the number of examples with class=c
23          self.class_counts = [0]*self.num_classes
24          # feature_sum[i][c] is the sum of the values of feature i for class c
25          self.feature_sum = [[0]*self.num_classes
26                              for feat in self.dataset.input_features]
27          for eg in self.dataset.train:
28              cl = random.randrange(self.num_classes) # assign eg to random class
29              self.class_counts[cl] += 1
30              for (ind,feat) in enumerate(self.dataset.input_features):
31                  self.feature_sum[ind][cl] += feat(eg)
32          self.num_iterations = 0
33          self.display(1,"Initial class counts: ",self.class_counts)
```

The distance from (the mean of) a class to an example is the sum, over all fratures, of the sum-of-squares differences of the class mean and the example value.

_____learnKMeans.py — (continued)_____

```
35      def distance(self,cl,eg):
36          """distance of the eg from the mean of the class"""
37          return sum( (self.class_prediction(ind,cl)-feat(eg))**2
38                          for (ind,feat) in enumerate(self.dataset.input_features))
39
40      def class_prediction(self,feat_ind,cl):
41          """prediction of the class cl on the feature with index feat_ind"""
42          if self.class_counts[cl] == 0:
43              return 0 # there are no examples so we can choose any value
44          else:
45              return self.feature_sum[feat_ind][cl]/self.class_counts[cl]
46
47      def class_of_eg(self,eg):
48          """class to which eg is assigned"""
49          return (min((self.distance(cl,eg),cl)
50                      for cl in range(self.num_classes)))[1]
51              # second element of tuple, which is a class with minimum distance
```

One step of k-means updates the *class_counts* and *feature_sum*. It uses the old values to determine the classes, and so the new values for *class_counts* and *feature_sum*. At the end it determines whether the values of these have changes, and then replaces the old ones with the new ones. It returns an indicator of whether the values are stable (have not changed).

_____learnKMeans.py — (continued)_____

```
53      def k_means_step(self):
54          """Updates the model with one step of k-means.
55          Returns whether the assignment is stable.
56          """
57          new_class_counts = [0]*self.num_classes
58          # feature_sum[i][c] is the sum of the values of feature i for class c
59          new_feature_sum = [[0]*self.num_classes
60                              for feat in self.dataset.input_features]
```

```
61          for eg in self.dataset.train:
62              cl = self.class_of_eg(eg)
63              new_class_counts[cl] += 1
64              for (ind,feat) in enumerate(self.dataset.input_features):
65                  new_feature_sum[ind][cl] += feat(eg)
66          stable = (new_class_counts == self.class_counts) and (self.feature_sum == new_feature_sum)
67          self.class_counts = new_class_counts
68          self.feature_sum = new_feature_sum
69          self.num_iterations += 1
70          return stable
71
72
73      def learn(self,n=100):
74          """do n steps of k-means, or until convergence"""
75          i=0
76          stable = False
77          while i<n and not stable:
78              stable = self.k_means_step()
79              i += 1
80              self.display(1,"Iteration",self.num_iterations,
81                              "class counts: ",self.class_counts," Stable=",stable)
82          return stable
83
84      def show_classes(self):
85          """sorts the data by the class and prints in order.
86          For visualizing small data sets
87          """
88          class_examples = [[] for i in range(self.num_classes)]
89          for eg in self.dataset.train:
90              class_examples[self.class_of_eg(eg)].append(eg)
91          print("Class","Example",sep='\t')
92          for cl in range(self.num_classes):
93              for eg in class_examples[cl]:
94                  print(cl,*eg,sep='\t')
95
96      def plot_error(self, maxstep=20):
97          """Plots the sum-of-suares error as a function of the number of steps"""
98          plt.ion()
99          plt.xlabel("step")
100         plt.ylabel("Ave sum-of-squares error")
101         train_errors = []
102         if self.dataset.test:
103             test_errors = []
104         for i in range(maxstep):
105             self.learn(1)
106             train_errors.append( sum(self.distance(self.class_of_eg(eg),eg)
107                                         for eg in self.dataset.train)
108                             /len(self.dataset.train))
109             if self.dataset.test:
110                 test_errors.append( sum(self.distance(self.class_of_eg(eg),eg)
```

```
111                                          for eg in self.dataset.test)
112                                /len(self.dataset.test))
113           plt.plot(range(1,maxstep+1),train_errors,
114                 label=str(self.num_classes)+" classes. Training set")
115           if self.dataset.test:
116               plt.plot(range(1,maxstep+1),test_errors,
117                     label=str(self.num_classes)+" classes. Test set")
118           plt.legend()
119           plt.draw()
120
121   %data = Data_from_file('data/emdata1.csv', num_train=10, target_index=2000) % trivial example
122   data = Data_from_file('data/emdata2.csv', num_train=10, target_index=2000)
123   %data = Data_from_file('data/emdata0.csv', num_train=14, target_index=2000) % example from textbook
124   kml = K_means_learner(data,2)
125   num_iter=4
126   print("Class assignment after",num_iter,"iterations:")
127   kml.learn(num_iter); kml.show_classes()
128
129   # Plot the error
130   # km2=K_means_learner(data,2); km2.plot_error(20) # 2 classes
131   # km3=K_means_learner(data,3); km3.plot_error(20) # 3 classes
132   # km13=K_means_learner(data,13); km13.plot_error(20) # 13 classes
133
134   # data = Data_from_file('data/carbool.csv', target_index=2000,boolean_features=True)
135   # kml = K_means_learner(data,3)
136   # kml.learn(20); kml.show_classes()
137   # km3=K_means_learner(data,3); km3.plot_error(20) # 3 classes
138   # km3=K_means_learner(data,30); km3.plot_error(20) # 30 classes
```

Exercise 10.1 Change *boolean_features = True* flag to allow for numerical features. K-means assumes the features are numerical, so we want to make non-numerical features into numerical features (using characteristic functions) but we probably don't want to change numerical features into Boolean.

Exercise 10.2 If there are many classes, some of the classes can become empty (e.g., try 100 classes with carbool.csv). Implement a way to put some examples into a class, if possible. Two ideas are:

(a) Initialize the classes with actual examples, so that the classes will not start empty. (Do the classes become empty?)

(b) In *class_prediction*, we test whether the code is empty, and make a prediction of 0 for an empty class. It is possible to make a different prediction to "steal" an example (but you should make sure that a class has a consistent value for each feature in a loop).

Make your own suggestions, and compare it with the original, and whichever of these you think may work better.

10.2 EM

In the following definition, a class, *c*, is a integer in range $[0, num\_classes)$. *i* is an index of a feature, so *feat*[*i*] is the *i*th feature, and a feature is a function from tuples to values. *val* is a value of a feature.

A model consists of 2 lists, which form the sufficient statistics:

- *class_counts* is a list such that *class_counts*[*c*] is the number of tuples with *class* = *c*, where each tuple is weighted by its probability, i.e.,

$$class\_counts[c] = \sum_{t:class(t)=c} P(t)$$

- *feature_counts* is a list such that *feature_counts*[*i*][*val*][*c*] is the weighted count of the number of tuples *t* with *feat*[*i*](*t*) = *val* and *class*(*t*) = *c*, each tuple is weighted by its probability, i.e.,

$$feature\_counts[i][val][c] = \sum_{t:feat[i](t)=val \text{ and} class(t)=c} P(t)$$

```
_____learnEM.py — EM Learning_____
11  from learnProblem import Data_set, Learner, Data_from_file
12  import random
13  import math
14  import matplotlib.pyplot as plt
15
16  class EM_learner(Learner):
17      def __init__(self,dataset, num_classes):
18          self.dataset = dataset
19          self.num_classes = num_classes
20          self.class_counts = None
21          self.feature_counts = None
```

The function *em_step* goes though the training examples, and updates these counts. The first time it is run, when there is no model, it uses random distributions.

```
_____learnEM.py — (continued)_____
23      def em_step(self, orig_class_counts, orig_feature_counts):
24          """updates the model."""
25          class_counts = [0]*self.num_classes
26          feature_counts = [{val:[0]*self.num_classes
27                             for val in feat.frange}
28                             for feat in self.dataset.input_features]
29          for tple in self.dataset.train:
30              if orig_class_counts: # a model exists
31                  tpl_class_dist = self.prob(tple, orig_class_counts, orig_feature_counts)
32              else:                  # initially, with no model, return a random distribution
```

```
33              tpl_class_dist = random_dist(self.num_classes)
34          for cl in range(self.num_classes):
35              class_counts[cl] += tpl_class_dist[cl]
36              for (ind,feat) in enumerate(self.dataset.input_features):
37                  feature_counts[ind][feat(tple)][cl] += tpl_class_dist[cl]
38      return class_counts, feature_counts
```

prob computes the probability of a class for a tuple, given the current statistics.

$$P(c \mid tple) \propto P(c) * \prod_i P(X_i{=}tple(i) \mid c)$$

$$= \frac{class\_counts[c]}{len(self.dataset)} * \prod_i \frac{feature\_counts[i][feat_i(tple)][c]}{class\_counts[c]}$$

$len(self.dataset)$ is a constant (independent of c). $class\_counts[c]$ can be taken out of the product, but needs to be raised to the power of the number of features, and one of them cancels.

_____learnEM.py — (continued)_____

```
40      def prob(self,tple,class_counts,feature_counts):
41          """returns a distribution over the classes for the original tuple in the current model
42          """
43          feats = self.dataset.input_features
44          unnorm = [prod(feature_counts[i][feat(tple)][c]
45                          for (i,feat) in enumerate(feats))/(class_counts[c]**(len(feats)-1))
46                      for c in range(self.num_classes)]
47          thesum = sum(unnorm)
48          return [un/thesum for un in unnorm]
```

learn does n steps of EM:

_____learnEM.py — (continued)_____

```
50      def learn(self,n):
51          """do n steps of em"""
52          for i in range(n):
53              self.class_counts,self.feature_counts = self.em_step(self.class_counts,
54                                                  self.feature_counts)
```

The following is for visualizing the classes. It prints the dataset ordered by the probability of class c.

_____learnEM.py — (continued)_____

```
56      def show_class(self,c):
57          """sorts the data by the class and prints in order.
58          For visualizing small data sets
59          """
60          sorted_data = sorted((self.prob(tpl,self.class_counts,self.feature_counts)[c],
61                          ind,   # preserve ordering for equal probabilities
62                          tpl)
63                      for (ind,tpl) in enumerate(self.dataset.train))
64          for cc,r,tpl in sorted_data:
65              print(cc,*tpl,sep='\t')
```

The following are for evaluating the classes.

The probability of a tuple can be evaluated by marginalizing over the classes:

$$P(tple) = \sum_c P(c) * \prod_i P(X_i{=}tple(i) \mid c)$$

$$= \sum_c \frac{cc[c]}{len(self.dataset)} * \prod_i \frac{fc[i][feat_i(tple)][c]}{cc[c]}$$

where cc is the class count and fc is feature count. $len(self.dataset)$ can be distributed out of the sum, and $cc[c]$ can be taken out of the product:

$$= \frac{1}{len(self.dataset)} \sum_c \frac{1}{cc[c]^{\#feats-1}} * \prod_i fc[i][feat_i(tple)][c]$$

Given the probability of each tuple, we can evaluate the logloss, as the negative of the log probability:

```python
_____learnEM.py — (continued)_____
67      def logloss(self,tple):
68          """returns the logloss of the prediction on tple, which is -log(P(tple))
69          based on the current class counts and feature counts
70          """
71          feats = self.dataset.input_features
72          res = 0
73          cc = self.class_counts
74          fc = self.feature_counts
75          for c in range(self.num_classes):
76              res += prod(fc[i][feat(tple)][c]
77                          for (i,feat) in enumerate(feats))/(cc[c]**(len(feats)-1))
78          if res>0:
79              return -math.log2(res/len(self.dataset.train))
80          else:
81              return float("inf") #infinity
82
83      def plot_error(self, maxstep=20):
84          """Plots the logloss error as a function of the number of steps"""
85          plt.ion()
86          plt.xlabel("step")
87          plt.ylabel("Ave Logloss (bits)")
88          train_errors = []
89          if self.dataset.test:
90              test_errors = []
91          for i in range(maxstep):
92              self.learn(1)
93              train_errors.append( sum(self.logloss(tple) for tple in self.dataset.train)
94                              /len(self.dataset.train))
95              if self.dataset.test:
96                  test_errors.append( sum(self.logloss(tple) for tple in self.dataset.test)
97                              /len(self.dataset.test))
98          plt.plot(range(1,maxstep+1),train_errors,
```

```
99                          label=str(self.num_classes)+" classes. Training set")
100             if self.dataset.test:
101                 plt.plot(range(1,maxstep+1),test_errors,
102                         label=str(self.num_classes)+" classes. Test set")
103             plt.legend()
104             plt.draw()
105
106  def prod(L):
107      """returns the product of the elements of L"""
108      res = 1
109      for e in L:
110          res *= e
111      return res
112
113  def random_dist(k):
114      """generate k random numbers that sum to 1"""
115      res = [random.random() for i in range(k)]
116      s = sum(res)
117      return [v/s for v in res]
118
119  data = Data_from_file('data/emdata2.csv', num_train=10, target_index=2000)
120  eml = EM_learner(data,2)
121  num_iter=2
122  print("Class assignment after",num_iter,"iterations:")
123  eml.learn(num_iter); eml.show_class(0)
124
125  # Plot the error
126  # em2=EM_learner(data,2); em2.plot_error(40) # 2 classes
127  # em3=EM_learner(data,3); em3.plot_error(40) # 3 classes
128  # em13=EM_learner(data,13); em13.plot_error(40) # 13 classes
129
130  # data = Data_from_file('data/carbool.csv', target_index=2000,boolean_features=False)
131  # [f.frange for f in data.input_features]
132  # eml = EM_learner(data,3)
133  # eml.learn(20); eml.show_class(0)
134  # em3=EM_learner(data,3); em3.plot_error(60) # 3 classes
135  # em3=EM_learner(data,30); em3.plot_error(60) # 30 classes
```

Exercise 10.3 For the EM data, where there are naturally 2 classes, 3 classes does better on the training set after a while than 2 classes, but worse on the test set. Explain why. Hint: look what the 3 classes are. Use "em3.show_class(i)" for each of the classes $i \in [0, 3)$.

Exercise 10.4 Write code to plot the logloss as a function of the number of classes (from 1 to say 15) for a fixed number of iterations. (From the experience with the existing code, think about how many iterations is appropriate.)

Chapter 11

Multiagent Systems

11.1 Minimax

Here we consider two-player zero-sum games. Here a player only wins when another player loses. This can be modeled as where there is a single utility which one agent (the maximizing agent) is trying minimize and the other agent (the minimizing agent) is trying to minimize.

11.1.1 Creating a two-player game

_____masProblem.py — A Multiagent Problem _____
```
11  from utilities import Displayable
12
13  class Node(Displayable):
14      """A node in a search tree. It has a
15      name a string
16      isMax is True if it is a maximizing node, otherwise it is minimizing node
17      children is the list of children
18      value is what it evaluates to if it is a leaf.
19      """
20      def __init__(self, name, isMax, value, children):
21          self.name = name
22          self.isMax = isMax
23          self.value = value
24          self.allchildren = children
25
26      def isLeaf(self):
27          """returns true of this is a leaf node"""
28          return self.allchildren is None
29
```

```
30    def children(self):
31        """returns the list of all children."""
32        return self.allchildren
33
34    def evaluate(self):
35        """returns the evaluation for this node if it is a leaf"""
36        return self.value
```

The following gives the tree from Figure 10.5 of the book. Note how 888 is used as a value here, but never appears in the trace.

_____masProblem.py — (continued) _____

```
38  fig10_5 = Node("a",True,None, [
39            Node("b",False,None, [
40              Node("d",True,None, [
41                Node("h",False,None, [
42                    Node("h1",True,7,None),
43                    Node("h2",True,9,None)]),
44                Node("i",False,None, [
45                    Node("i1",True,6,None),
46                    Node("i2",True,888,None)])]),
47              Node("e",True,None, [
48                Node("j",False,None, [
49                    Node("j1",True,11,None),
50                    Node("j2",True,12,None)]),
51                Node("k",False,None, [
52                    Node("k1",True,888,None),
53                    Node("k2",True,888,None)])])]),
54          Node("c",False,None, [
55            Node("f",True,None, [
56              Node("l",False,None, [
57                  Node("l1",True,5,None),
58                  Node("l2",True,888,None)]),
59              Node("m",False,None, [
60                  Node("m1",True,4,None),
61                  Node("m2",True,888,None)])]),
62            Node("g",True,None, [
63              Node("n",False,None, [
64                  Node("n1",True,888,None),
65                  Node("n2",True,888,None)]),
66              Node("o",False,None, [
67                  Node("o1",True,888,None),
68                  Node("o2",True,888,None)])])])])
```

The following is a representation of a **magic-sum game**, where players take turns picking a number in the range $[1, 9]$, and the first player to have 3 numbers that sum to 15 wins. Note that this is a syntactic variant of **tic-tac-toe** or **naughts and crosses**. To see this, consider the numbers on a **magic square** (Figure 11.1); 3 numbers that add to 15 correspond exactly to the winning positions of tic-tac-toe played on the magic square.

Note that we do not remove symmetries. (What are the symmetries? How

6	1	8
7	5	3
2	9	4

Figure 11.1: Magic Square

do the symmetries of tic-tac-toe translate here?)

```
                                ____masProblem.py — (continued)____
70
71  class Magic_sum(Node):
72      def __init__(self, xmove=True, last_move=None,
73                  available=[1,2,3,4,5,6,7,8,9], x=[], o=[]):
74          """This is a node in the search for the magic-sum game.
75          xmove is True if the next move belongs to X.
76          last_move is the number selected in the last move
77          available is the list of numbers that are available to be chosen
78          x is the list of numbers already chosen by x
79          o is the list of numbers already chosen by o
80          """
81          self.isMax = self.xmove = xmove
82          self.last_move = last_move
83          self.available = available
84          self.x = x
85          self.o = o
86          self.allchildren = None #computed on demand
87          lm = str(last_move)
88          self.name = "start" if not last_move else "o="+lm if xmove else "x="+lm
89
90      def children(self):
91          if self.allchildren is None:
92              if self.xmove:
93                  self.allchildren = [
94                      Magic_sum(xmove = not self.xmove,
95                              last_move = sel,
96                              available = [e for e in self.available if e is not sel],
97                              x = self.x+[sel],
98                              o = self.o)
99                          for sel in self.available]
100             else:
101                 self.allchildren = [
102                     Magic_sum(xmove = not self.xmove,
103                             last_move = sel,
104                             available = [e for e in self.available if e is not sel],
105                             x = self.x,
106                             o = self.o+[sel])
107                         for sel in self.available]
108         return self.allchildren
109
```

```
110     def isLeaf(self):
111         """A leaf has no numbers available or is a win for one of the players.
112         We only need to check for a win for o if it is currently x's turn,
113         and only check for a win for x if it is o's turn (otherwise it would
114         have been a win earlier).
115         """
116         return (self.available == [] or
117                 (sum_to_15(self.last_move,self.o)
118                  if self.xmove
119                  else sum_to_15(self.last_move,self.x)))
120
121     def evaluate(self):
122         if self.xmove and sum_to_15(self.last_move,self.o):
123             return -1
124         elif not self.xmove and sum_to_15(self.last_move,self.x):
125             return 1
126         else:
127             return 0
128
129 def sum_to_15(last,selected):
130     """is true if last, toegether with two other elements of selected sum to 15.
131     """
132     return any(last+a+b == 15
133               for a in selected if a != last
134               for b in selected if b != last and b != a)
```

11.1.2   Minimax and α-β Pruning

This is a naive depth-first **minimax algorithm**:

```
_____masMiniMax.py — Minimax search with alpha-beta pruning _____
11 def minimax(node):
12     """returns the value of node, and a best path for the agents
13     """
14     if node.isLeaf():
15         return node.evaluate(),None
16     elif node.isMax:
17         max_score = -999
18         max_path = None
19         for C in node.children():
20             score,path = minimax(C,depth+1)
21             if score > max_score:
22                 max_score = score
23                 max_path = C.name,path
24         return max_score,max_path
25     else:
26         min_score = 999
27         min_path = None
28         for C in node.children():
29             score,path = minimax(C,depth+1)
30             if score < min_score:
```

```
31            min_score = score
32            min_path = C.name,path
33      return min_score,min_path
```

The following is a depth-first minimax with α-β **pruning**. It returns the value for a node as well as a best path for the agents.

———————————— masMiniMax.py — (continued) ————————————
```
35  def minimax_alpha_beta(node,alpha,beta,depth=0):
36      """node is a Node, alpha and beta are cutoffs, depth is the depth
37      returns value, path
38      where path is a sequence of nodes that results in the value"""
39      node.display(2," "*depth,"minimax_alpha_beta(",node.name,", ",alpha, ", ", beta,")")
40      best=None      # only used if it will be pruned
41      if node.isLeaf():
42          node.display(2," "*depth,"returning leaf value",node.evaluate())
43          return node.evaluate(),None
44      elif node.isMax:
45          for C in node.children():
46              score,path = minimax_alpha_beta(C,alpha,beta,depth+1)
47              if score >= beta:  # beta pruning
48                  node.display(2," "*depth,"pruned due to beta=",beta,"C=",C.name)
49                  return score, None
50              if score > alpha:
51                  alpha = score
52                  best = C.name, path
53          node.display(2," "*depth,"returning max alpha",alpha,"best",best)
54          return alpha,best
55      else:
56          for C in node.children():
57              score,path = minimax_alpha_beta(C,alpha,beta,depth+1)
58              if score <= alpha:   # alpha pruning
59                  node.display(2," "*depth,"pruned due to alpha=",alpha,"C=",C.name)
60                  return score, None
61              if score < beta:
62                  beta=score
63                  best = C.name,path
64          node.display(2," "*depth,"returning min beta",beta,"best=",best)
65          return beta,best
```

Testing:

———————————— masMiniMax.py — (continued) ————————————
```
67  from masProblem import fig10_5, Magic_sum, Node
68
69  # Node.max_display_level=2 # print detailed trace
70  # minimax_alpha_beta(fig10_5, -9999, 9999,0)
71  # minimax_alpha_beta(Magic_sum(), -9999, 9999,0)
72
73  #To see how much time alpha-beta pruning can save over minimax, uncomment the following:
74  ## import timeit
75  ## timeit.Timer("minimax(Magic_sum())",setup="from __main__ import minimax, Magic_sum"
```

```
76  ##              ).timeit(number=1)
77  ## trace=False
78  ## timeit.Timer("minimax_alpha_beta(Magic_sum(), -9999, 9999,0)",
79  ##              setup="from __main__ import minimax_alpha_beta, Magic_sum"
80  ##              ).timeit(number=1)
```

Chapter 12

Reinforcement Learning

12.1 Representing Agents and Environments

When the learning agent does an action in the environment, it observes a (*state*, *reward*) pair from the environment. The *state* is the world state; this is the fully observable assumption.

An RL environment implements a *do*(*action*) method that returns a (*state*, *reward*) pair.

_____rlProblem.py — Representations for Reinforcement Learning _____

```
11  import random
12  from utilities import Displayable, flip
13
14  class RL_env(Displayable):
15      def __init__(self,actions,state):
16          self.actions = actions # set of actions
17          self.state = state     # initial state
18
19      def do(self, action):
20          """do action
21          returns state,reward
22          """
23          raise NotImplementedError("RL_env.do") # abstract method
```

Here is the definition of the simple 2-state, 2-action party/relax decision.

_____rlProblem.py — (continued) _____

```
25  class Healthy_env(RL_env):
26      def __init__(self):
27          RL_env.__init__(self,["party","relax"], "healthy")
28
29      def do(self, action):
```

```
30          """updates the state based on the agent doing action.
31          returns state,reward
32          """
33          if self.state=="healthy":
34              if action=="party":
35                  self.state = "healthy" if flip(0.7) else "sick"
36                  reward = 10
37              else: # action=="relax"
38                  self.state = "healthy" if flip(0.95) else "sick"
39                  reward = 7
40          else: # self.state=="sick"
41              if action=="party":
42                  self.state = "healthy" if flip(0.1) else "sick"
43                  reward = 2
44              else:
45                  self.state = "healthy" if flip(0.5) else "sick"
46                  reward = 0
47          return self.state,reward
```

12.1.1   Simulating an environment from an MDP

Given the definition for an MDP (page 172), *Env_from_MDP* takes in an MDP and simulates the environment with those dynamics.

Note that the MDP does not contain enough information to simulate a system, because it loses any dependency between the rewards and the resulting state; here we assume the agent always received the average reward for the state and action.

```
                        ___rlProblem.py — (continued) ___
49  class Env_from_MDP(RL_env):
50      def __init__(self, mdp):
51          initial_state = mdp.states[0]
52          RL_env.__init__(self,mdp.actions, initial_state)
53          self.mdp = mdp
54          self.action_index = {action:index for (index,action) in enumerate(mdp.actions)}
55          self.state_index = {state:index for (index,state) in enumerate(mdp.states)}
56
57      def do(self, action):
58          """updates the state based on the agent doing action.
59          returns state,reward
60          """
61          action_ind = self.action_index[action]
62          state_ind = self.state_index[self.state]
63          self.state = pick_from_dist(self.mdp.trans[state_ind][action_ind], self.mdp.states)
64          reward = self.mdp.reward[state_ind][action_ind]
65          return self.state, reward
66
67  def pick_from_dist(dist,values):
68      """
```

Figure 12.1: Monster game

```
69      e.g. pick_from_dist([0.3,0.5,0.2],['a','b','c']) should pick 'a' with probability 0.3, etc.
70      """
71      ran = random.random()
72      i=0
73      while ran>dist[i]:
74          ran -= dist[i]
75          i += 1
76      return values[i]
```

12.1.2 Simple Game

This is for the game depicted in Figure 12.1.

```
                          __rlSimpleEnv.py — Simple game __
11  import random
12  from utilities import flip
13  from rlProblem import RL_env
14
15  class Simple_game_env(RL_env):
16      xdim = 5
17      ydim = 5
18
19      vwalls = [(0,3), (0,4), (1,4)] # vertical walls right of these locations
20      hwalls = [] # not implemented
21      crashed_reward = -1
22
23      prize_locs = [(0,0), (0,4), (4,0), (4,4)]
24      prize_apears_prob = 0.3
25      prize_reward = 10
26
```

```
27      monster_locs = [(0,1), (1,1), (2,3), (3,1), (4,2)]
28      monster_appears_prob = 0.4
29      monster_reward_when_damaged = -10
30      repair_stations = [(1,4)]
31
32      actions = ["up","down","left","right"]
33
34      def __init__(self):
35          # State:
36          self.x = 2
37          self.y = 2
38          self.damaged = False
39          self.prize = None
40          # Statistics
41          self.number_steps = 0
42          self.total_reward = 0
43          self.min_reward = 0
44          self.min_step = 0
45          self.zero_crossing = 0
46          RL_env.__init__(self, Simple_game_env.actions,
47                       (self.x, self.y, self.damaged, self.prize))
48          self.display(2,"","Step","Tot Rew","Ave Rew",sep="\t")
49
50      def do(self,action):
51          """updates the state based on the agent doing action.
52          returns state,reward
53          """
54          reward = 0.0
55          # A prize can appear:
56          if self.prize is None and flip(self.prize_apears_prob):
57                  self.prize = random.choice(self.prize_locs)
58          # Actions can be noisy
59          if flip(0.4):
60              actual_direction = random.choice(self.actions)
61          else:
62              actual_direction = action
63          # Modeling the actions given the actual direction
64          if actual_direction == "right":
65              if self.x==self.xdim-1 or (self.x,self.y) in self.vwalls:
66                  reward += self.crashed_reward
67              else:
68                  self.x += 1
69          elif actual_direction == "left":
70              if self.x==0 or (self.x-1,self.y) in self.vwalls:
71                  reward += self.crashed_reward
72              else:
73                  self.x += -1
74          elif actual_direction == "up":
75              if self.y==self.ydim-1:
76                  reward += self.crashed_reward
```

```
77              else:
78                  self.y += 1
79          elif actual_direction == "down":
80              if self.y==0:
81                  reward += self.crashed_reward
82              else:
83                  self.y += -1
84          else:
85              raise RuntimeError("unknown_direction "+str(direction))

87          # Monsters
88          if (self.x,self.y) in self.monster_locs and flip(self.monster_appears_prob):
89              if self.damaged:
90                  reward += self.monster_reward_when_damaged
91              else:
92                  self.damaged = True
93          if (self.x,self.y) in self.repair_stations:
94              self.damaged = False

96          # Prizes
97          if (self.x,self.y) == self.prize:
98              reward += self.prize_reward
99              self.prize = None

101         # Statistics
102         self.number_steps += 1
103         self.total_reward += reward
104         if self.total_reward < self.min_reward:
105             self.min_reward = self.total_reward
106             self.min_step = self.number_steps
107         if self.total_reward>0 and reward>self.total_reward:
108             self.zero_crossing = self.number_steps
109         self.display(2,"",self.number_steps,self.total_reward,
110                     self.total_reward/self.number_steps,sep="\t")

112         return (self.x, self.y, self.damaged, self.prize), reward
```

12.1.3   Evaluation and Plotting

_____ rlPlot.py — RL Plotter _____
```
11  import matplotlib.pyplot as plt
12
13  def plot_rl(ag, label=None, yplot='Total', step_size=None,
14              steps_explore=1000, steps_exploit=1000, xscale='linear'):
15      """
16      plots the agent ag
17      label is the label for the plot
18      yplot is 'Average' or 'Total'
19      step_size is the number of steps between each point plotted
```

```
20       steps_explore is the number of steps the agent spends exploring
21       steps_exploit is the number of steps the agent spends exploiting
22       xscale is 'log' or 'linear'
23
24       returns total reward when exploring, total reward when exploiting
25       """
26       assert yplot in ['Average','Total']
27       if step_size is None:
28           step_size = max(1,(steps_explore+steps_exploit)//500)
29       if label is None:
30           label = ag.label
31       ag.max_display_level,old_mdl = 1,ag.max_display_level
32       plt.ion()
33       plt.xscale(xscale)
34       plt.xlabel("step")
35       plt.ylabel(yplot+" reward")
36       steps = []        # steps
37       rewards = []      # return
38       ag.restart()
39       step = 0
40       while step < steps_explore:
41           ag.do(step_size)
42           step += step_size
43           steps.append(step)
44           if yplot == "Average":
45               rewards.append(ag.acc_rewards/step)
46           else:
47               rewards.append(ag.acc_rewards)
48       acc_rewards_exploring = ag.acc_rewards
49       ag.explore,explore_save = 0,ag.explore
50       while step < steps_explore+steps_exploit:
51           ag.do(step_size)
52           step += step_size
53           steps.append(step)
54           if yplot == "Average":
55               rewards.append(ag.acc_rewards/step)
56           else:
57               rewards.append(ag.acc_rewards)
58       plt.plot(steps,rewards,label=label)
59       plt.legend(loc="upper left")
60       plt.draw()
61       ag.max_display_level = old_mdl
62       ag.explore=explore_save
63       return acc_rewards_exploring, ag.acc_rewards-acc_rewards_exploring
```

12.2 Q Learning

> To run the Q-learning demo, in folder "aipython", load "rlQTest.py",
> and copy and paste the example queries at the bottom of that file. This
> assumes Python 3.

_____ rlQLearner.py — Q Learning _____

```python
11  import random
12  from utilities import Displayable, argmax, flip
13
14  class RL_agent(Displayable):
15      """An RL_Agent
16      has percepts (s, r) for some state s and real reward r
17      """
```

_____ rlQLearner.py — (continued) _____

```python
19  class Q_learner(RL_agent):
20      """A Q-learning agent has
21      belief-state consisting of
22          state is the previous state
23          q is a {(state,action):value} dict
24          visits is a {(state,action):n} dict. n is how many times action was done in state
25          acc_rewards is the accumulated reward
26
27      it observes (s, r) for some world-state s and real reward r
28      """
```

_____ rlQLearner.py — (continued) _____

```python
30      def __init__(self, env, discount, explore=0.1, fixed_alpha=True, alpha=0.2,
31                   alpha_fun=lambda k:1/k,
32                   qinit=0, label="Q_learner"):
33          """env is the environment to interact with.
34          discount is the discount factor
35          explore is the proportion of time the agent will explore
36          fixed_alpha specifies whether alpha is fixed or varies with the number of visits
37          alpha is the weight of new experiences compared to old experiences
38          alpha_fun is a function that computes alpha from the number of visits
39          qinit is the initial value of the Q's
40          label is the label for plotting
41          """
42          RL_agent.__init__(self)
43          self.env = env
44          self.actions = env.actions
45          self.discount = discount
46          self.explore = explore
47          self.fixed_alpha = fixed_alpha
48          self.alpha = alpha
49          self.alpha_fun = alpha_fun
```

```
50          self.qinit = qinit
51          self.label = label
52          self.restart()
```

restart is used to make the learner relearn everything. This is used by the plotter to create new plots.

```
_____rlQLearner.py — (continued) _____

54      def restart(self):
55          """make the agent relearn, and reset the accumulated rewards
56          """
57          self.acc_rewards = 0
58          self.state = self.env.state
59          self.q = {}
60          self.visits = {}
```

do takes in the number of steps.

```
_____rlQLearner.py — (continued) _____

62      def do(self,num_steps=100):
63          """do num_steps of interaction with the environment"""
64          self.display(2,"s\ta\tr\ts'\tQ")
65          alpha = self.alpha
66          for i in range(num_steps):
67              action = self.select_action(self.state)
68              next_state,reward = self.env.do(action)
69              if not self.fixed_alpha:
70                  k = self.visits[(self.state, action)] = self.visits.get((self.state, action),0)+1
71                  alpha = self.alpha_fun(k)
72              self.q[(self.state, action)] = (
73                  (1-alpha) * self.q.get((self.state, action),self.qinit)
74                  + alpha * (reward + self.discount
75                                      * max(self.q.get((next_state, next_act),self.qinit)
76                                          for next_act in self.actions)))
77              self.display(2,self.state, action, reward, next_state,
78                      self.q[(self.state, action)], sep='\t')
79              self.state = next_state
80              self.acc_rewards += reward
```

select_action us used to select the next action to perform. This can be reimplemented to give a different exploration strategy.

```
_____rlQLearner.py — (continued) _____

82      def select_action(self, state):
83          """returns an action to carry out for the current agent
84          given the state, and the q-function
85          """
86          if flip(self.explore):
87              return random.choice(self.actions)
88          else:
89              return argmax((next_act, self.q.get((state, next_act),self.qinit))
90                              for next_act in self.actions)
```

Exercise 12.1 Implement a soft-max action selection. Choose a temperature that works well for the domain. Explain how you picked this temperature. Compare the epsilon-greedy, soft-max and optimism in the face of uncertainty.

Exercise 12.2 Implement SARSA. Hint: it does not do a *max* in *do*. Instead it needs to choose *next_act* before it does the update.

12.2.1 Testing Q-learning

The first tests are for the 2-action 2-state

```
_____rlQTest.py — RL Q Tester _____
11  from rlProblem import Healthy_env
12  from rlQLearner import Q_learner
13  from rlPlot import plot_rl
14
15  env = Healthy_env()
16  ag = Q_learner(env, 0.7)
17  ag_opt = Q_learner(env, 0.7, qinit=100, label="optimistic" ) # optimistic agent
18  ag_exp_l = Q_learner(env, 0.7, explore=0.01, label="less explore")
19  ag_exp_m = Q_learner(env, 0.7, explore=0.5, label="more explore")
20  ag_disc = Q_learner(env, 0.9, qinit=100, label="disc 0.9")
21  ag_va = Q_learner(env, 0.7, qinit=100,fixed_alpha=False,alpha_fun=lambda k:10/(9+k),label="alpha=1
22
23  # ag.max_display_level = 2
24  # ag.do(20)
25  # ag.q   # get the learned q-values
26  # ag.max_display_level = 1
27  # ag.do(1000)
28  # ag.q   # get the learned q-values
29  # plot_rl(ag,yplot="Average")
30  # plot_rl(ag_opt,yplot="Average")
31  # plot_rl(ag_exp_l,yplot="Average")
32  # plot_rl(ag_exp_m,yplot="Average")
33  # plot_rl(ag_disc,yplot="Average")
34  # plot_rl(ag_va,yplot="Average")
35
36  from mdpExamples import mdpt
37  from rlProblem import Env_from_MDP
38  envt = Env_from_MDP(mdpt)
39  agt = Q_learner(envt, 0.8)
40  # agt.do(20)
41
42  from rlSimpleEnv import Simple_game_env
43  senv = Simple_game_env()
44  sag1 = Q_learner(senv,0.9,explore=0.2,fixed_alpha=True,alpha=0.1)
45  # plot_rl(sag1,steps_explore=100000,steps_exploit=100000,label="alpha="+str(sag1.alpha))
46  sag2 = Q_learner(senv,0.9,explore=0.2,fixed_alpha=False)
47  # plot_rl(sag2,steps_explore=100000,steps_exploit=100000,label="alpha=1/k")
48  sag3 = Q_learner(senv,0.9,explore=0.2,fixed_alpha=False,alpha_fun=lambda k:10/(9+k))
49  # plot_rl(sag3,steps_explore=100000,steps_exploit=100000,label="alpha=10/(9+k)")
```

12.3 Model-based Reinforcement Learner

> To run the demo, in folder "aipython", load "`rlModelLearner.py`", and
> copy and paste the example queries at the bottom of that file. This
> assumes Python 3.

A model-based reinforcement learner builds a Markov decision process model
of the domain, simultaneously learns the model and plans with that model.

The model-based reinforcement learner used the following data structures:

- $q[s,a]$ is dictionary that, given a (s,a) pair returns the Q-value, the estimate of the future (discounted) value of being in state s and doing action a.

- $r[s,a]$ is dictionary that, given a (s,a) pair returns the average reward from doing a in state s.

- $t[s,a,s']$ is dictionary that, given a (s,a,s') tuple returns the number of times a was done in state s, with the result being state s'.

- *visits*$[s,a]$ is dictionary that, given a (s,a) pair returns the number of times action a was carried out in state s.

- *res_states*$[s,a]$ is dictionary that, given a (s,a) pair returns the list of resulting states that have occurred when action a was carried out in state s. This is used in the asynchronous value iteration to determine the s' states to sum over.

- *visits_list* is a list of (s,a) pair that have been carried out. This is used to ensure there is no divide-by zero in the asynchronous value iteration. Note that this could be constructed from *r*, *visits* or *res_states* by enumerating the keys, but needs to be a list for *random.choice*, and we don't want to keep recreating it.

_____ rlModelLearner.py — Model-based Reinforcement Learner _____

```
11  import random
12  from rlQLearner import RL_agent
13  from utilities import Displayable, argmax, flip
14
15  class Model_based_reinforcement_learner(RL_agent):
16      """A Model-based reinforcement learner
17      """
18
19      def __init__(self, env, discount, explore=0.1, qinit=0,
20                   updates_per_step=10, label="MBR_learner"):
21          """env is the environment to interact with.
22          discount is the discount factor
23          explore is the proportion of time the agent will explore
```

```
24          qinit is the initial value of the Q's
25          updates_per_step is the number of AVI updates per action
26          label is the label for plotting
27          """
28          RL_agent.__init__(self)
29          self.env = env
30          self.actions = env.actions
31          self.discount = discount
32          self.explore = explore
33          self.qinit = qinit
34          self.updates_per_step = updates_per_step
35          self.label = label
36          self.restart()
```

_____rlModelLearner.py — (continued)_____

```
38      def restart(self):
39          """make the agent relearn, and reset the accumulated rewards
40          """
41          self.acc_rewards = 0
42          self.state = self.env.state
43          self.q = {}            # {(st,action):q_value} map
44          self.r = {}            # {(st,action):reward} map
45          self.t = {}            # {(st,action,st_next):count} map
46          self.visits = {}       # {(st,action):count} map
47          self.res_states = {}  # {(st,action):set_of_states} map
48          self.visits_list = [] # list of (st,action)
49          self.previous_action = None
```

_____rlModelLearner.py — (continued)_____

```
51      def do(self,num_steps=100):
52          """do num_steps of interaction with the environment
53          for each action, do updates_per_step iterations of asynchronous value iteration
54          """
55          for step in range(num_steps):
56              pst = self.state    # previous state
57              action = self.select_action(pst)
58              self.state,reward = self.env.do(action)
59              self.acc_rewards += reward
60              self.t[(pst,action,self.state)] = self.t.get((pst, action,self.state),0)+1
61              if (pst,action) in self.visits:
62                  self.visits[(pst,action)] += 1
63                  self.r[(pst,action)] += (reward-self.r[(pst,action)])/self.visits[(pst,action)]
64                  self.res_states[(pst,action)].add(self.state)
65              else:
66                  self.visits[(pst,action)] = 1
67                  self.r[(pst,action)] = reward
68                  self.res_states[(pst,action)] = {self.state}
69                  self.visits_list.append((pst,action))
70              st,act = pst,action    #initial state-action pair for AVI
71              for update in range(self.updates_per_step):
```

```
72              self.q[(st,act)] = self.r[(st,act)]+self.discount*(
73                  sum(self.t[st,act,rst]/self.visits[st,act]*
74                      max(self.q.get((rst,nact),self.qinit) for nact in self.actions)
75                      for rst in self.res_states[(st,act)]))
76              st,act = random.choice(self.visits_list)
```

──────────── rlModelLearner.py — (continued) ────────────

```
78      def select_action(self, state):
79          """returns an action to carry out for the current agent
80          given the state, and the q-function
81          """
82          if flip(self.explore):
83              return random.choice(self.actions)
84          else:
85              return argmax((next_act, self.q.get((state, next_act),self.qinit))
86                              for next_act in self.actions)
```

──────────── rlModelLearner.py — (continued) ────────────

```
88  from rlQTest import senv # simple game environment
89  mbl1 = Model_based_reinforcement_learner(senv,0.9,updates_per_step=10)
90  # plot_rl(mbl1,steps_explore=100000,steps_exploit=100000,label="model-based(10)")
91  mbl2 = Model_based_reinforcement_learner(senv,0.9,updates_per_step=1)
92  # plot_rl(mbl2,steps_explore=100000,steps_exploit=100000,label="model-based(1)")
```

Exercise 12.3 If there was only one update per step, the algorithm can be made simpler and use less space. Explain how. Does it make it more efficient? Is it worthwhile having more than one update per step for the games implemented here?

Exercise 12.4 It is possible to implement the model-based reinforcement learner by replacing *q*, *r*, *visits*, *res_states* with a single dictionary that returns a tuple (q, r, v, tm) where *q*, *r* and *v* are numbers, and *tm* is a map from resulting states into counts. Does this make the algorithm easier to understand? Does this make the algorithm more efficient?

Exercise 12.5 If the states and the actions were mapped into integers, the dictionaries could be implemented more efficiently as arrays. This entails an extra step in specifying problems. Implement this for the simple game. Is it more efficient?

12.4   Reinforcement Learning with Features

> To run the demo, in folder "aipython", load "rlFeatures.py", and copy and paste the example queries at the bottom of that file. This assumes Python 3.

12.4.1   Representing Features

A feature is a function from state and action. To construct the features for a domain, we construct a function that takes a state and an action and returns the

list of all feature values for that state and action. This feature set is redesigned for each problem.

get_features(*state*, *action*) returns the feature values appropriate for the simple game.

```python
from rlSimpleEnv import Simple_game_env
from rlProblem import RL_env

def get_features(state,action):
    """returns the list of feature values for the state-action pair
    """
    assert action in Simple_game_env.actions
    (x,y,d,p) = state
    # f1: would go to a monster
    f1 = monster_ahead(x,y,action)
    # f2: would crash into wall
    f2 = wall_ahead(x,y,action)
    # f3: action is towards a prize
    f3 = towards_prize(x,y,action,p)
    # f4: damaged and action is toward repair station
    f4 = towards_repair(x,y,action) if d else 0
    # f5: damaged and towards monster
    f5 = 1 if d and f1 else 0
    # f6: damaged
    f6 = 1 if d else 0
    # f7: not damaged
    f7 = 1-f6
    # f8: damaged and prize ahead
    f8 = 1 if d and f3 else 0
    # f9: not damaged and prize ahead
    f9 = 1 if not d and f3 else 0
    features = [1,f1,f2,f3,f4,f5,f6,f7,f8,f9]
    for pr in Simple_game_env.prize_locs+[None]:
        if p==pr:
            features += [x, 4-x, y, 4-y]
        else:
            features += [0, 0, 0, 0]
    # fp04 feature for y when prize is at 0,4
    # this knows about the wall to the right of the prize
    if p==(0,4):
        if x==0:
            fp04 = y
        elif y<3:
            fp04 = y
        else:
            fp04 = 4-y
    else:
        fp04 = 0
    features.append(fp04)
```

```
55        return features
56
57   def monster_ahead(x,y,action):
58       """returns 1 if the location expected to get to by doing
59       action from (x,y) can contain a monster.
60       """
61       if action == "right" and (x+1,y) in Simple_game_env.monster_locs:
62           return 1
63       elif action == "left" and (x-1,y) in Simple_game_env.monster_locs:
64           return 1
65       elif action == "up" and (x,y+1) in Simple_game_env.monster_locs:
66           return 1
67       elif action == "down" and (x,y-1) in Simple_game_env.monster_locs:
68           return 1
69       else:
70           return 0
71
72   def wall_ahead(x,y,action):
73       """returns 1 if there is a wall in the direction of action from (x,y).
74       This is complicated by the internal walls.
75       """
76       if action == "right" and (x==Simple_game_env.xdim-1 or (x,y) in Simple_game_env.vwalls):
77           return 1
78       elif action == "left" and (x==0 or (x-1,y) in Simple_game_env.vwalls):
79           return 1
80       elif action == "up" and y==Simple_game_env.ydim-1:
81           return 1
82       elif action == "down" and y==0:
83           return 1
84       else:
85           return 0
86
87   def towards_prize(x,y,action,p):
88       """action goes in the direction of the prize from (x,y)"""
89       if p is None:
90           return 0
91       elif p==(0,4): # take into account the wall near the top-left prize
92           if action == "left" and (x>1 or x==1 and y<3):
93               return 1
94           elif action == "down" and (x>0 and y>2):
95               return 1
96           elif action == "up" and (x==0 or y<2):
97               return 1
98           else:
99               return 0
100      else:
101          px,py = p
102          if p==(4,4) and x==0:
103              if (action=="right" and y<3) or (action=="down" and y>2) or (action=="up" and y<2):
104                  return 1
```

```
105              else:
106                  return 0
107          if (action == "up" and y<py) or (action == "down" and py<y):
108              return 1
109          elif (action == "left" and px<x) or (action == "right" and x<px):
110              return 1
111          else:
112              return 0
113
114  def towards_repair(x,y,action):
115      """returns 1 if action is towards the repair station.
116      """
117      if action == "up" and (x>0 and y<4 or x==0 and y<2):
118          return 1
119      elif action == "left" and x>1:
120          return 1
121      elif action == "right" and x==0 and y<3:
122          return 1
123      elif action == "down" and x==0 and y>2:
124          return 1
125      else:
126          return 0
127
128  def simp_features(state,action):
129      """returns a list of feature values for the state-action pair
130      """
131      assert action in Simple_game_env.actions
132      (x,y,d,p) = state
133      # f1: would go to a monster
134      f1 = monster_ahead(x,y,action)
135      # f2: would crash into wall
136      f2 = wall_ahead(x,y,action)
137      # f3: action is towards a prize
138      f3 = towards_prize(x,y,action,p)
139      return [1,f1,f2,f3]
```

12.4.2 Feature-based RL learner

This learns a linear function approximation of the Q-values. It requires the function *get_features* that given a state and an action returns a list of values for all of the features. Each environment requires this function to be provided.

rlFeatures.py — Feature-based Reinforcement Learner

```
11  import random
12  from rlQLearner import RL_agent
13  from utilities import Displayable, argmax, flip
14
15  class SARSA_LFA_learner(RL_agent):
16      """A SARSA_LFA learning agent has
17      belief-state consisting of
```

```
18          state is the previous state
19          q is a {(state,action):value} dict
20          visits is a {(state,action):n} dict. n is how many times action was done in state
21          acc_rewards is the accumulated reward
22
23      it observes (s, r) for some world-state s and real reward r
24      """
25      def __init__(self, env, get_features, discount, explore=0.2, step_size=0.01,
26                  winit=0, label="SARSA_LFA"):
27          """env is the feature environment to interact with
28          get_features is a function get_features(state,action) that returns the list of feature values
29          discount is the discount factor
30          explore is the proportion of time the agent will explore
31          step_size is gradient descent step size
32          winit is the initial value of the weights
33          label is the label for plotting
34          """
35          RL_agent.__init__(self)
36          self.env = env
37          self.get_features = get_features
38          self.actions = env.actions
39          self.discount = discount
40          self.explore = explore
41          self.step_size = step_size
42          self.winit = winit
43          self.label = label
44          self.restart()
```

restart() is used to make the learner relearn everything. This is used by the plotter to create new plots.

———————————————— rlFeatures.py — (continued) ————————————————

```
46      def restart(self):
47          """make the agent relearn, and reset the accumulated rewards
48          """
49          self.acc_rewards = 0
50          self.state = self.env.state
51          self.features = self.get_features(self.state, list(self.env.actions)[0])
52          self.weights = [self.winit for f in self.features]
53          self.action = self.select_action(self.state)
```

do takes in the number of steps.

———————————————— rlFeatures.py — (continued) ————————————————

```
55      def do(self,num_steps=100):
56          """do num_steps of interaction with the environment"""
57          self.display(2,"s\ta\tr\ts'\tQ\tdelta")
58          for i in range(num_steps):
59              next_state,reward = self.env.do(self.action)
60              self.acc_rewards += reward
61              next_action = self.select_action(next_state)
62              feature_values = self.get_features(self.state,self.action)
```

```
63                 oldQ = dot_product(self.weights, feature_values)
64                 nextQ = dot_product(self.weights, self.get_features(next_state,next_action))
65                 delta = reward + self.discount * nextQ - oldQ
66                 for i in range(len(self.weights)):
67                     self.weights[i] += self.step_size * delta * feature_values[i]
68                 self.display(2,self.state, self.action, reward, next_state,
69                             dot_product(self.weights, feature_values), delta, sep='\t')
70                 self.state = next_state
71                 self.action = next_action
72
73         def select_action(self, state):
74             """returns an action to carry out for the current agent
75             given the state, and the q-function.
76             This implements an epsilon-greedy approach
77             where self.explore is the probability of exploring.
78             """
79             if flip(self.explore):
80                 return random.choice(self.actions)
81             else:
82                 return argmax((next_act, dot_product(self.weights,
83                                                 self.get_features(state,next_act)))
84                             for next_act in self.actions)
85
86         def show_actions(self,state=None):
87             """prints the value for each action in a state.
88             This may be useful for debugging.
89             """
90             if state is None:
91                 state = self.state
92             for next_act in self.actions:
93                 print(next_act,dot_product(self.weights, self.get_features(state,next_act)))
94
95  def dot_product(l1,l2):
96      return sum(e1*e2 for (e1,e2) in zip(l1,l2))
```

Test code:

```
──────────────────── rlFeatures.py — (continued) ────────────────────
99   from rlQTest import senv # simple game environment
100  from rlSimpleGameFeatures import get_features, simp_features
101  from rlPlot import plot_rl
102
103  fa1 = SARSA_LFA_learner(senv, get_features, 0.9, step_size=0.01)
104  #fa1.max_display_level = 2
105  #fa1.do(20)
106  #plot_rl(fa1,steps_explore=10000,steps_exploit=10000,label="SARSA_LFA(0.01)")
107  fas1 = SARSA_LFA_learner(senv, simp_features, 0.9, step_size=0.01)
108  #plot_rl(fas1,steps_explore=10000,steps_exploit=10000,label="SARSA_LFA(simp)")
```

Exercise 12.6 How does the step-size affect performance? Try different step sizes
(e.g., 0.1, 0.001, other sizes in between). Explain the behaviour you observe. Which

step size works best for this example. Explain what evidence you are basing your prediction on.

Exercise 12.7 Does having extra features always help? Does it sometime help? Does whether it helps depend on the step size? Give evidence for your claims.

Exercise 12.8 For each of the following first predict, then plot, then explain the behaviour you observed:

(a) SARSA_LFA, Model-based learning (with 1 update per step) and Q-learning for 10,000 steps 20% exploring followed by 10,000 steps 100% exploiting

(b) SARSA_LFA, model-based learning and Q-learning for

   i) 100,000 steps 20% exploring followed by 100,000 steps 100% exploit

   ii) 10,000 steps 20% exploring followed by 190,000 steps 100% exploit

(c) Suppose your goal was to have the best accumulated reward after 200,000 steps. You are allowed to change the exploration rate at a fixed number of steps. For each of the methods, which is the best position to start exploiting more? Which method is better? What if you wanted to have the best reward after 10,000 or 1,000 steps?

Based on this evidence, explain when it is preferable to use SARSA_LFA, Model-based learner, or Q-learning.

Important: you need to run each algorithm more than once. Your explanation should include the variability as well as the typical behavior.

12.5   Learning to coordinate – UNFINISHED!!!!

Coordinating agents should implement the agent architecture. However, in that architecture, an agent calls the environment. That architecture was chosen because it was simple. However, it does not really work when there are multiple agents. In such cases, a coroutining architecture is more appropriate.

We assume there is an x-player, and a y-player. *game*[*xa*][*ya*][*ag*] gives value to the agent *ag* (ag=for the x-player) of the strategy of the x-agent doing *xa* and the y-agent doing *ya*.

```
_____learnCoordinate.py — Learning to Coordinate _____

11  from learnProblem import Learner
12
13  soccer = [[(-0.6,0.6),(-0.3,0.3)],[(-0.2,0.2),(-0.9,0.9)]]]
14  football = [[(2,1),(0,0)],[(0,0),(1,2)]]
15  prisoners_game = [[(100,100),(0,1100)],[(1100,0),(1000,1000)]]]
16
17  class Policy_hill_climbing(Learner):
18      def __init__(self,game)
```

Chapter 13

Relational Learning

13.1 Collaborative Filtering

Based on gradient descent algorithm of Koren, Y., Bell, R. and Volinsky, C.,
Matrix Factorization Techniques for Recommender Systems, IEEE Computer
2009.

This assumes the form of the dataset from movielens (`http://grouplens.`
`org/datasets/movielens/`). The rating are a set of ($user, item, rating, timestamp$)
tuples.

```
11  import random
12  import matplotlib.pyplot as plt
13  import urllib.request
14  from learnProblem import Learner
15  from utilities import Displayable
16
17  class CF_learner(Learner):
18      def __init__(self,
19                  rating_set,         # a Rating_set object
20                  rating_subset = None, # subset of ratings to be used as training ratings
21                  test_subset = None, # subset of ratings to be used as test ratings
22                  step_size = 0.01,   # gradient descent step size
23                  reglz = 1.0,        # the weight for the regularization terms
24                  num_properties = 10, # number of hidden properties
25                  property_range = 0.02 # properties are initialized to be between
26                                      # -property_range and property_range
27                  ):
28          self.rating_set = rating_set
29          self.ratings = rating_subset or rating_set.training_ratings # whichever is not empty
30          if test_subset is None:
```

```
31              self.test_ratings = self.rating_set.test_ratings
32          else:
33              self.test_ratings = test_subset
34          self.step_size = step_size
35          self.reglz = reglz
36          self.num_properties = num_properties
37          self.num_ratings = len(self.ratings)
38          self.ave_rating = (sum(r for (u,i,r,t) in self.ratings)
39                              /self.num_ratings)
40          self.users = {u for (u,i,r,t) in self.ratings}
41          self.items = {i for (u,i,r,t) in self.ratings}
42          self.user_bias = {u:0 for u in self.users}
43          self.item_bias = {i:0 for i in self.items}
44          self.user_prop = {u:[random.uniform(-property_range,property_range)
45                                for p in range(num_properties)]
46                             for u in self.users}
47          self.item_prop = {i:[random.uniform(-property_range,property_range)
48                                for p in range(num_properties)]
49                             for i in self.items}
50          self.zeros = [0 for p in range(num_properties)]
51          self.iter=0
52
53      def stats(self):
54          self.display(1,"ave sumsq error of mean for training=",
55                  sum((self.ave_rating-rating)**2 for (user,item,rating,timestamp)
56                     in self.ratings)/len(self.ratings))
57          self.display(1,"ave sumsq error of mean for test=",
58                  sum((self.ave_rating-rating)**2 for (user,item,rating,timestamp)
59                     in self.test_ratings)/len(self.test_ratings))
60          self.display(1,"error on training set",
61                     self.evaluate(self.ratings))
62          self.display(1,"error on test set",
63                     self.evaluate(self.test_ratings))
```

learn carries out *num_iter* steps of gradient descent.

_____ relnCollFilt.py — (continued) _____

```
65      def prediction(self,user,item):
66          """Returns prediction for this user on this item.
67          The use of .get() is to handle users or items not in the training set.
68          """
69          return (self.ave_rating
70                  + self.user_bias.get(user,0) #self.user_bias[user]
71                  + self.item_bias.get(item,0) #self.item_bias[item]
72                  + sum([self.user_prop.get(user,self.zeros)[p]*self.item_prop.get(item,self.zeros)[p]
73                        for p in range(self.num_properties)]))
74
75      def learn(self, num_iter = 50):
76          """ do num_iter iterations of gradient descent."""
77          for i in range(num_iter):
78              self.iter += 1
```

```
79              abs_error=0
80              sumsq_error=0
81              for (user,item,rating,timestamp) in random.sample(self.ratings,len(self.ratings)):
82                  error = self.prediction(user,item) - rating
83                  abs_error += abs(error)
84                  sumsq_error += error * error
85                  self.user_bias[user] -= self.step_size*error
86                  self.item_bias[item] -= self.step_size*error
87                  for p in range(self.num_properties):
88                      self.user_prop[user][p] -= self.step_size*error*self.item_prop[item][p]
89                      self.item_prop[item][p] -= self.step_size*error*self.user_prop[user][p]
90              for user in self.users:
91                  self.user_bias[user] -= self.step_size*self.reglz* self.user_bias[user]
92                  for p in range(self.num_properties):
93                      self.user_prop[user][p] -= self.step_size*self.reglz*self.user_prop[user][p]
94              for item in self.items:
95                  self.item_bias[item] -= self.step_size*self.reglz*self.item_bias[item]
96                  for p in range(self.num_properties):
97                      self.item_prop[item][p] -= self.step_size*self.reglz*self.item_prop[item][p]
98          self.display(1,"Iteration",self.iter,
99              "(Ave Abs,AveSumSq) training =",self.evaluate(self.ratings),
100             "test =",self.evaluate(self.test_ratings))
```

evaluate evaluates current predictions on the rating set:

```
_____ relnCollFilt.py — (continued) _____
102     def evaluate(self,ratings):
103         """returns (avergage_absolute_error, average_sum_squares_error) for ratings
104         """
105         abs_error = 0
106         sumsq_error = 0
107         if not ratings: return (0,0)
108         for (user,item,rating,timestamp) in ratings:
109             error = self.prediction(user,item) - rating
110             abs_error += abs(error)
111             sumsq_error += error * error
112         return abs_error/len(ratings), sumsq_error/len(ratings)
```

13.1.1 Alternative Formulation

An alternative formulation is to regularize after each update.

13.1.2 Plotting

```
_____ relnCollFilt.py — (continued) _____
114     def plot_predictions(self, examples="test"):
115         """
116         examples is either "test" or "training" or the actual examples
117         """
```

```
118         if examples == "test":
119             examples = self.test_ratings
120         elif examples == "training":
121             examples = self.ratings
122         plt.ion()
123         plt.xlabel("prediction")
124         plt.ylabel("cumulative proportion")
125         self.actuals = [[] for r in range(0,6)]
126         for (user,item,rating,timestamp) in examples:
127             self.actuals[rating].append(self.prediction(user,item))
128         for rating in range(1,6):
129             self.actuals[rating].sort()
130             numrat=len(self.actuals[rating])
131             yvals = [i/numrat for i in range(numrat)]
132             plt.plot(self.actuals[rating], yvals, label="rating="+str(rating))
133         plt.legend()
134         plt.draw()
```

This plots a single property. Each $(user, item, rating)$ is plotted where the x-value is the value of the property for the user, the y-value is the value of the property for the item, and the rating is plotted at this (x, y) position. That is, *rating* is plotted at the (x, y) position $(p(user), p(item))$.

```
                    relnCollFilt.py — (continued)
136     def plot_property(self,
137                   p,                # property
138                   plot_all=False, # true if all points should be plotted
139                   num_points=200 # number of random points plotted if not all
140                   ):
141         """plot some of the user-movie ratings,
142         if plot_all is true
143         num_points is the number of points selected at random plotted.
144
145         the plot has the users on the x-axis sorted by their value on property p and
146         with the items on the y-axis sorted by their value on property p and
147         the ratings plotted at the corresponding x-y position.
148         """
149         plt.ion()
150         plt.xlabel("users")
151         plt.ylabel("items")
152         user_vals = [self.user_prop[u][p]
153                     for u in self.users]
154         item_vals = [self.item_prop[i][p]
155                     for i in self.items]
156         plt.axis([min(user_vals)-0.02,
157                 max(user_vals)+0.05,
158                 min(item_vals)-0.02,
159                 max(item_vals)+0.05])
160         if plot_all:
161             for (u,i,r,t) in self.ratings:
162                 plt.text(self.user_prop[u][p],
```

```
163                           self.item_prop[i][p],
164                           str(r))
165             else:
166                 for i in range(num_points):
167                     (u,i,r,t) = random.choice(self.ratings)
168                     plt.text(self.user_prop[u][p],
169                              self.item_prop[i][p],
170                              str(r))
171         plt.show()
```

13.1.3 Creating Rating Sets

A rating set can be read from the Internet or read from a local file. The default is to read the Movielens 100K dataset from the Internet. It would be more efficient to save the dataset as a local file, and then set *local_file = True*, as then it will not need to download the dataset every time the program is run.

—————————————relnCollFilt.py — (continued) —————————————

```
173  class Rating_set(Displayable):
174      def __init__(self,
175                   date_split=892000000,
176                   local_file=False,
177                   url="http://files.grouplens.org/datasets/movielens/ml-100k/u.data",
178                   file_name="u.data"):
179          self.display(1,"reading...")
180          if local_file:
181              lines = open(file_name,'r')
182          else:
183              lines = (line.decode('utf-8') for line in urllib.request.urlopen(url))
184          all_ratings = (tuple(int(e) for e in line.strip().split('\t'))
185                          for line in lines)
186          self.training_ratings = []
187          self.training_stats = {1:0, 2:0, 3:0, 4:0 ,5:0}
188          self.test_ratings = []
189          self.test_stats = {1:0, 2:0, 3:0, 4:0 ,5:0}
190          for rate in all_ratings:
191              if rate[3] < date_split: # rate[3] is timestamp
192                  self.training_ratings.append(rate)
193                  self.training_stats[rate[2]] += 1
194              else:
195                  self.test_ratings.append(rate)
196                  self.test_stats[rate[2]] += 1
197          self.display(1,"...read:", len(self.training_ratings),"training ratings and",
198                  len(self.test_ratings),"test ratings")
199          tr_users = {user for (user,item,rating,timestamp) in self.training_ratings}
200          test_users = {user for (user,item,rating,timestamp) in self.test_ratings}
201          self.display(1,"users:",len(tr_users),"training,",len(test_users),"test,",
202                      len(tr_users & test_users),"in common")
203          tr_items = {item for (user,item,rating,timestamp) in self.training_ratings}
204          test_items = {item for (user,item,rating,timestamp) in self.test_ratings}
205          self.display(1,"items:",len(tr_items),"training,",len(test_items),"test,",
```

```
206                     len(tr_items & test_items),"in common")
207           self.display(1,"Rating statistics for training set: ",self.training_stats)
208           self.display(1,"Rating statistics for test set: ",self.test_stats)
```

Sometimes it is useful to plot a property for all (*user*, *item*, *rating*) triples. There are too many such triples in the data set. The method *create_top_subset* creates a much smaller dataset where this makes sense. It picks the most rated items, then picks the users who have the most ratings on these items. It is designed for depicting the meaning of properties, and may not be useful for other purposes.

———————————— relnCollFilt.py — (continued) ————————————

```
210     def create_top_subset(self, num_items = 30, num_users = 30):
211         """Returns a subset of the ratings by picking the most rated items,
212         and then the users that have most ratings on these, and then all of the
213         ratings that involve these users and items.
214         """
215         items = {item for (user,item,rating,timestamp) in self.training_ratings}
216
217         item_counts = {i:0 for i in items}
218         for (user,item,rating,timestamp) in self.training_ratings:
219             item_counts[item] += 1
220
221         items_sorted = sorted((item_counts[i],i) for i in items)
222         top_items = items_sorted[-num_items:]
223         set_top_items = set(item for (count, item) in top_items)
224
225         users = {user for (user,item,rating,timestamp) in self.training_ratings}
226         user_counts = {u:0 for u in users}
227         for (user,item,rating,timestamp) in self.training_ratings:
228             if item in set_top_items:
229                 user_counts[user] += 1
230
231         users_sorted = sorted((user_counts[u],u)
232                               for u in users)
233         top_users = users_sorted[-num_users:]
234         set_top_users = set(user for (count, user) in top_users)
235         used_ratings = [ (user,item,rating,timestamp)
236                         for (user,item,rating,timestamp) in self.training_ratings
237                         if user in set_top_users and item in set_top_items]
238         return used_ratings
239
240 movielens = Rating_set()
241 learner0 = CF_learner(movielens, num_properties = 1)
242 #learner0.learn(50)
243 # learner0.plot_predictions(examples = "training")
244 # learner0.plot_predictions(examples = "test")
245 #learner0.plot_property(0)
246 #movielens_subset = movielens.create_top_subset(num_items = 20, num_users = 20)
247 #learner1 = CF_learner(movielens, rating_subset=movielens_subset, test_subset=[], num_properties=1)
248 #learner1.learn(1000)
```

```
249  #learner1.plot_property(0,plot_all=True)
```

Index